

Data Visualization in R Programming: A Practical Introduction

Nith Kosal

May 12, 2021

Data Visualization Notes

This is an outcome in PDF file from RMarkdown project template to accompany training with *Data Visualization*. You can use it to take notes, write your code, and produce a good-looking, reproducible document that records the work you have done. At the very top of the file is a section of *metadata*, or information about what the file is and what it does. The metadata is delimited by three dashes at the start and another three at the end. You should change the title, author, and date to the values that suit you. Keep the **output** line as it is for now, however. Each line in the metadata has a structure. First the *key* (“title”, “author”, etc), then a colon, and then the *value* associated with the key.

This Document is an RMarkdown File

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. A *code chunk* is a specially delimited section of the file. You can add one by moving the cursor to a blank line choosing Code > Insert Chunk from the RStudio menu. When you do, an empty chunk will appear in your file.

Code chunks are delimited by three backticks (found to the left of the 1 key on US and UK keyboards) at the start and end. The opening backticks also have a pair of braces and the letter **r**, to indicate what language the chunk is written in. You write your code inside the code chunks. Write your notes and other material around them, as here.

Things to know about R

Any new piece of software takes a bit of getting used to. This is especially true when using an IDE to work in a language like R.

Everything has a name

In R, everything you deal with has a name. You refer to things by their names as you examine, use, or modify them. Named entities include variables (like **x**, or **y**), data that you have loaded (like **my_data**), and functions that you use. (More about functions momentarily.) You will spend a lot of time talking about, creating, referring to, and modifying things with names.

Some names are forbidden. These include reserved words like **FALSE** and **TRUE**, core programming words like **Inf**, **for**, **else**, **break**, **function**, and words for special entities like **NA** and **NaN**. (These last two are codes designating missing data and “Not a Number”, respectively.) You probably won’t use these names by accident, but it’s good to know that they are not allowed.

Some names you should not use, even if they are technically permitted. These are mostly words that are already in use for objects or functions that form part of the core of R. These include the names of

basic functions like `q()` or `c()`, common statistical functions like `mean()`, `range()` or `var()`, and built-in mathematical constants like `pi`.

Names in R are case sensitive, not much difficult from others programming languages. The object `my_data` is not the same as the object `My_Data`. When choosing names for things, be concise, consistent, and informative. Follow the style of the tidyverse and name things in lower case, separating words with the underscore character, `_`, as needed. Do not use spaces when naming things, including variables in your data.

Everything is an object

Some objects are built in to R, some are added via libraries, and some are created by the user. But almost everything is some kind of object. The code you write will create, manipulate, and use named objects as a matter of course. We can start immediately. Let's create a vector of numbers. The command `c()` is a function. It's short for "combine" or "concatenate". It will take a sequence of comma-separated things inside the parentheses and join them together into a vector where each element is still individually accessible.

```
c(1, 2, 3, 1, 3, 2020, 2021)

## [1] 1 2 3 1 3 2020 2021

# create object's name
my_numbers <- c(1, 2, 3, 1, 3, 2020, 2021)

futureforum_team <- c("Dara", "Soriya", "Sokhouy", "Kimly", "Heang", "Samnang")
y <- 1:9

# call for object that we want to see
futureforum_team

## [1] "Dara" "Soriya" "Sokhouy" "Kimly" "Heang" "Samnang"

# you do things using functions
x = my_numbers

# call for show function x
x

## [1] 1 2 3 1 3 2020 2021

# calculate mean for function x
mean(x)

## [1] 578.7143

# sum function x
sum(x)

## [1] 4051

# log in `x`
log(x)

## [1] 0.0000000 0.6931472 1.0986123 0.0000000 1.0986123 7.6108528 7.6113477

# exponential function `x`
exp(x)

## [1] 2.718282 7.389056 20.085537 2.718282 20.085537 Inf Inf

x^6

## [1] 1.000000e+00 6.400000e+01 7.290000e+02 1.000000e+00 7.290000e+02
```

```
## [6] 6.793729e+19 6.813933e+19
# sin and computes arc-sine or sine inverse of `x`
sin(x)

## [1] 0.84147098 0.90929743 0.14112001 0.84147098 0.14112001 0.04406199
## [7] -0.81684695

asin(x)

## Warning in asin(x): NaNs produced
## [1] 1.570796      NaN      NaN 1.570796      NaN      NaN      NaN
# summary x
x_summary <- summary(x)
x_summary

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.0     1.5     3.0   578.7  1011.5  2021.0
# functions come in libraries
table(my_numbers)

## my_numbers
##      1      2      3 2020 2021
##      2      1      2      1      1
# calculate standard deviation of `my_numbers`
sd(my_numbers)

## [1] 984.9275
# `my_numbers` multiply with 2021
my_numbers * 2021

## [1]      2021      4042      6063      2021      6063 4082420 4084441
# `my_numbers` plus 1
my_numbers + 1

## [1]      2      3      4      2      4 2021 2022
# `my_numbers` divide with 12
my_numbers / 12

## [1] 0.08333333 0.16666667 0.25000000 0.08333333 0.25000000
## [6] 168.33333333 168.41666667
# `my_numbers` plus `my_numbers`
my_numbers + my_numbers

## [1]      2      4      6      2      6 4040 4042
# if you're not sure what an object is, ask for its class
class(my_numbers)

## [1] "numeric"
class(futureforum_team)

## [1] "character"
```

```

class(x_summary)

## [1] "summaryDefault" "table"

class(summary)

## [1] "function"

my_new_vector <- c(my_numbers, "Apple")
my_new_vector

## [1] "1"      "2"      "3"      "1"      "3"      "2020"   "2021"   "Apple"

class(my_new_vector)

## [1] "character"

```

Load Libraries

To begin we must load some libraries we will be using. If we do not load them, R will not be able to find the functions contained in these libraries. But for the first time with *R Programming*, you should install library packages that common use.

Here, the braces at the start of the code chunk have some additional options set in them. There is the language, `r`, as before. This is required. Then there is the word `setup`, which is a label for your code chunk. Labels are useful to briefly say what the chunk does. Label names must be unique (no two chunks in the same document can have the same label) and cannot contain spaces. Then, after the comma, an option is set: `include=FALSE`. This tells R to run this code but not to include the output in the final document.

If you have not installed these required libraries yet, make sure you have an internet connection and install them now.

```

# to install these packages, in the line above.

# basic method that common use for new learner
install.packages("tidyverse")
install.packages("broom")

# advanced method
my_packages <- c("tidyverse", "broom", "coefplot", "cowplot", "drat", "fs",
                "gapminder", "GGally", "ggrepel", "ggribbles", "gridExtra",
                "here", "interplot", "margins", "maps", "mapproj",
                "mapdata", "MASS", "quantreg", "rlang", "scales",
                "survey", "srvyr", "viridis", "viridisLite", "devtools")

# R Studio should then download and install these packages for you.
install.packages(my_packages, repos = "http://cran.rstudio.com")

# load libraries for use
# basic method
library(gapminder)
library(here)
library(tidyverse)

# advanced method
Packages <- my_packages

```

```

Packages <- c("tidyverse", "broom", "coefplot", "cowplot", "drat", "fs",
              "gapminder", "maps", "mapproj", "survey", "srvyr", "viridis",
              "viridisLite", "devtools")

lapply(Packages, library, character.only = TRUE)

# to update our package
update.packages()

# to know which pack need an update
old.packages()

# to know which packages are being loaded
search()

# request for help
?hist
help(package = "tidyverse")
example("hist")

# use ?? to search by keyword
??regression

```

So let we go with some datasets on R

```

# call the dataset from R
mtcars

```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1

```

## Porsche 914-2      26.0   4 120.3  91 4.43 2.140 16.70  0  1   5   2
## Lotus Europa      30.4   4  95.1 113 3.77 1.513 16.90  1  1   5   2
## Ford Pantera L    15.8   8 351.0 264 4.22 3.170 14.50  0  1   5   4
## Ferrari Dino      19.7   6 145.0 175 3.62 2.770 15.50  0  1   5   6
## Maserati Bora     15.0   8 301.0 335 3.54 3.570 14.60  0  1   5   8
## Volvo 142E        21.4   4 121.0 109 4.11 2.780 18.60  1  1   4   2

mymtcars <- mtcars

class(mymtcars)

## [1] "data.frame"

# list the variables in `mymtcars`
names(mymtcars)

## [1] "mpg" "cyl" "disp" "hp" "drat" "wt" "qsec" "vs" "am" "gear"
## [11] "carb"

# or list objects in the working environment
ls(mymtcars)

## [1] "am" "carb" "cyl" "disp" "drat" "gear" "hp" "mpg" "qsec" "vs"
## [11] "wt"

# list the structure of `mymtcars`
str(mymtcars)

## 'data.frame': 32 obs. of 11 variables:
## $ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
## $ disp: num 160 160 108 258 360 ...
## $ hp : num 110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt : num 2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num 16.5 17 18.6 19.4 17 ...
## $ vs : num 0 0 1 1 0 1 0 1 1 1 ...
## $ am : num 1 1 1 0 0 0 0 0 0 0 ...
## $ gear: num 4 4 4 3 3 3 3 4 4 4 ...
## $ carb: num 4 4 1 1 2 1 4 2 2 4 ...

# print first 10 rows of `mymtcars`
head(mymtcars, n=7)

##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4    21.0   6  160 110 3.90 2.620 16.46  0  1    4    4
## Mazda RX4 Wag 21.0   6  160 110 3.90 2.875 17.02  0  1    4    4
## Datsun 710    22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive 21.4   6  258 110 3.08 3.215 19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0    3    2
## Valiant      18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
## Duster 360    14.3   8  360 245 3.21 3.570 15.84  0  0    3    4

# print last 5 rows of `mymtcars`
tail(mymtcars, n=4)


##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Ford Pantera L 15.8   8  351 264 4.22 3.17 14.5  0  1    5    4
## Ferrari Dino   19.7   6  145 175 3.62 2.77 15.5  0  1    5    6

```

```
## Maserati Bora 15.0 8 301 335 3.54 3.57 14.6 0 1 5 8
## Volvo 142E 21.4 4 121 109 4.11 2.78 18.6 1 1 4 2
```

we can also create plots:

```
plot(mymtcars$mpg, mtcars$disp)
```



Now, we try with another example and use the dataset from link

```
url <- "https://cdn.rawgit.com/kjhealy/viz-organdata/master/organdonation.csv"
organs <- read_csv(file = url)
```

Or locally:

```
organs <- read.csv(file = "Data/organdonation.csv")
```

tibble package to show only the first 10 of the dataset

```
library(tibble)
```

```
## Warning: package 'tibble' was built under R version 3.6.3
```

```
as_tibble(organs)
```

```
## # A tibble: 238 x 21
##   country year donors pop pop.dens gdp gdp.lag health health.lag pubhealth
##   <fct>   <int> <dbl> <int>   <dbl> <int>   <int>   <dbl>       <dbl>      <dbl>
## 1 Austra~    NA    NA  17065  0.220 16774  16591   1300       1224      4.8
## 2 Austra~  1991  12.1  17284  0.223 17171  16774   1379       1300      5.4
## 3 Austra~  1992  12.4  17495  0.226 17914  17171   1455       1379      5.4
```

```
## 4 Austra~ 1993 12.5 17667 0.228 18883 17914 1540 1455 5.4
## 5 Austra~ 1994 10.2 17855 0.231 19849 18883 1626 1540 5.4
## 6 Austra~ 1995 10.2 18072 0.233 21079 19849 1737 1626 5.5
## 7 Austra~ 1996 10.6 18311 0.237 21923 21079 1846 1737 5.6
## 8 Austra~ 1997 10.3 18518 0.239 22961 21923 1948 1846 5.7
## 9 Austra~ 1998 10.5 18711 0.242 24148 22961 2077 1948 5.9
## 10 Austra~ 1999 8.67 18926 0.244 25445 24148 2231 2077 6.1
## # ... with 228 more rows, and 11 more variables: roads <dbl>, cerebvas <int>,
## # assault <int>, external <int>, txp.pop <dbl>, world <fct>, opt <fct>,
## # consent.law <fct>, consent.practice <fct>, consistent <fct>, ccode <fct>
```

Make your first figure

```
gapminder
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

```
p <- ggplot(data = gapminder)
```

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
                          y = lifeExp))
```

```
p
```

```
p + geom_point()
```

Build your plots layer by layer

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
                          y=lifeExp))
p + geom_smooth()
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
                          y=lifeExp))
p + geom_point() + geom_smooth()
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

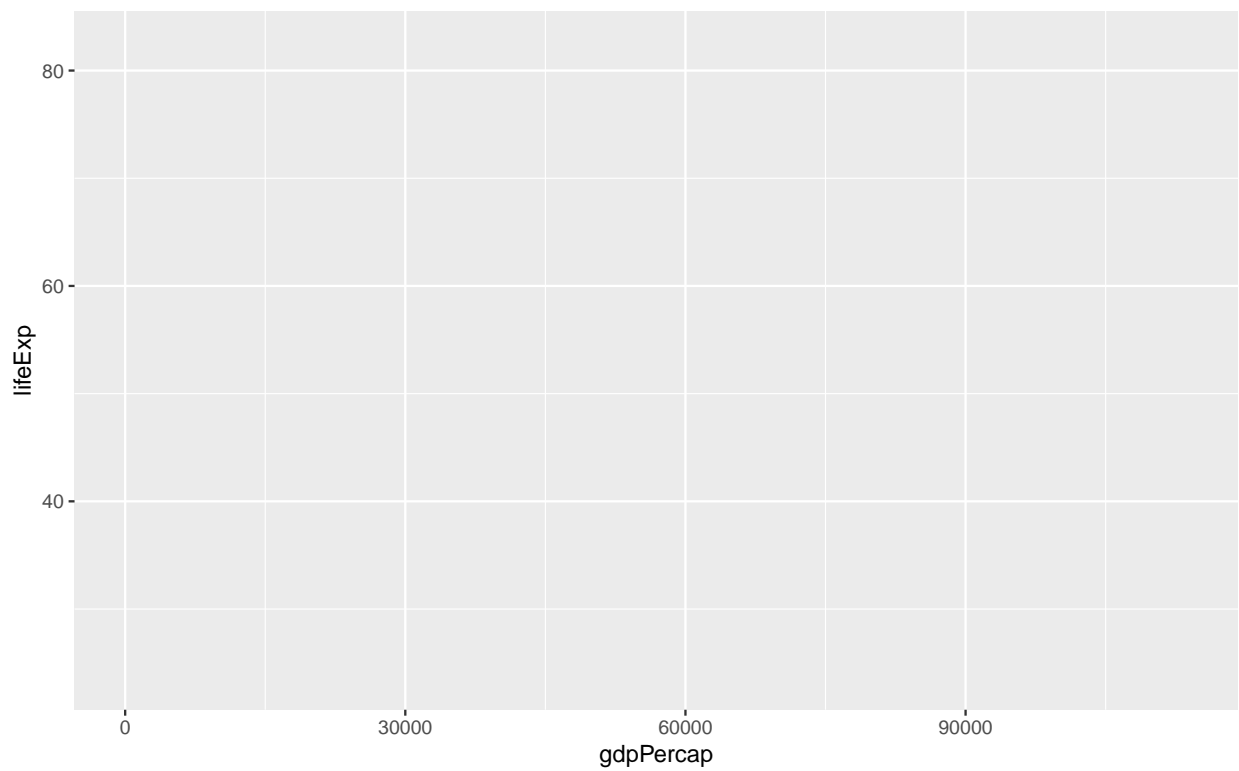



Figure 1: This empty plot has no geoms.

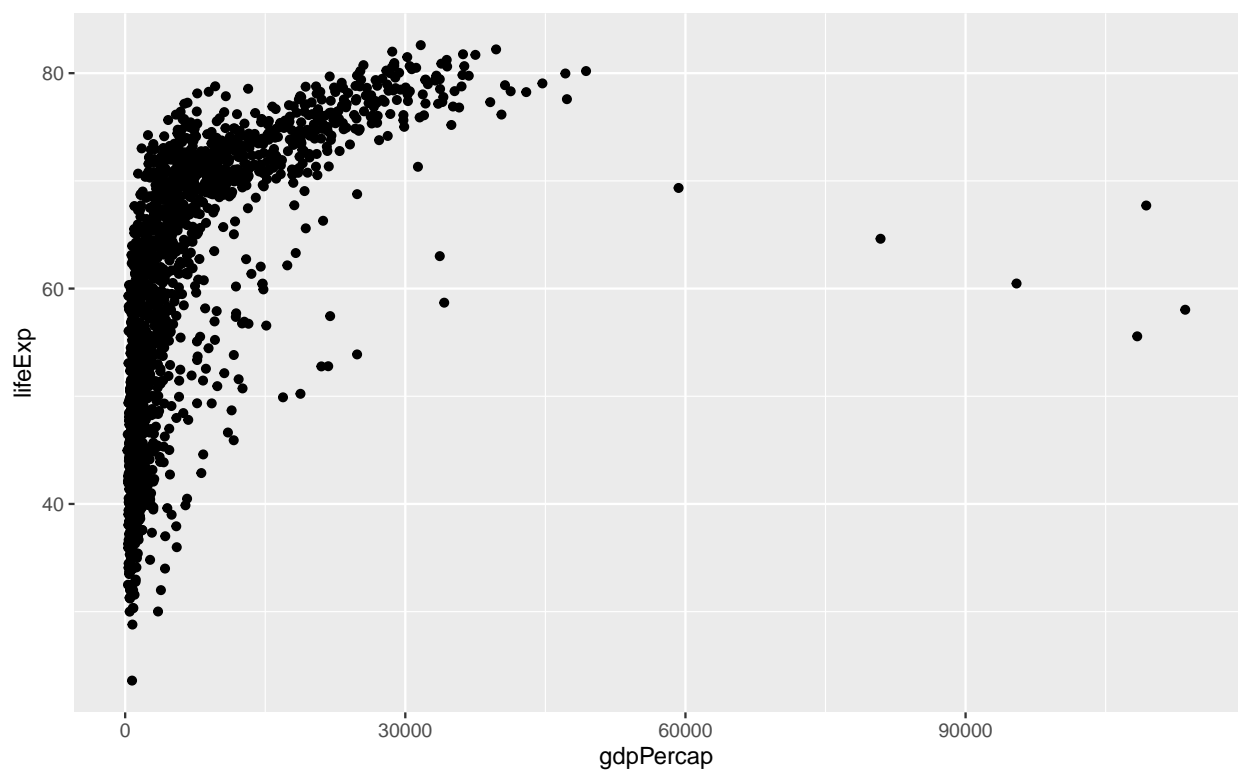


Figure 2: A scatterplot of Life Expectancy vs GDP

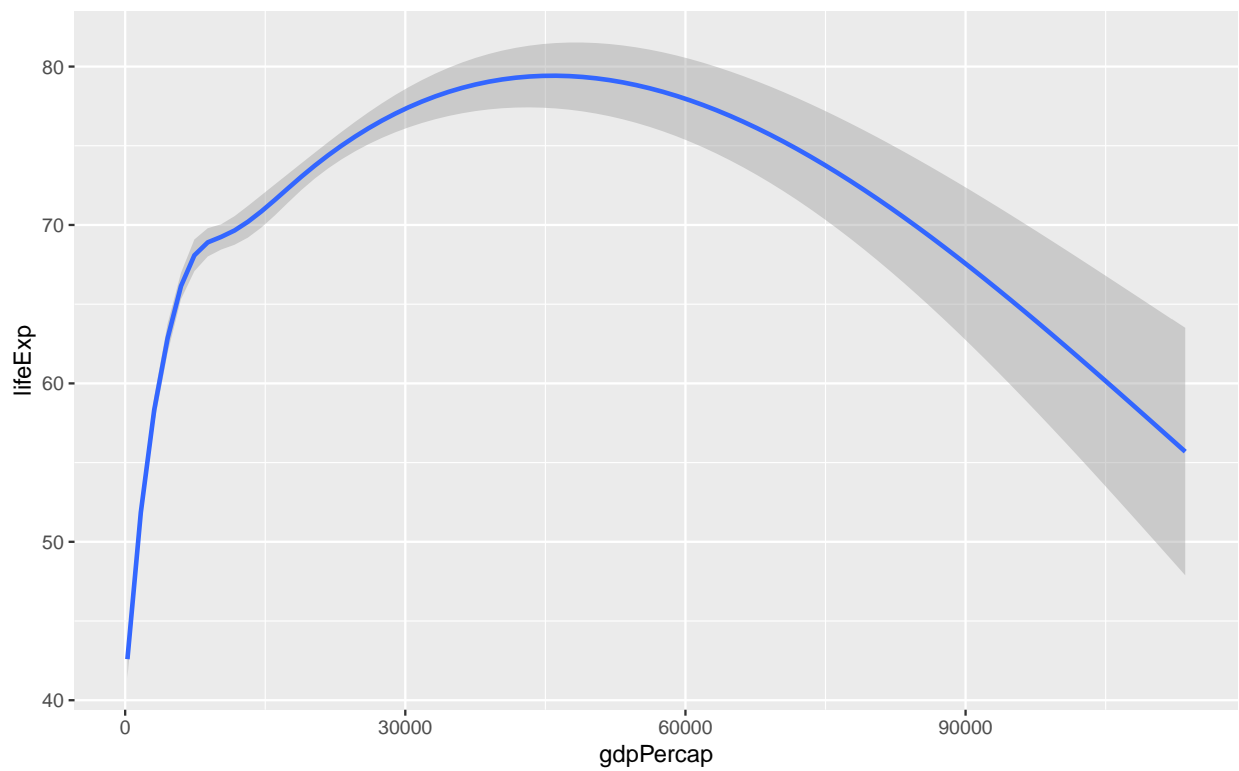


Figure 3: Life Expectancy vs GDP, using a smoother.

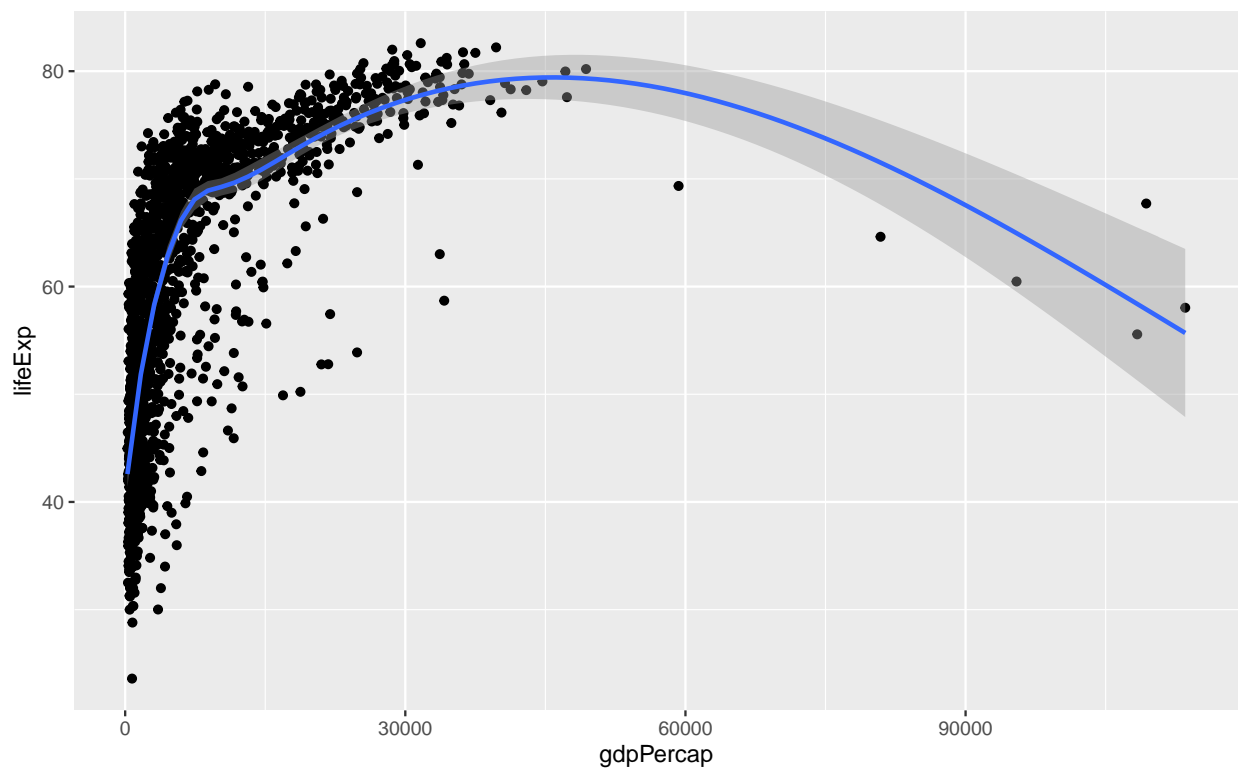


Figure 4: Life Expectancy vs GDP, showing both points and a GAM smoother.

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
                          y=lifeExp))
p + geom_point() + geom_smooth(method = "lm")

## `geom_smooth()` using formula 'y ~ x'
```




Figure 5: Life Expectancy vs GDP, points and an ill-advised linear fit.

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
                          y=lifeExp))
p + geom_point() +
  geom_smooth(method = "gam") +
  scale_x_log10()

## `geom_smooth()` using formula 'y ~ s(x, bs = "cs")'

p <- ggplot(data = gapminder, mapping = aes(x = gdpPercap, y=lifeExp))
p + geom_point() +
  geom_smooth(method = "gam") +
  scale_x_log10(labels = scales::dollar)

## `geom_smooth()` using formula 'y ~ s(x, bs = "cs")'
```

Mapping aesthetics vs setting them

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
```

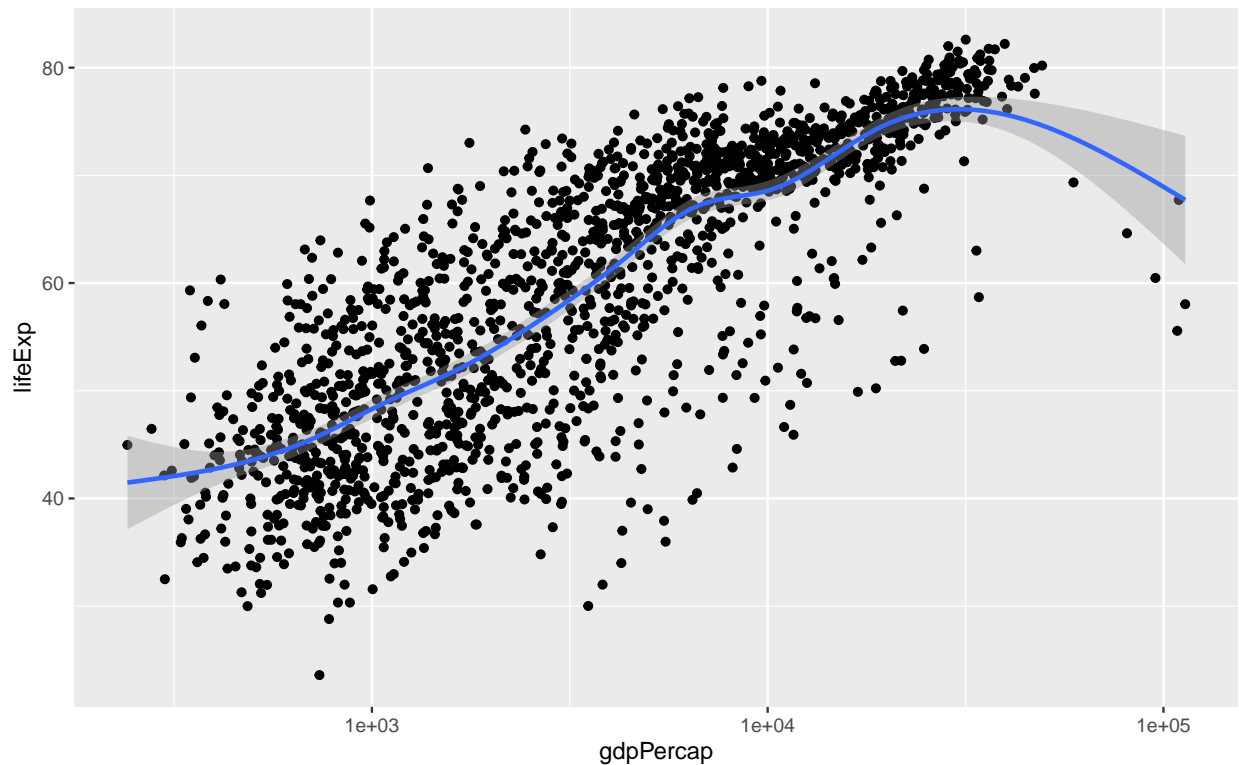


Figure 6: Life Expectancy vs GDP scatterplot, with a GAM smoother and a log scale on the x-axis.

```

                                y = lifeExp,
                                color = "purple"))
p + geom_point() +
  geom_smooth(method = "loess") +
  scale_x_log10()

## `geom_smooth()` using formula 'y ~ x'
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPerCap,
                          y = lifeExp))
p + geom_point(color = "purple") +
  geom_smooth(method = "loess") +
  scale_x_log10()

## `geom_smooth()` using formula 'y ~ x'
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPerCap,
                          y = lifeExp))
p + geom_point(alpha = 0.3) +
  geom_smooth(color = "orange", se = FALSE, size = 8, method = "lm") +
  scale_x_log10()

## `geom_smooth()` using formula 'y ~ x'

```

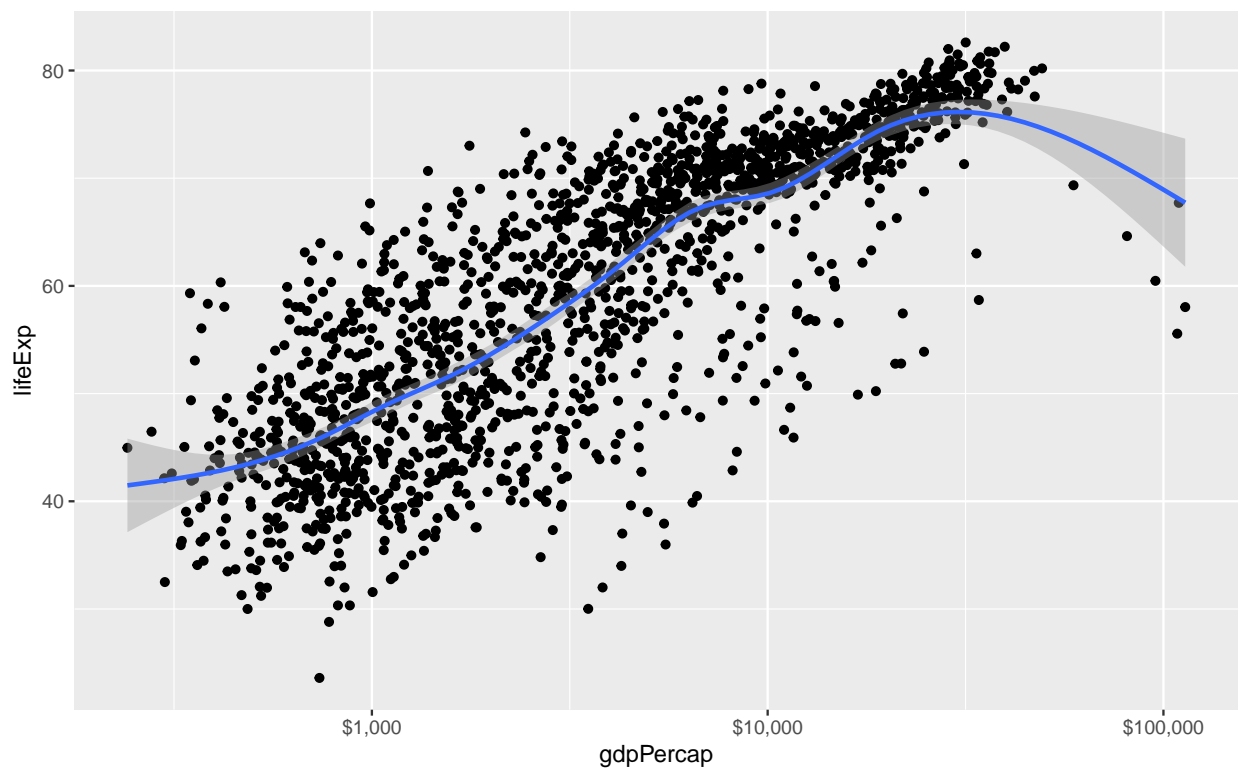


Figure 7: Life Expectancy vs GDP scatterplot, with a GAM smoother and a log scale on the x-axis, with better labels on the tick marks.




Figure 8: What has gone wrong here?




Figure 9: Setting the color attribute of the points directly.

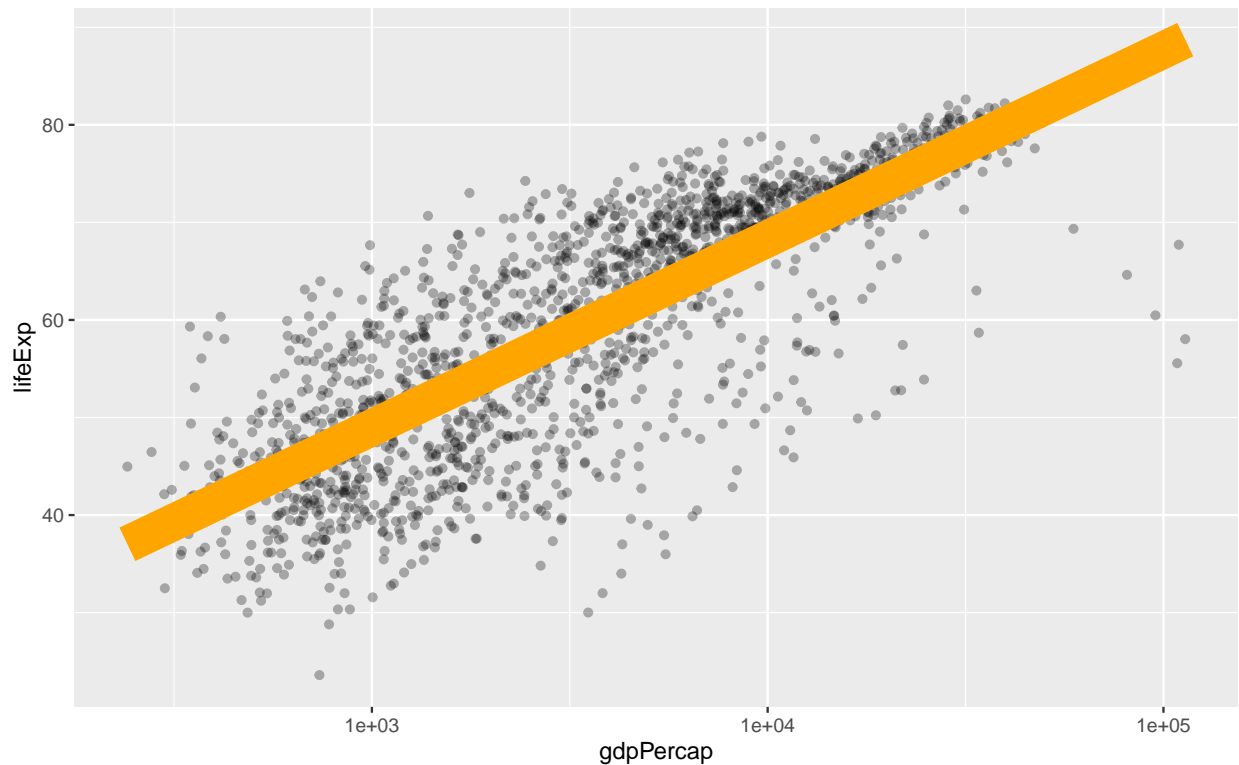


Figure 10: Setting some other arguments.

```
p <- ggplot(data = gapminder, mapping = aes(x = gdpPerCap, y=lifeExp))
p + geom_point(alpha = 0.3) + geom_smooth(method = "gam") +
  scale_x_log10(labels = scales::dollar) +
  labs(x = "GDP Per Capita", y = "Life Expectancy in Years",
       title = "Economic Growth and Life Expectancy",
       subtitle = "Data points are country-years",
       caption = "Source: Gapminder.")
```

```
## `geom_smooth()` using formula 'y ~ s(x, bs = "cs")'
```

```
p <- ggplot(data = gapminder,
           mapping = aes(x = gdpPerCap,
                         y = lifeExp,
                         color = continent))
p + geom_point() +
  geom_smooth(method = "loess") +
  scale_x_log10()
```

```
## `geom_smooth()` using formula 'y ~ x'
```

```
p <- ggplot(data = gapminder,
           mapping = aes(x = gdpPerCap,
                         y = lifeExp,
                         color = continent,
                         fill = continent))
p + geom_point() +
  geom_smooth(method = "loess") +
```




Figure 11: A more polished plot of Life Expectancy vs GDP.

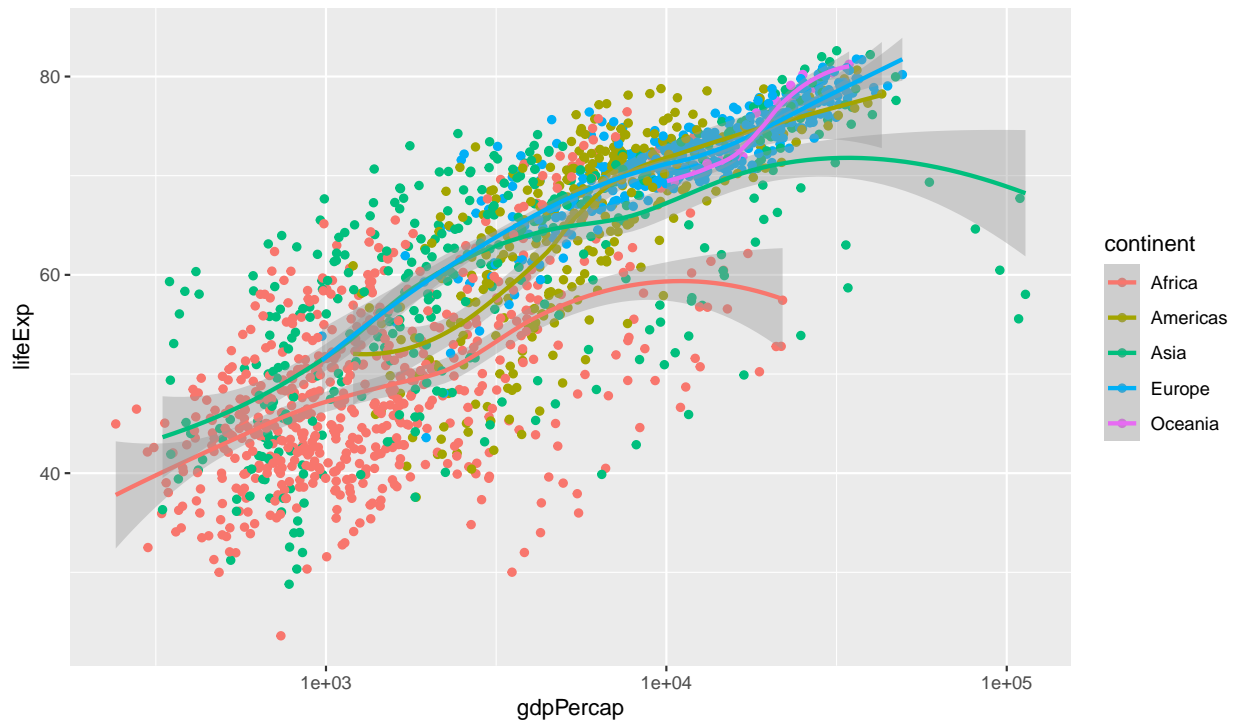


Figure 12: Mapping the continent variable to the color aesthetic.


```
scale_x_log10()
```

```
## `geom_smooth()` using formula 'y ~ x'
```




Figure 13: Mapping the continent variable to the color aesthetic, and correcting the error bars using the fill aesthetic.

Aesthetics can be mapped per geom

```
p <- ggplot(data = gapminder, mapping = aes(x = gdpPercap, y = lifeExp))
p + geom_point(mapping = aes(color = continent)) +
  geom_smooth(method = "loess") +
  scale_x_log10()
```

```
## `geom_smooth()` using formula 'y ~ x'
```

```
p <- ggplot(data = gapminder,
            mapping = aes(x = gdpPercap,
                          y = lifeExp))
p + geom_point(mapping = aes(color = log(pop))) +
  scale_x_log10()
```

Save your work

```
knitr::opts_chunk$set(fig.width=8, fig.height=5)
```

```
ggsave(filename = "Figure/figure1.png")
```

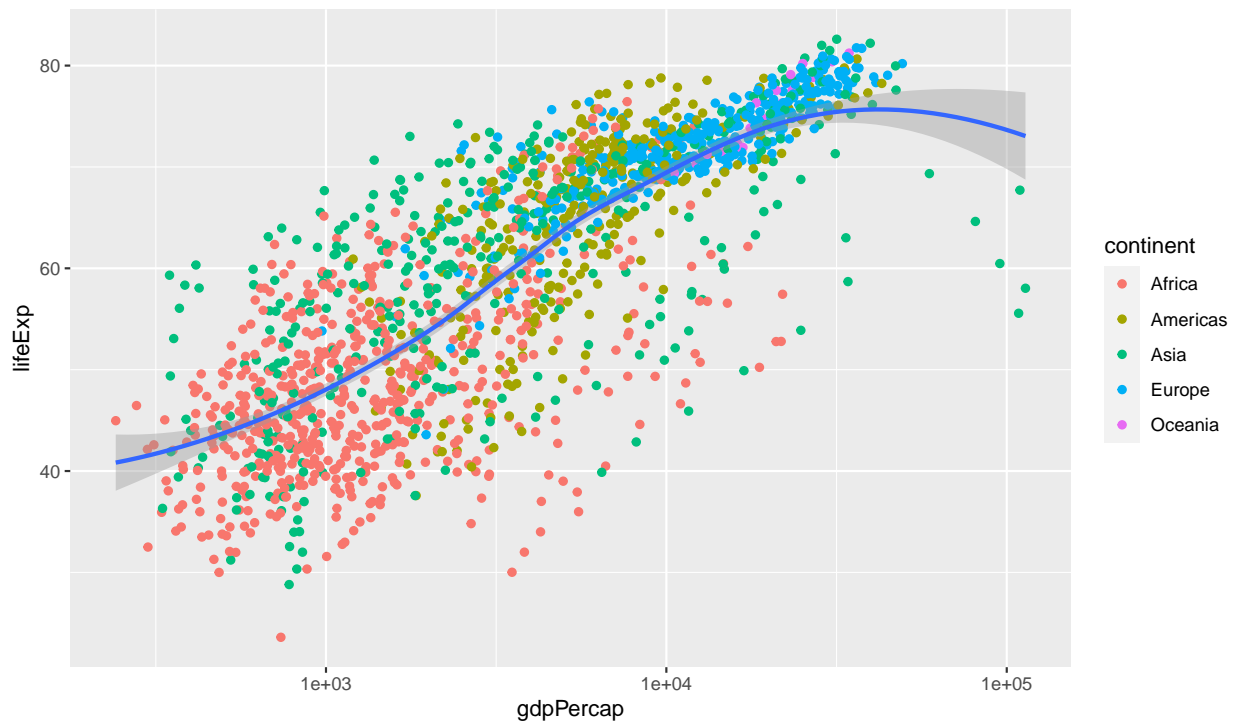


Figure 14: Mapping aesthetics on a per-geom basis. Here color is mapped to continent for the points but not the smoother.

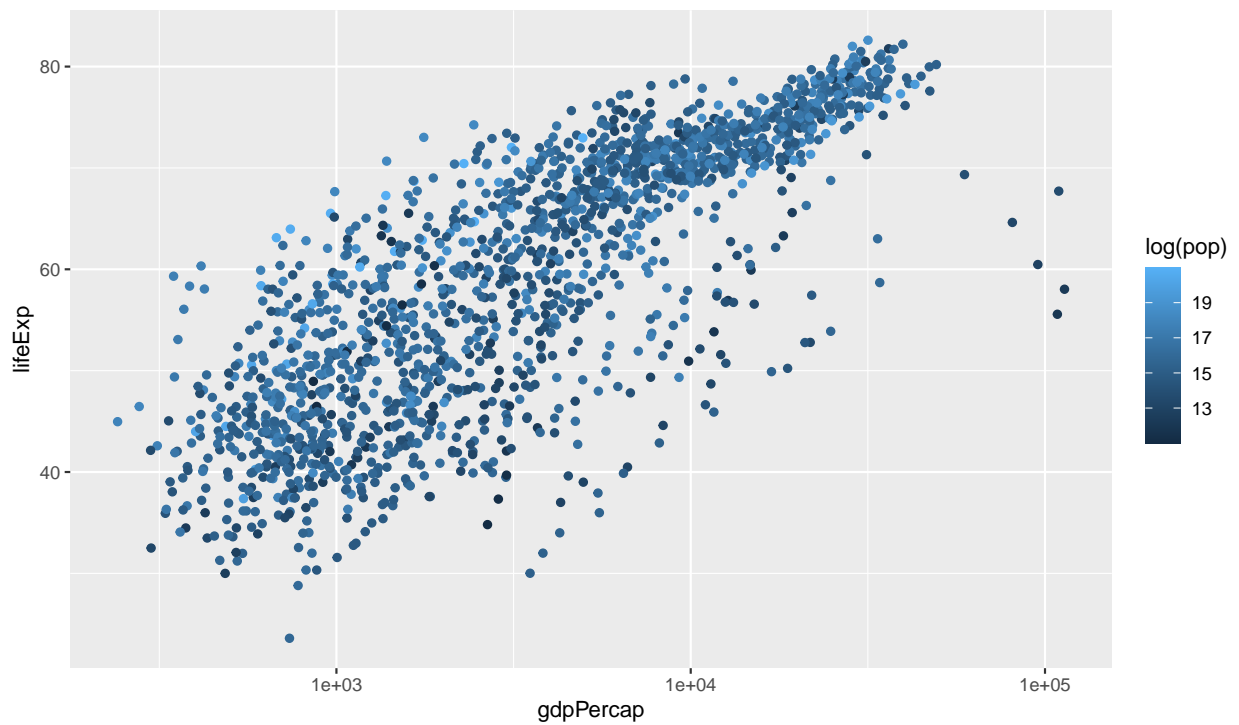


Figure 15: Mapping a continuous variable to color.

```

here()

## [1] "C:/Users/Nith Kosal/Documents/Kosal Documents/Kosal Teaching/DataVisualization"

p_out <- p + geom_point(mapping = aes(color = log(pop))) +
  scale_x_log10()

ggsave(here("Figure", "lifexp_vs_gdp_gradient.pdf"), plot = p_out)

ggsave(here("Figure", "lifexp_vs_gdp_gradient.png"), plot = p_out)


```

Grouped data and the group aesthetic

```

p <- ggplot(data = gapminder,
            mapping = aes(x = year,
                          y = gdpPercap))
p + geom_line()

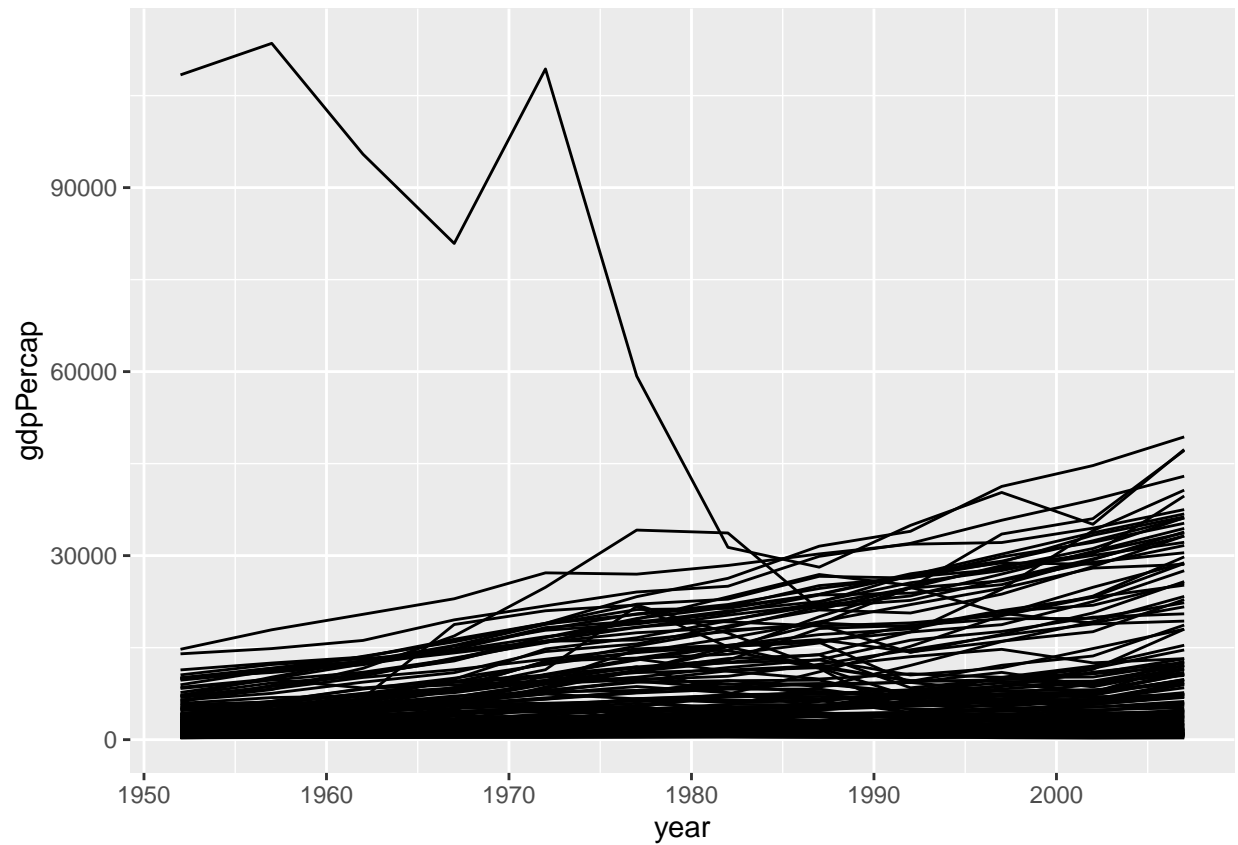
```



```

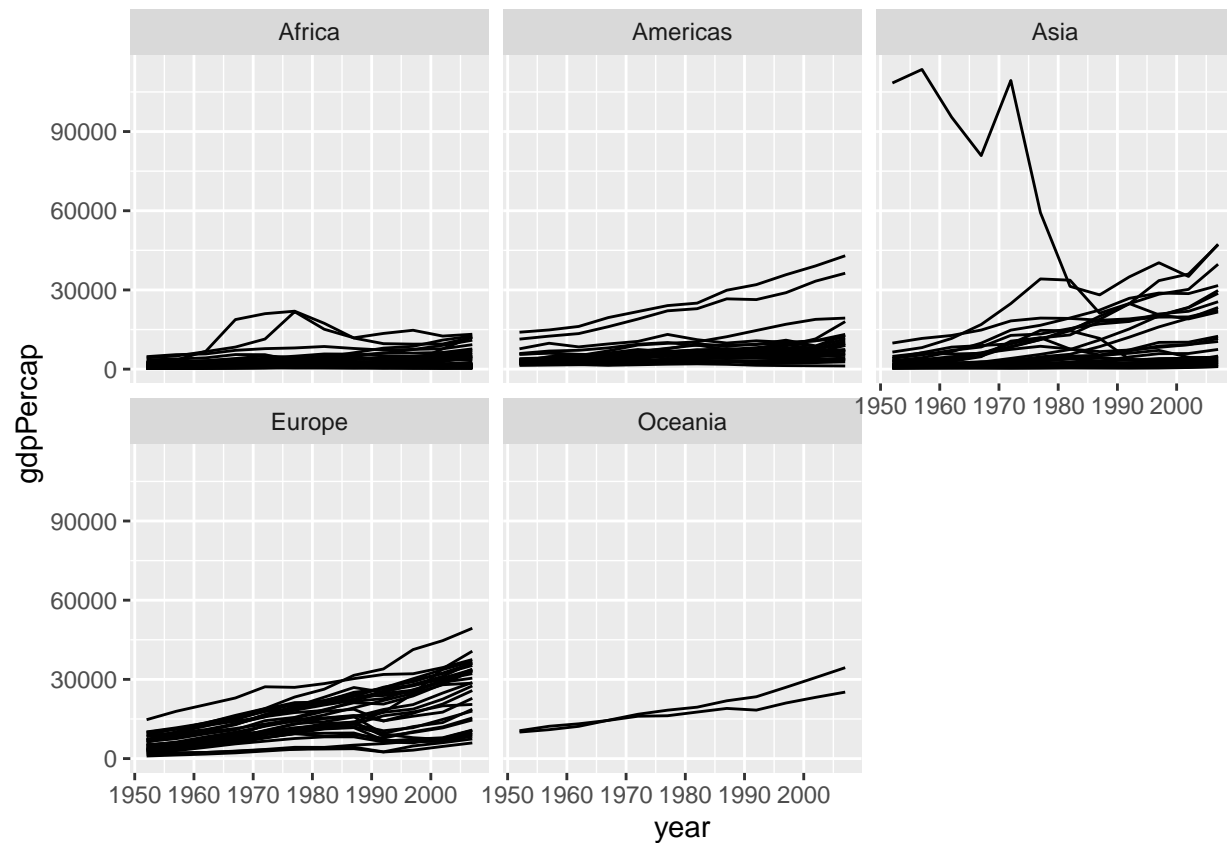
p <- ggplot(data = gapminder,
            mapping = aes(x = year,
                          y = gdpPercap))
p + geom_line(mapping =
              aes(group = country))

```



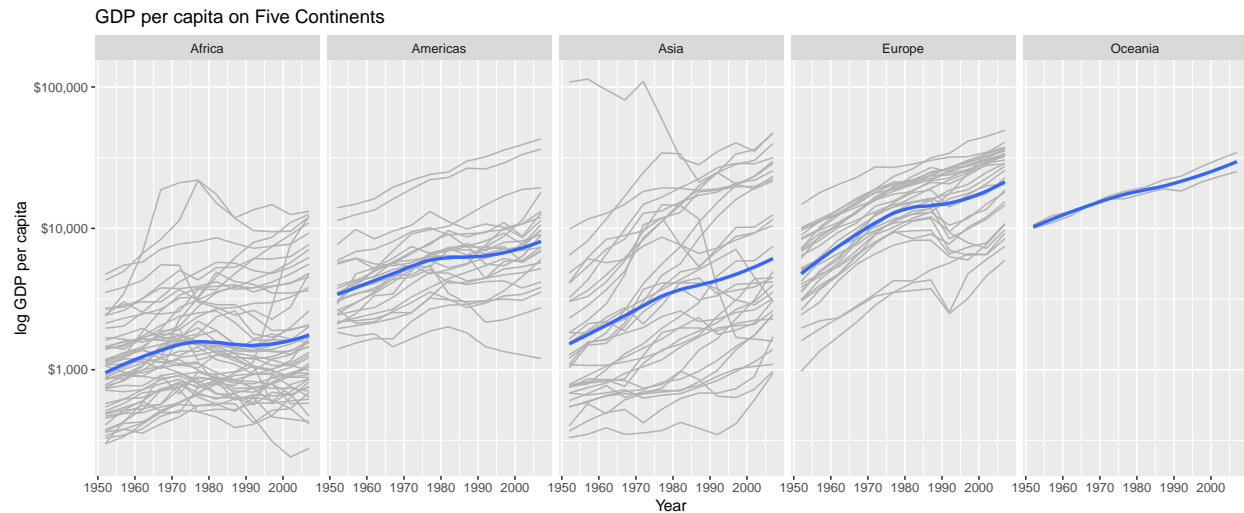
Faceting

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))  
  
p + geom_line(mapping =  
              aes(group = country)) +  
facet_wrap(~ continent)
```



```
p + geom_line(color="gray70",
              mapping=aes(group = country)) +
  geom_smooth(size = 1.1,
              method = "loess",
              se = FALSE) +
  scale_y_log10(labels=scales::dollar) +
  facet_wrap(~ continent, ncol = 5) +
  labs(x = "Year",
       y = "log GDP per capita",
       title = "GDP per capita on Five Continents")
```

```
## `geom_smooth()` using formula 'y ~ x'
```

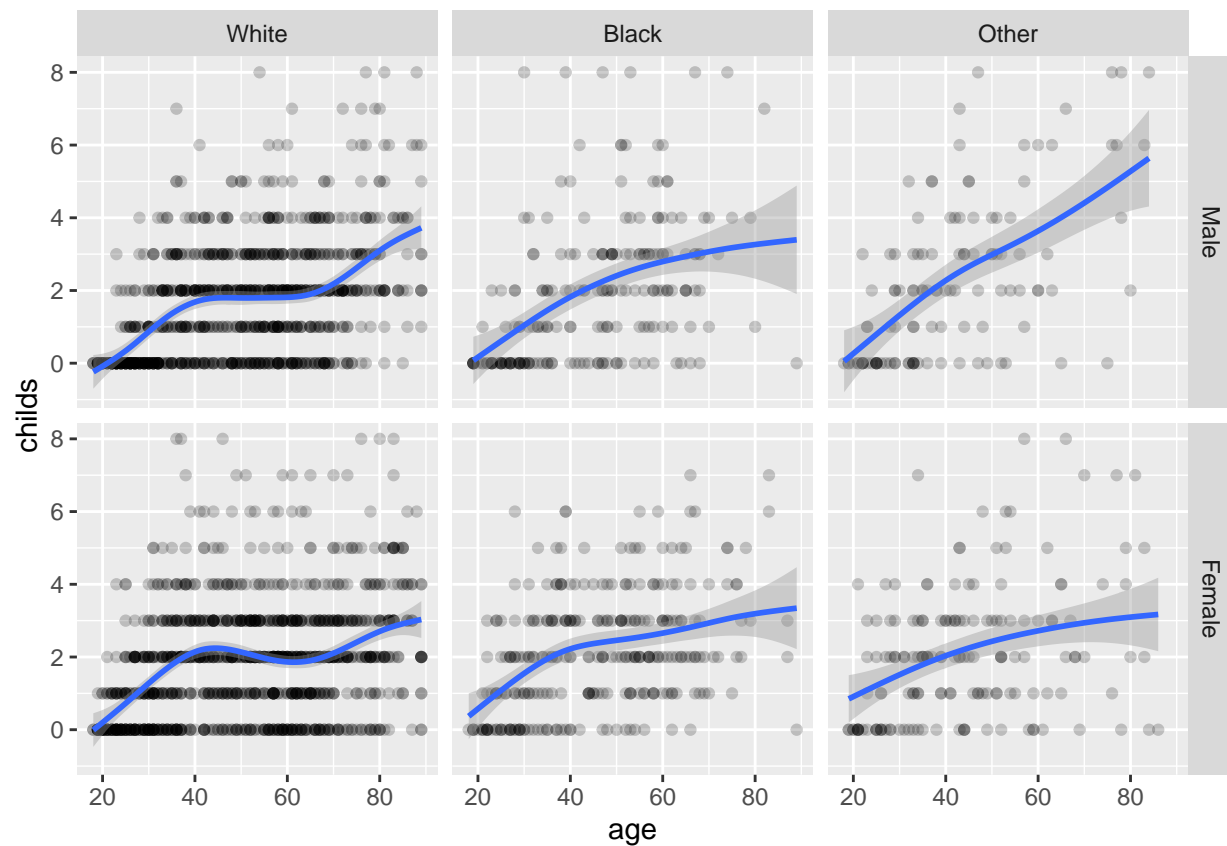


```
p <- ggplot(data = gss_sm,
            mapping = aes(x = age, y = childs))
p + geom_point(alpha = 0.2) + geom_smooth() +
  facet_grid(sex ~ race)
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

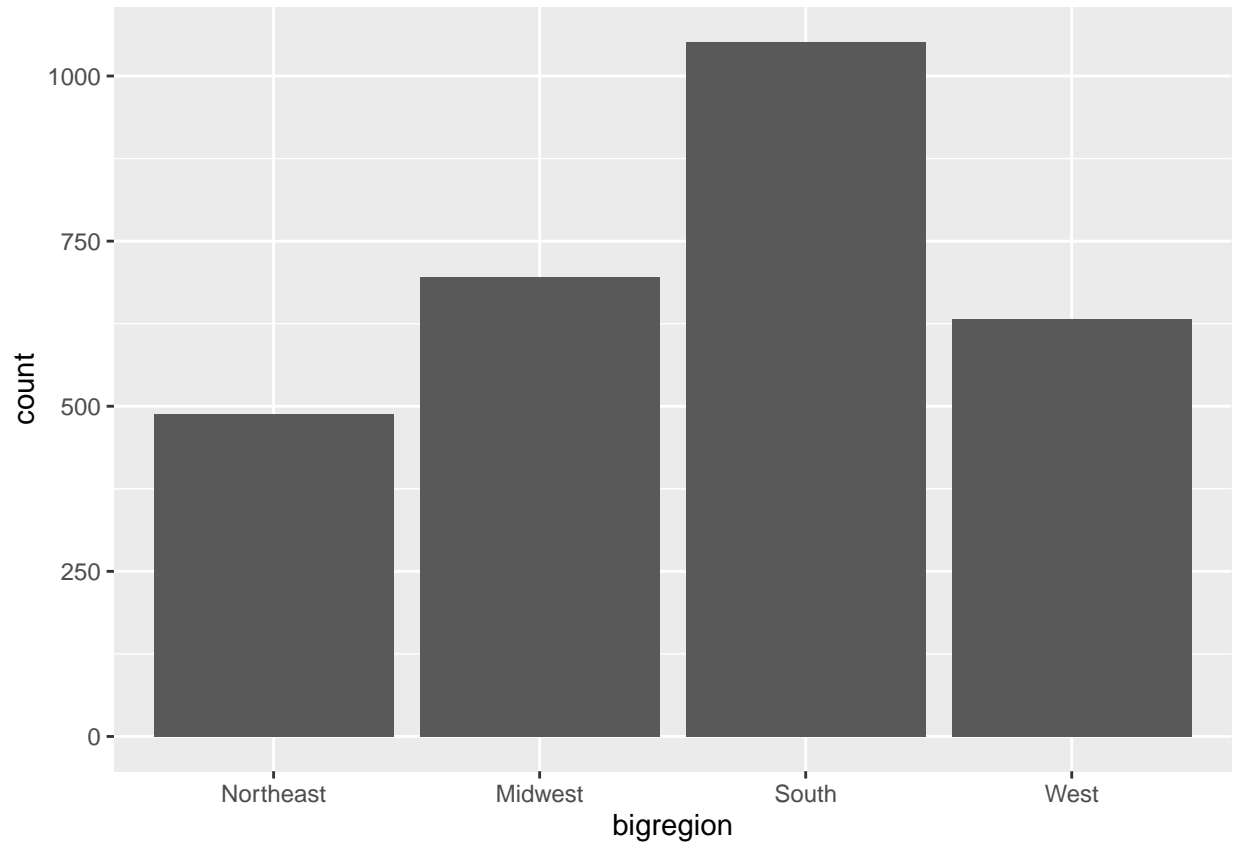
```
## Warning: Removed 18 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 18 rows containing missing values (geom_point).
```

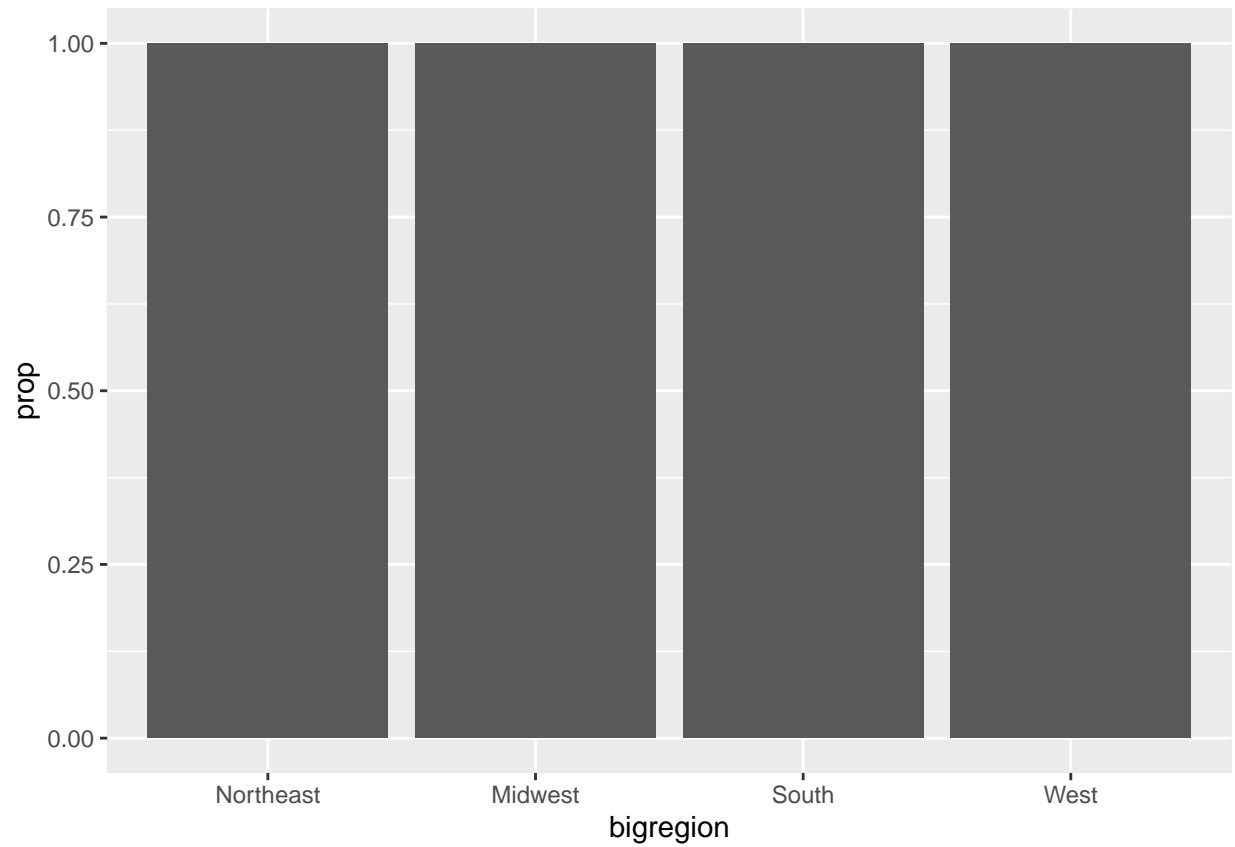


Geoms can transform data

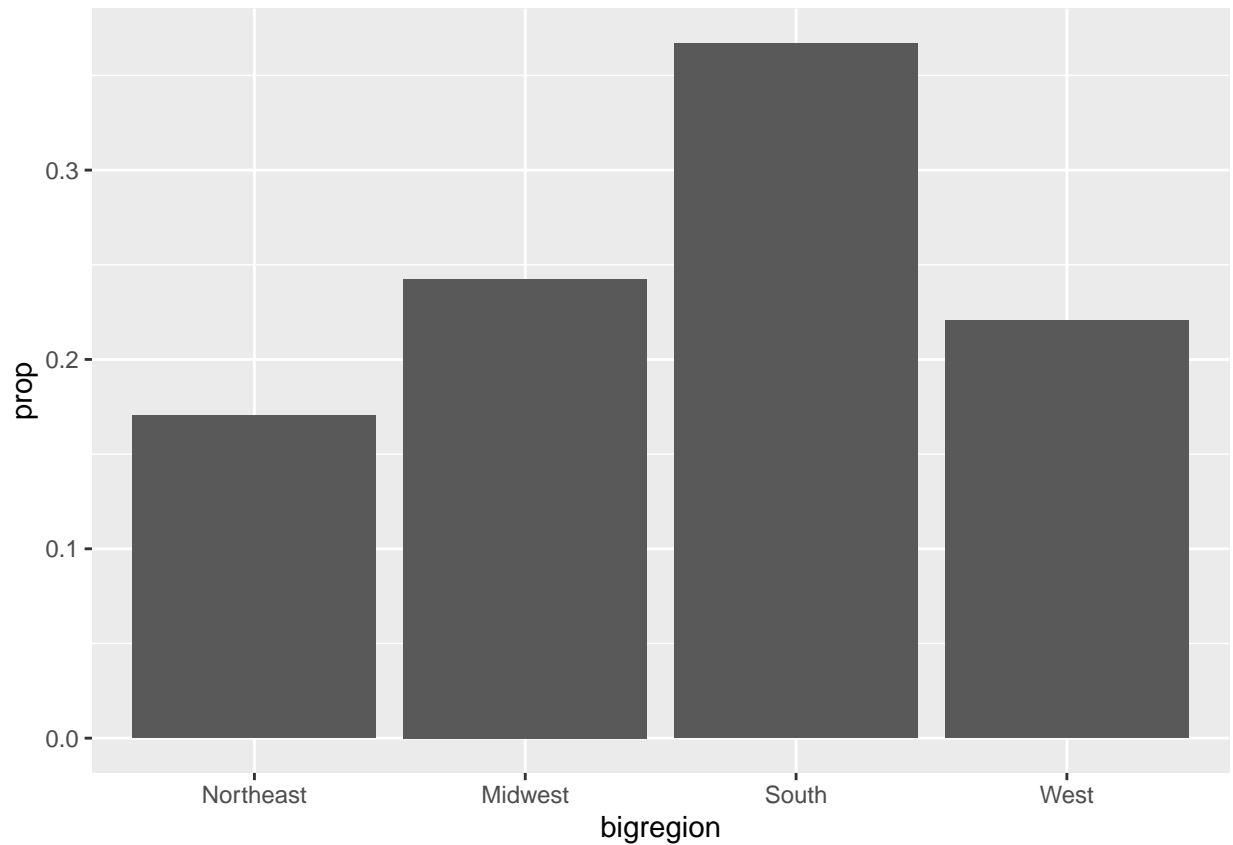
```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = bigregion))  
p + geom_bar()
```



```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = bigregion))  
p + geom_bar(mapping = aes(y = ..prop..))
```



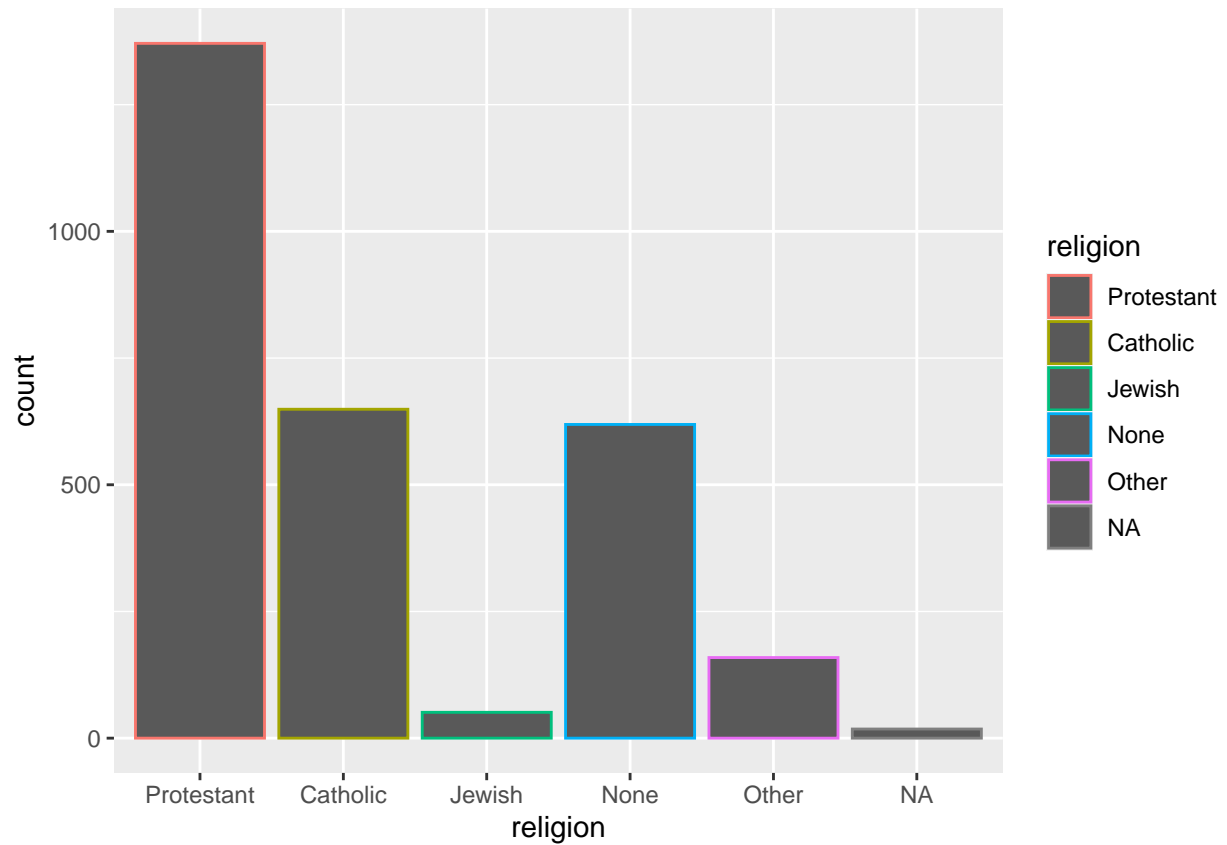
```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = bigregion))  
p + geom_bar(mapping = aes(y = ..prop.., group = 1))
```


Frequency Plots the Slightly Awkward Way

```
table(gss_sm$religion)
```

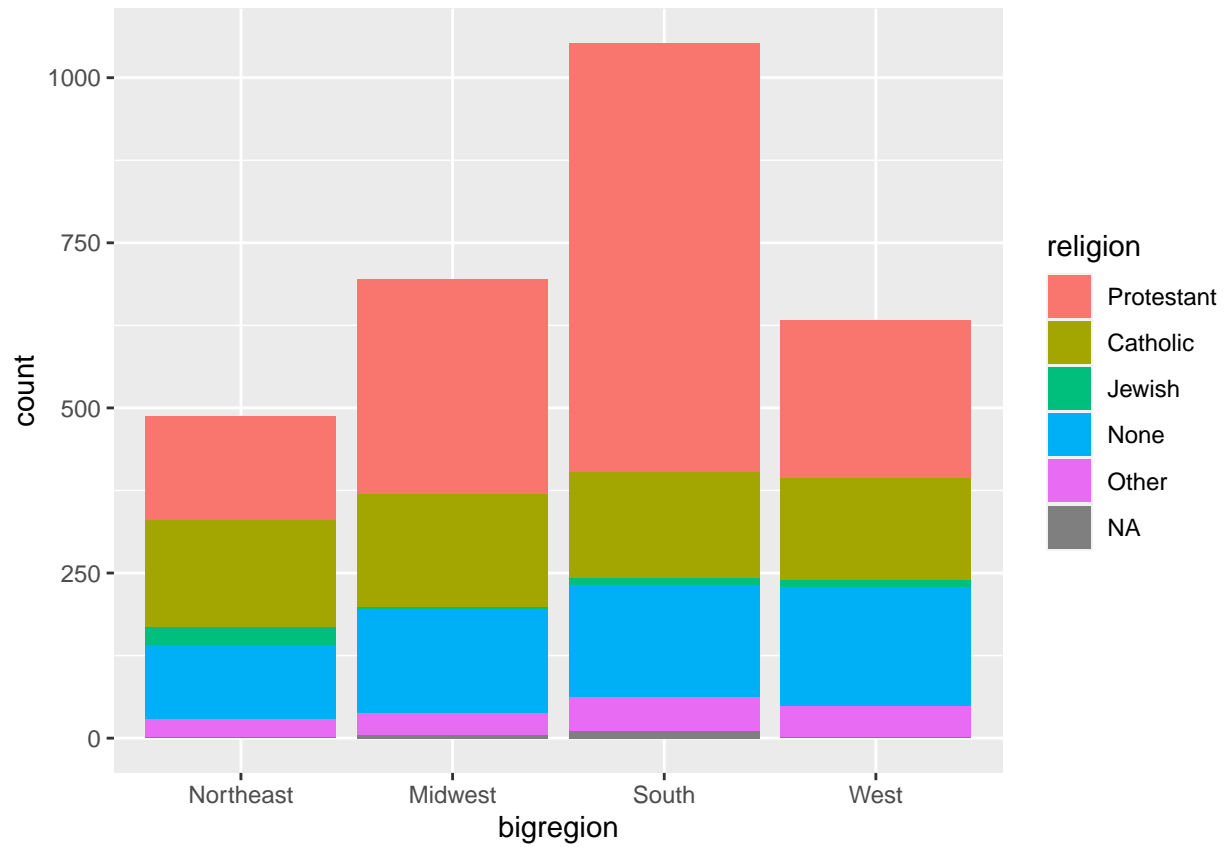
```
##  
## Protestant    Catholic    Jewish      None      Other  
##      1371         649         51      619      159  
  
p <- ggplot(data = gss_sm,  
            mapping = aes(x = religion, color = religion))  
p + geom_bar()
```



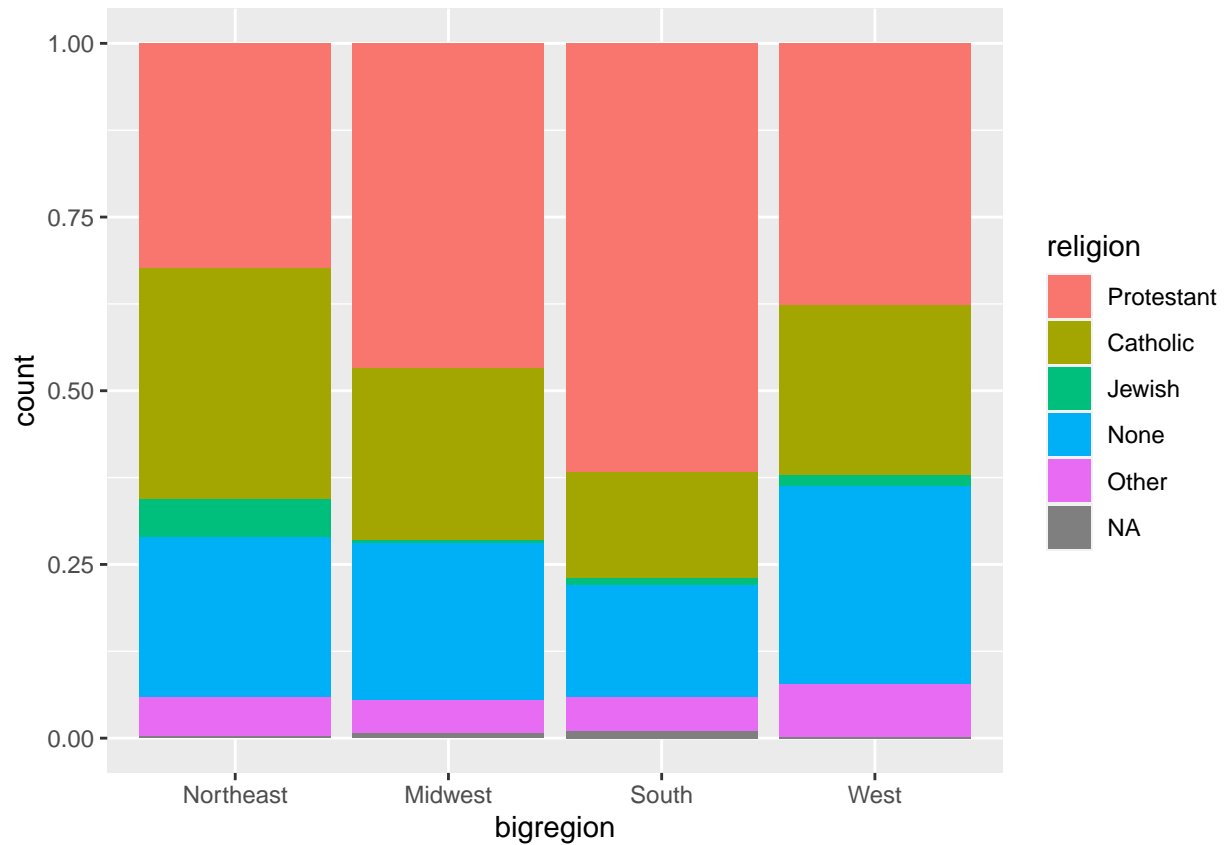
```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = religion, fill = religion))  
p + geom_bar() + guides(fill = FALSE)
```



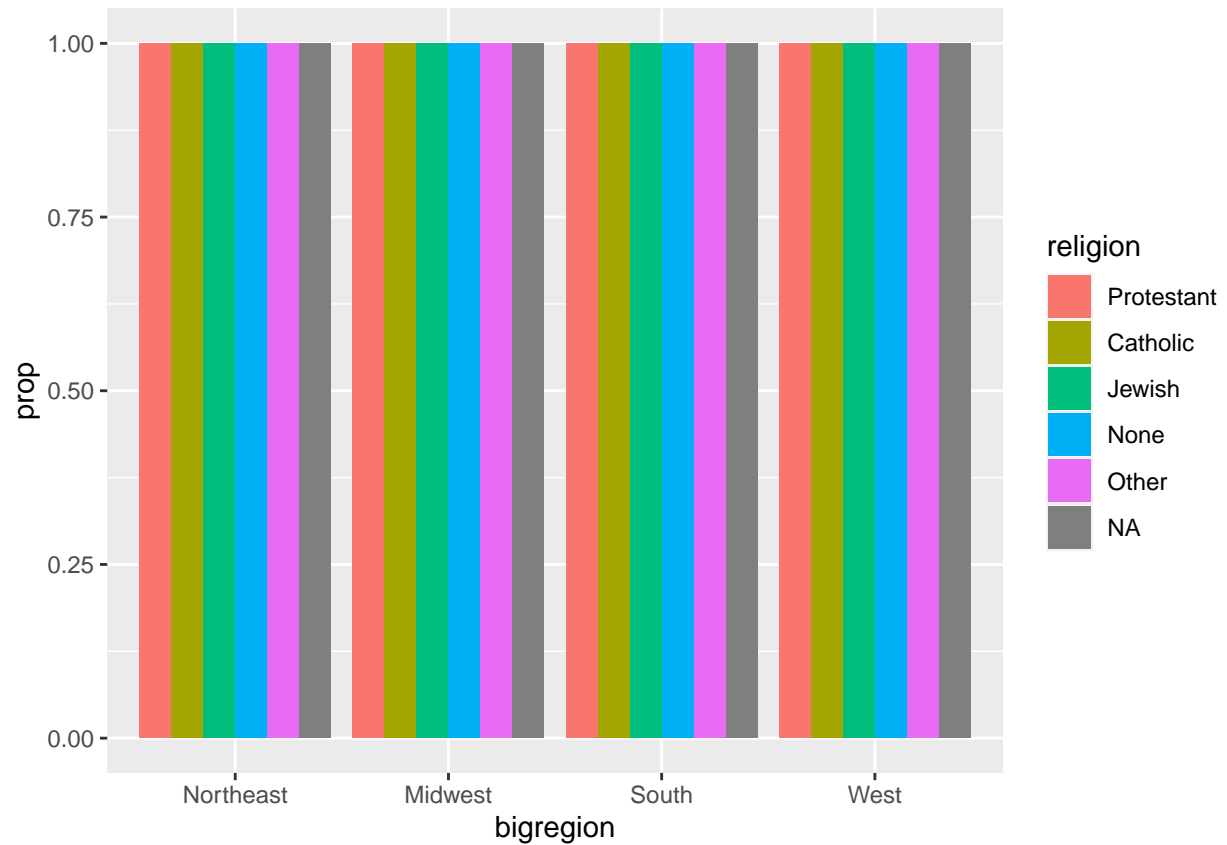
```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = bigregion,  
                          fill = religion))  
p + geom_bar()
```



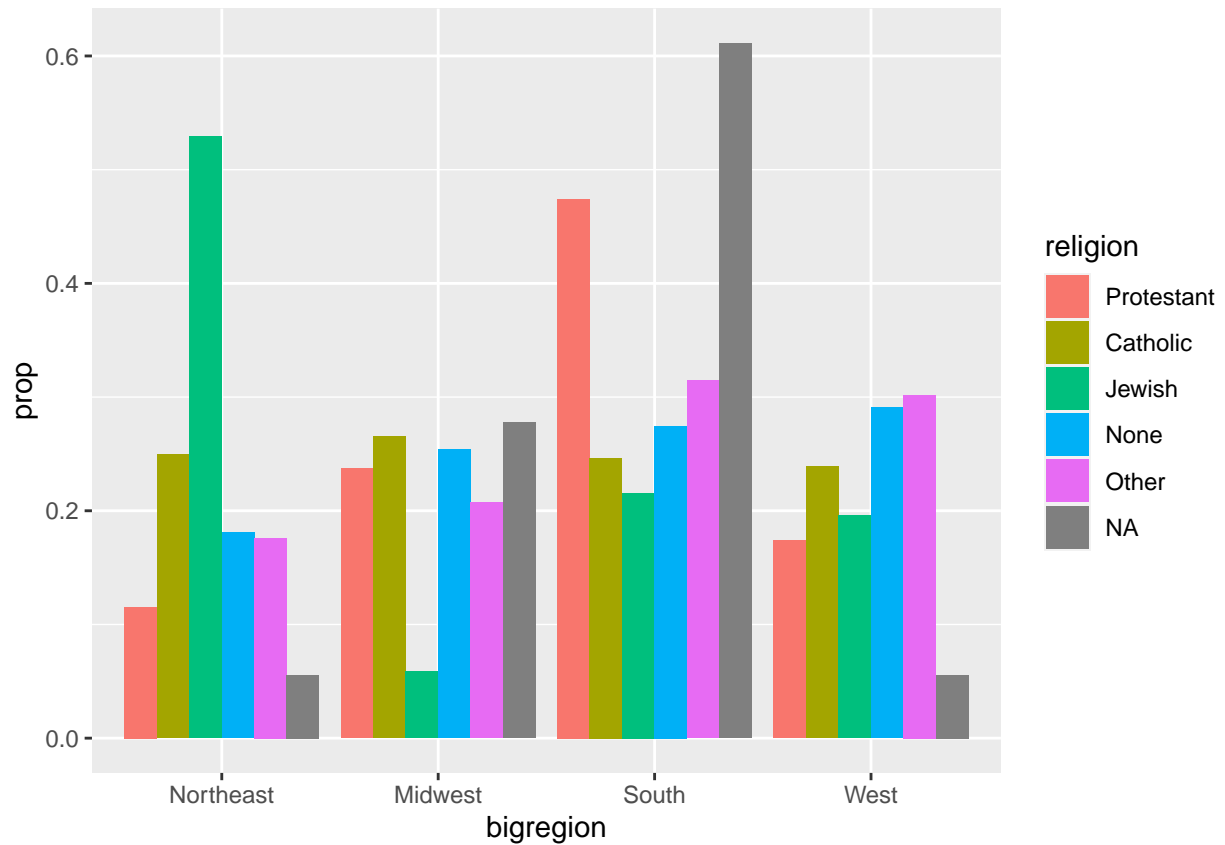
```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = bigregion,  
                          fill = religion))  
p + geom_bar(position = "fill")
```



```
p <- ggplot(data = gss_sm,
            mapping = aes(x = bigregion,
                          fill = religion))
p + geom_bar(position = "dodge",
            mapping = aes(y = ..prop..))
```



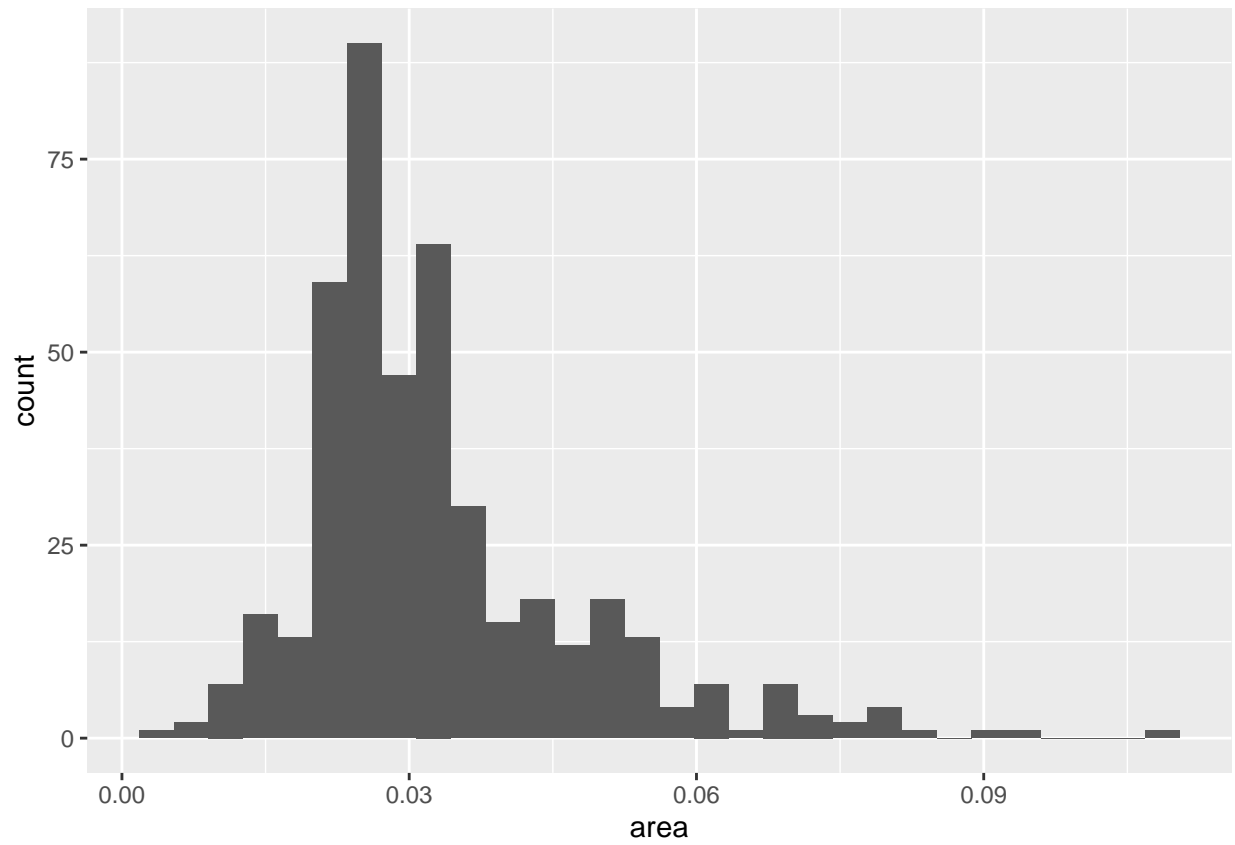
```
p <- ggplot(data = gss_sm,  
            mapping = aes(x = bigregion,  
                          fill = religion))  
p + geom_bar(position = "dodge",  
            mapping = aes(y = ..prop..,  
                          group = religion))
```



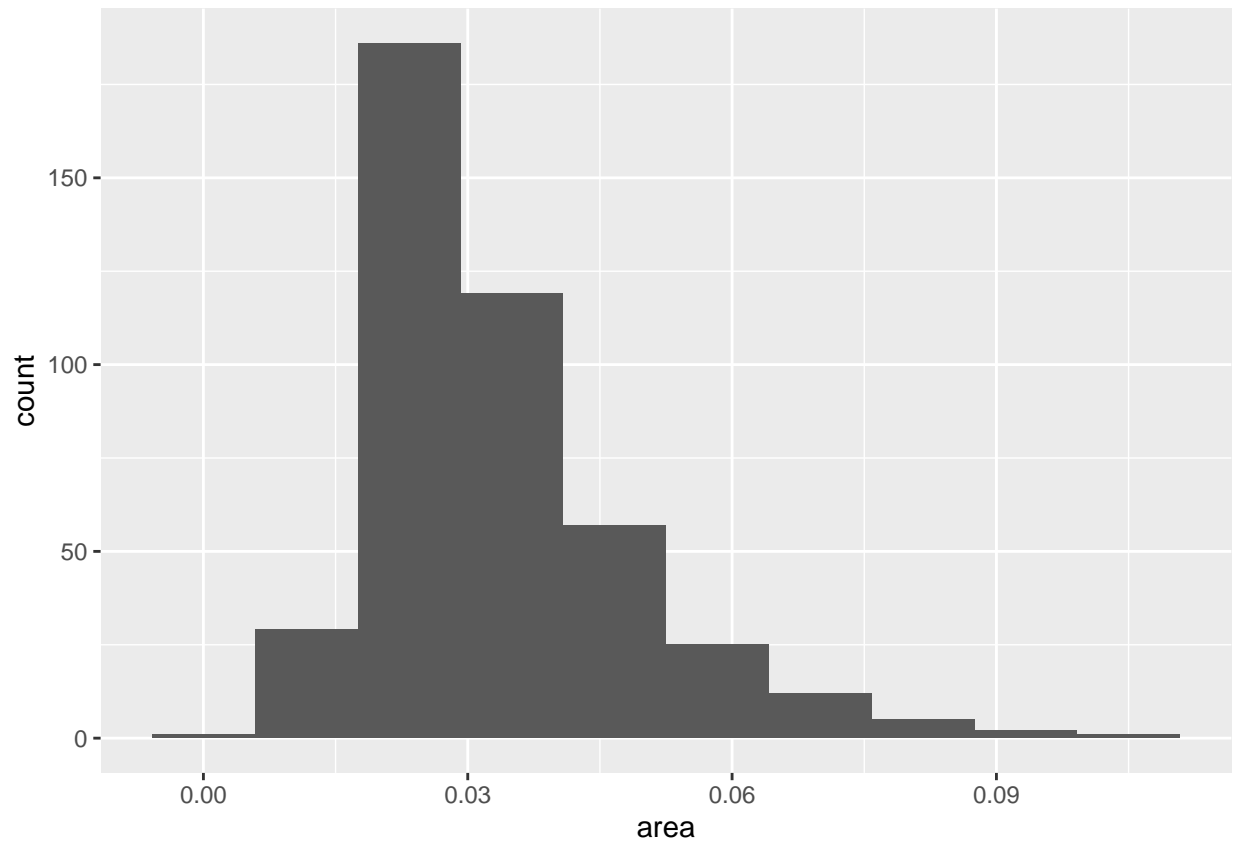
Histograms and Density Plots

```
p <- ggplot(data = midwest,
            mapping = aes(x = area))
p + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

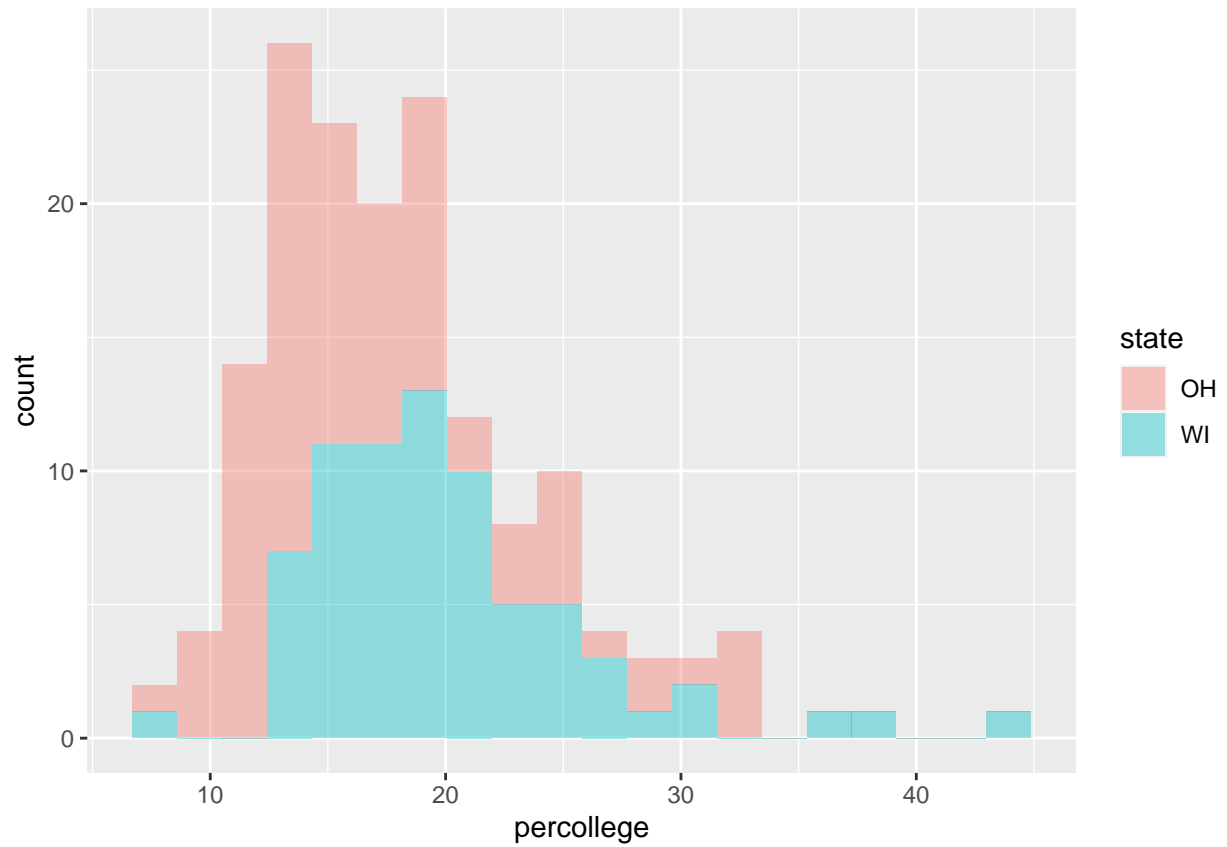


```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
p + geom_histogram(bins = 10)
```

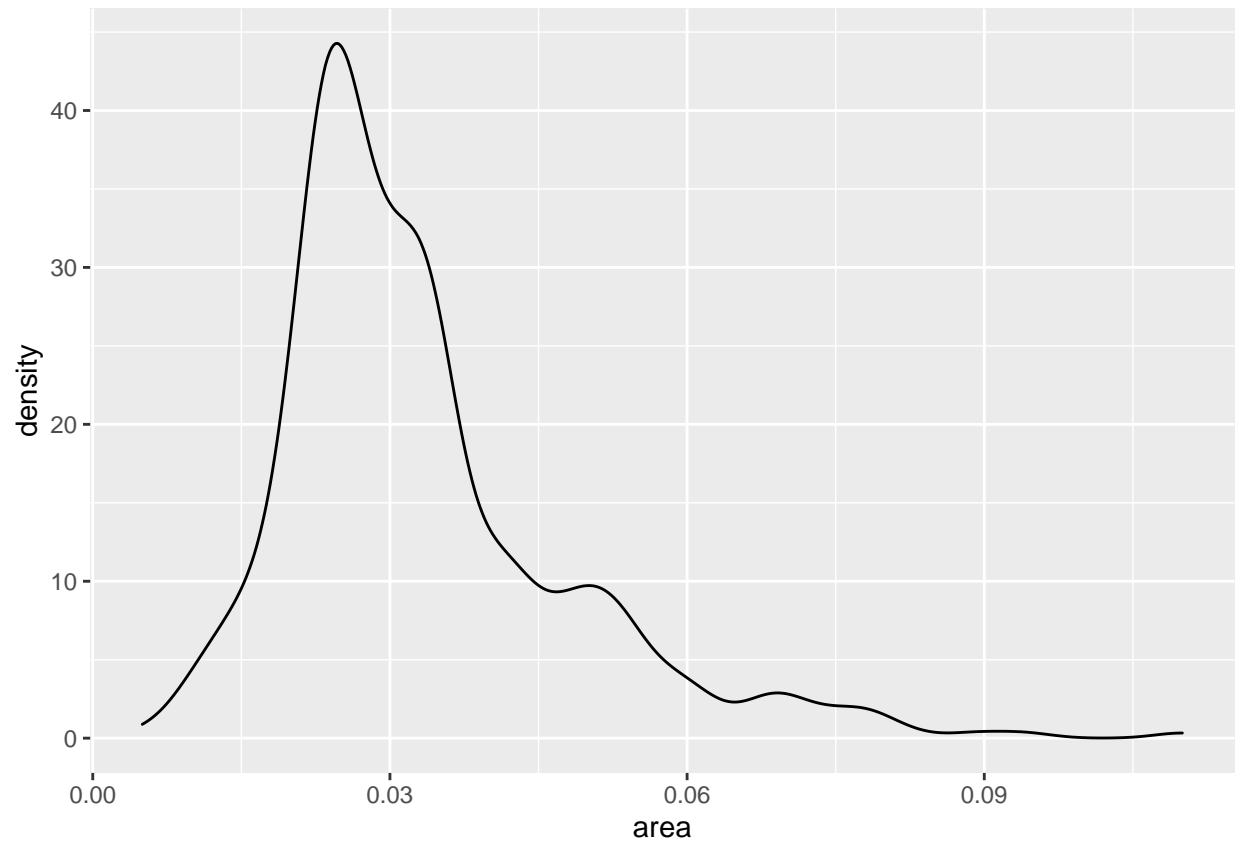



```
oh_wi <- c("OH", "WI")

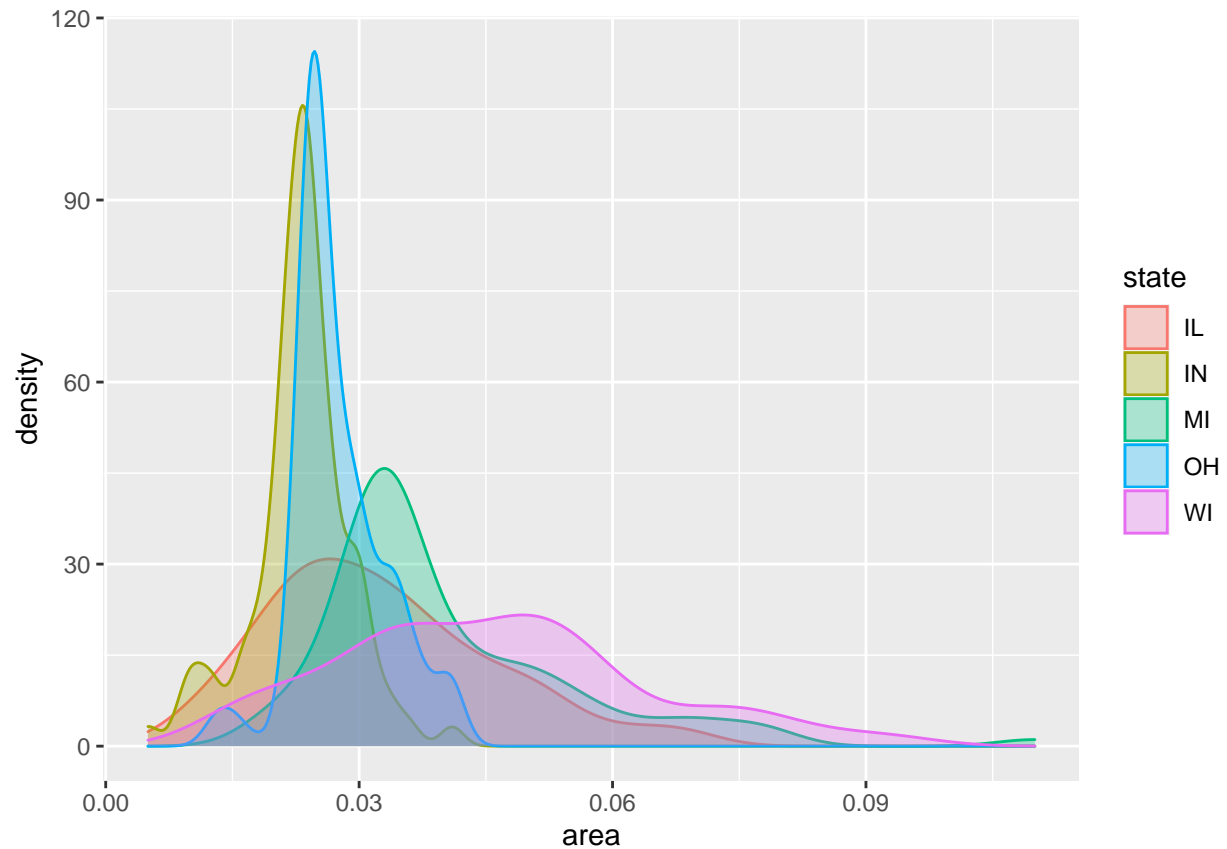
p <- ggplot(data = subset(midwest, subset = state %in% oh_wi),
            mapping = aes(x = percollege, fill = state))
p + geom_histogram(alpha = 0.4, bins = 20)
```



```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
p + geom_density()
```



```
p <- ggplot(data = midwest,  
            mapping = aes(x = area, fill = state, color = state))  
p + geom_density(alpha = 0.3)
```




Avoid transformations when necessary

titanic

```
##      fate    sex    n percent
## 1 perished  male 1364    62.0
## 2 perished female  126     5.7
## 3 survived  male  367    16.7
## 4 survived female  344    15.6
```

```
p <- ggplot(data = titanic,
            mapping = aes(x = fate, y = percent, fill = sex))
p + geom_bar(position = "dodge", stat = "identity") + theme(legend.position = "top")
```

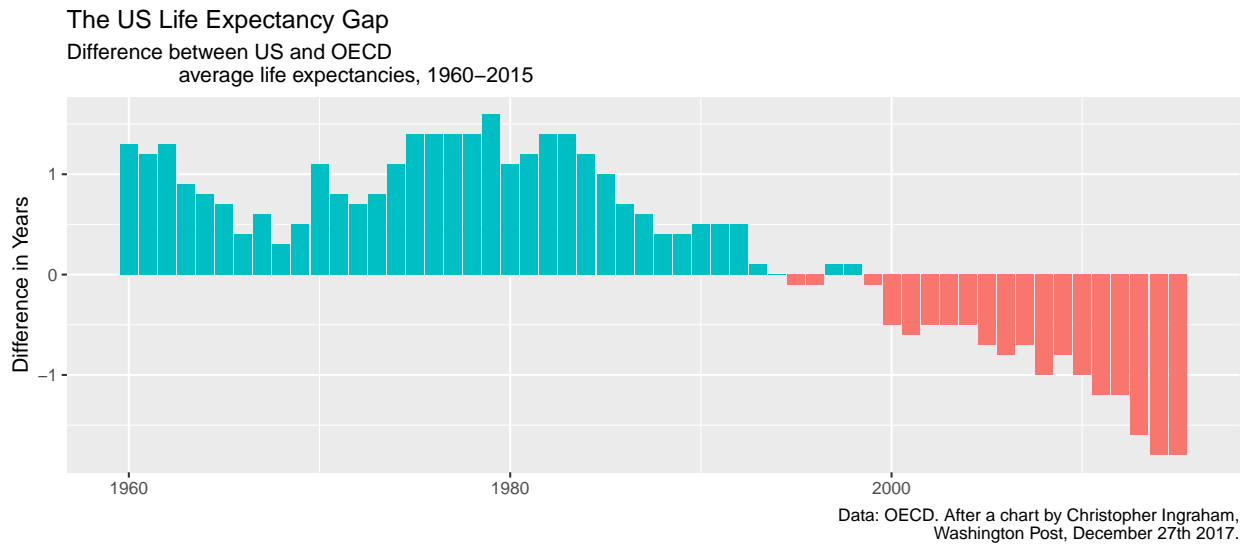


oecd_sum

```
## # A tibble: 57 x 5
## # Groups:   year [57]
##   year other  usa diff hi_lo
##   <int> <dbl> <dbl> <dbl> <chr>
## 1 1960 68.6 69.9 1.30 Below
## 2 1961 69.2 70.4 1.20 Below
## 3 1962 68.9 70.2 1.30 Below
## 4 1963 69.1 70   0.900 Below
## 5 1964 69.5 70.3 0.800 Below
## 6 1965 69.6 70.3 0.700 Below
## 7 1966 69.9 70.3 0.400 Below
## 8 1967 70.1 70.7 0.600 Below
## 9 1968 70.1 70.4 0.300 Below
## 10 1969 70.1 70.6 0.5   Below
## # ... with 47 more rows
```

```
p <- ggplot(data = oecd_sum,
            mapping = aes(x = year, y = diff, fill = hi_lo))
p + geom_col() + guides(fill = FALSE) +
  labs(x = NULL, y = "Difference in Years",
       title = "The US Life Expectancy Gap",
       subtitle = "Difference between US and OECD
                  average life expectancies, 1960-2015",
       caption = "Data: OECD. After a chart by Christopher Ingraham,
                  Washington Post, December 27th 2017.")
```

```
## Warning: Removed 1 rows containing missing values (position_stack).
```



Dplyr pipelines

```
rel_by_region <- gss_sm %>%
  group_by(bigregion, religion) %>%
  summarize(N = n()) %>%
  mutate(freq = N / sum(N),
         pct = round((freq*100), 0))
```

`summarise()` has grouped output by 'bigregion'. You can override using the `.groups` argument.

```
rel_by_region
```

```
## # A tibble: 24 x 5
## # Groups:   bigregion [4]
##   bigregion religion      N    freq    pct
##   <fct>      <fct>    <int>  <dbl>  <dbl>
## 1 Northeast Protestant   158  0.324    32
## 2 Northeast Catholic    162  0.332    33
## 3 Northeast Jewish       27  0.0553     6
## 4 Northeast None        112  0.230    23
## 5 Northeast Other        28  0.0574     6
## 6 Northeast <NA>         1  0.00205    0
## 7 Midwest Protestant   325  0.468    47
## 8 Midwest Catholic    172  0.247    25
## 9 Midwest Jewish        3  0.00432     0
## 10 Midwest None        157  0.226    23
## # ... with 14 more rows
```

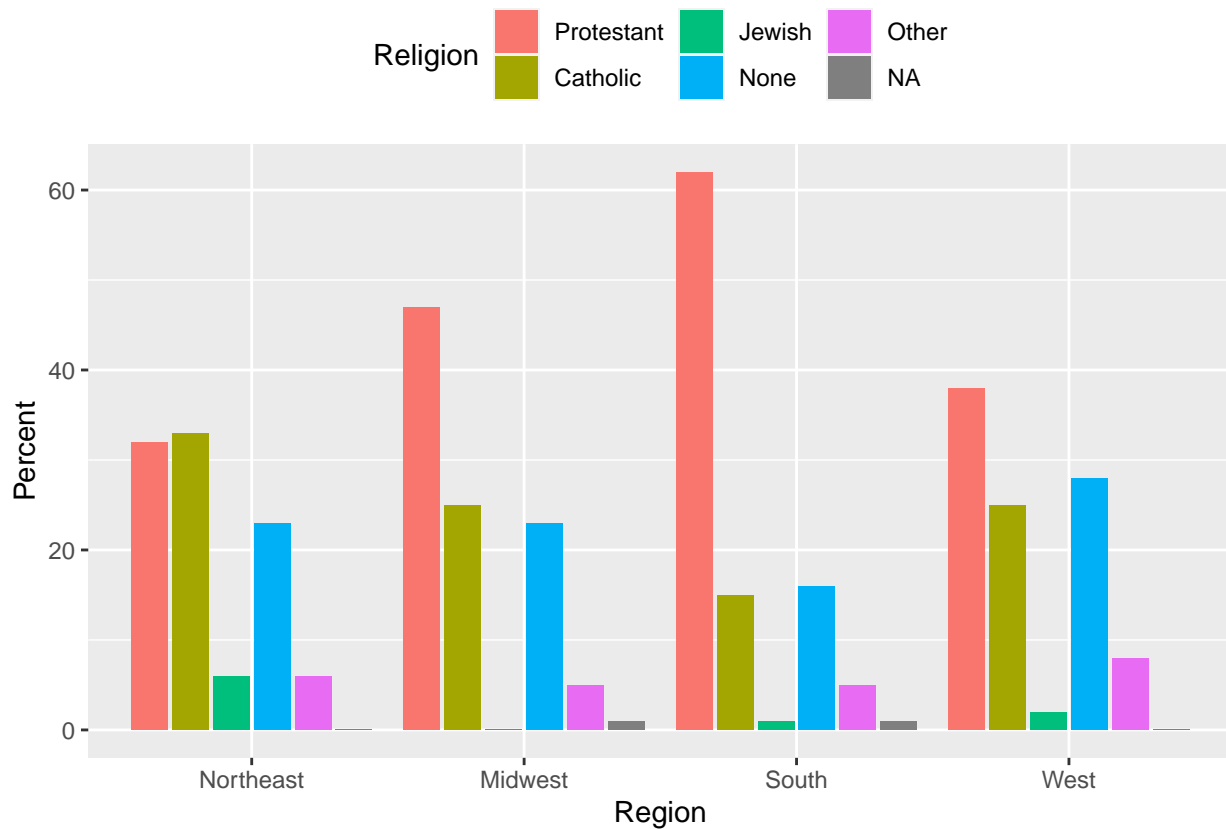
```
rel_by_region %>% group_by(bigregion) %>%
  summarize(total = sum(pct))
```

```
## # A tibble: 4 x 2
##   bigregion total
##   <fct>      <dbl>
## 1 Northeast    100
```


```
## 2 Midwest      101
## 3 South        100
## 4 West         101
```

```
## dodge2 presently requires the development version of ggplot
## devtools::install_github("tidyverse/ggplot2")
```

```
p <- ggplot(rel_by_region, aes(x = bigregion, y = pct, fill = religion))
p + geom_col(position = "dodge2") +
  labs(x = "Region", y = "Percent", fill = "Religion") +
  theme(legend.position = "top")
```



```
p <- ggplot(rel_by_region, aes(x = religion, y = pct, fill = religion))
p + geom_col(position = "dodge") +
  labs(x = NULL, y = "Percent", fill = "Religion") +
  guides(fill = FALSE) +
  coord_flip() +
  facet_grid(~ bigregion)
```



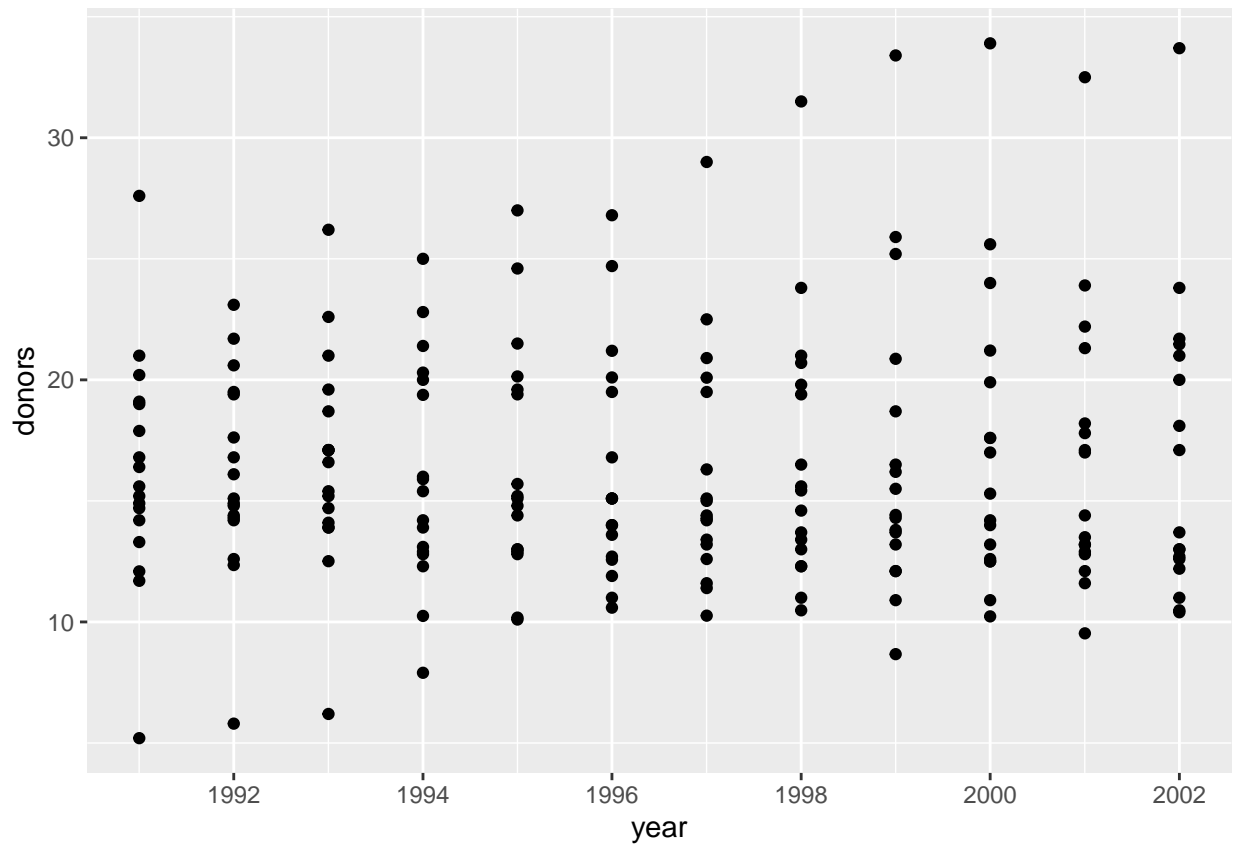
Continuous variables by category

```
organdata %>% select(1:6) %>% sample_n(size = 10)
```

```
## # A tibble: 10 x 6
##   country      year      donors    pop pop_dens    gdp
##   <chr>      <date>      <dbl>  <int>    <dbl>  <int>
## 1 Denmark  1996-01-01      14    5263     12.2  23548
## 2 Canada   1991-01-01     14.7  28031     0.281  19101
## 3 France   1992-01-01     16.8  57240     10.4   19566
## 4 United States 1996-01-01    20.1 269394     2.80  28772
## 5 Norway    NA           NA      NA      NA      NA
## 6 Belgium   1991-01-01     21   10005     30.2   18796
## 7 Belgium   1993-01-01     21   10085     30.5   19733
## 8 Denmark   1992-01-01     16.1   5171     12.0   19644
## 9 Ireland   1997-01-01     20.9   3673     5.23   22017
## 10 Denmark  1998-01-01     11    5304     12.3   25537
```

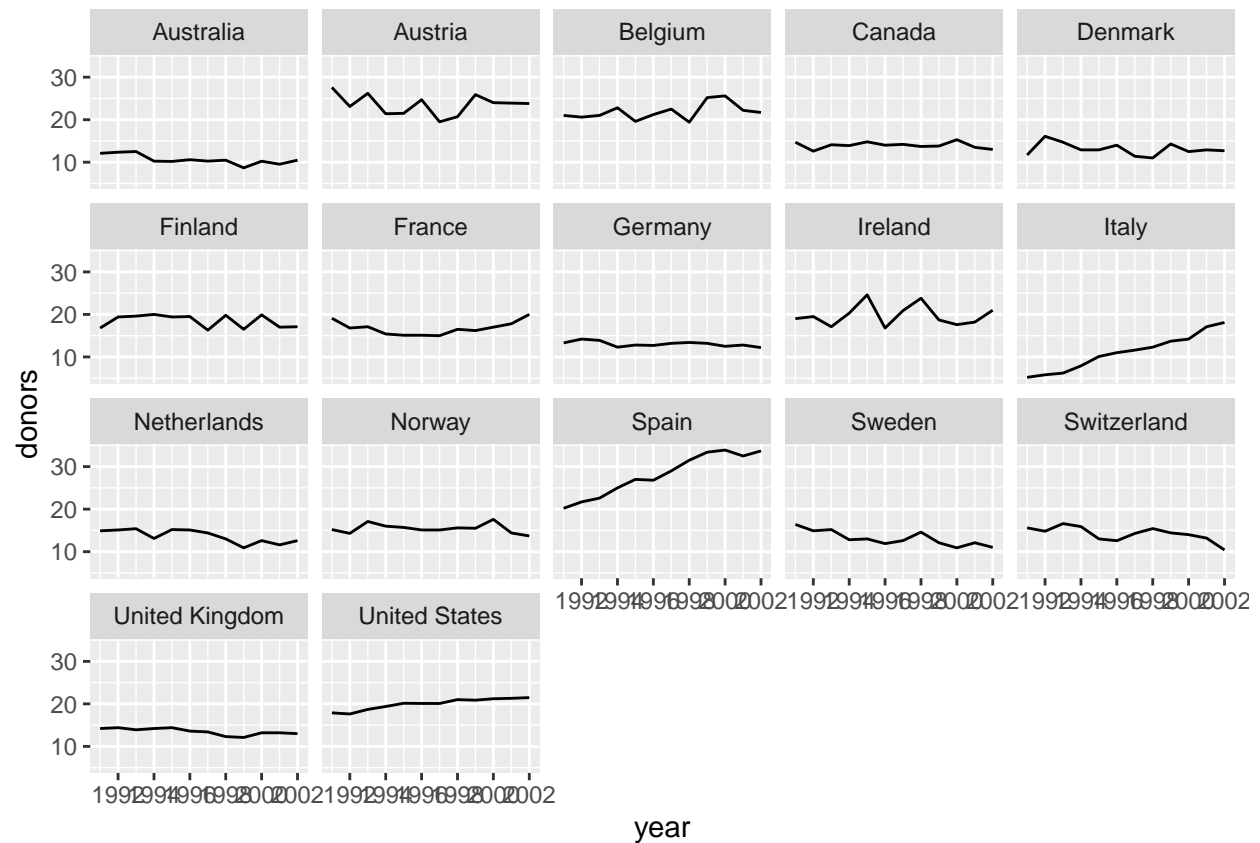
```
p <- ggplot(data = organdata,
            mapping = aes(x = year, y = donors))
p + geom_point()
```

```
## Warning: Removed 34 rows containing missing values (geom_point).
```

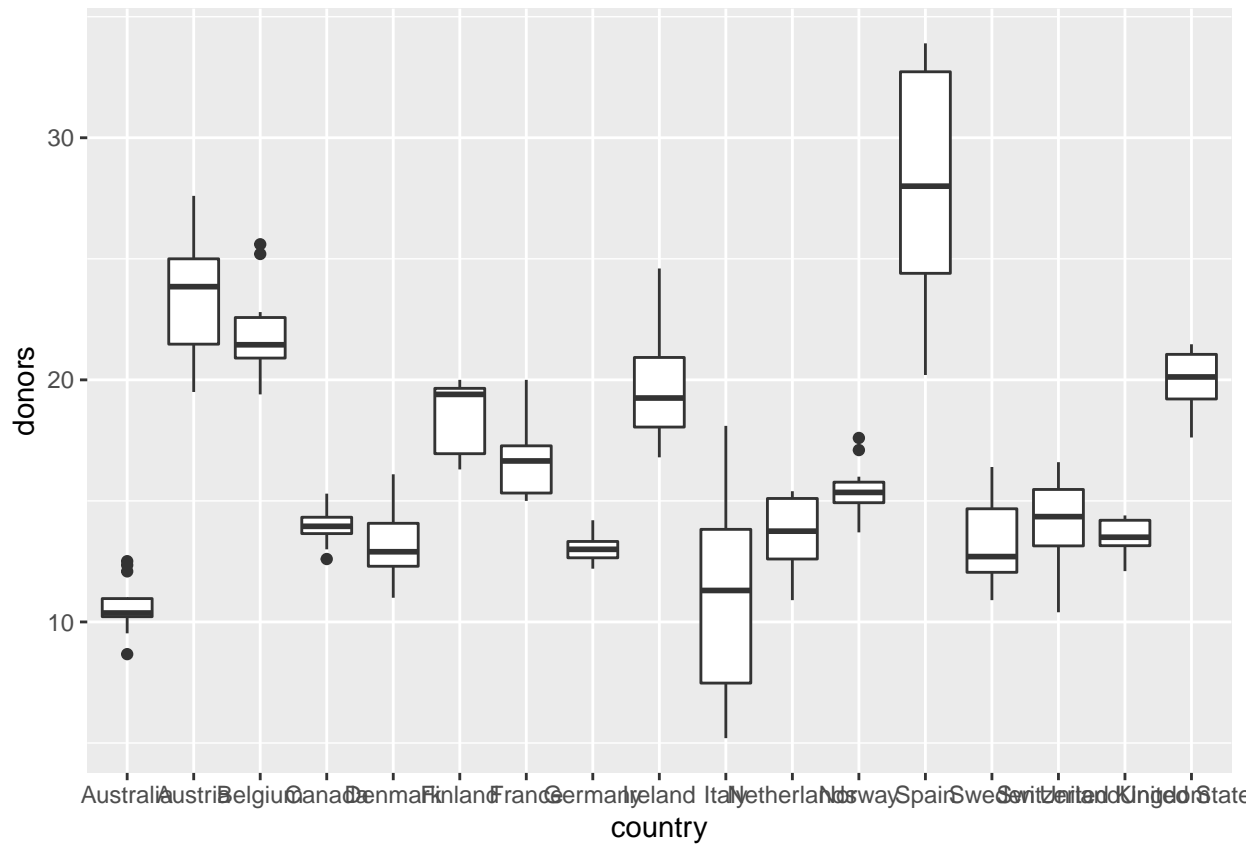
```
p <- ggplot(data = organdata,
            mapping = aes(x = year, y = donors))
p + geom_line(aes(group = country)) +
  facet_wrap(~ country)
```

```
## Warning: Removed 34 row(s) containing missing values (geom_path).
```



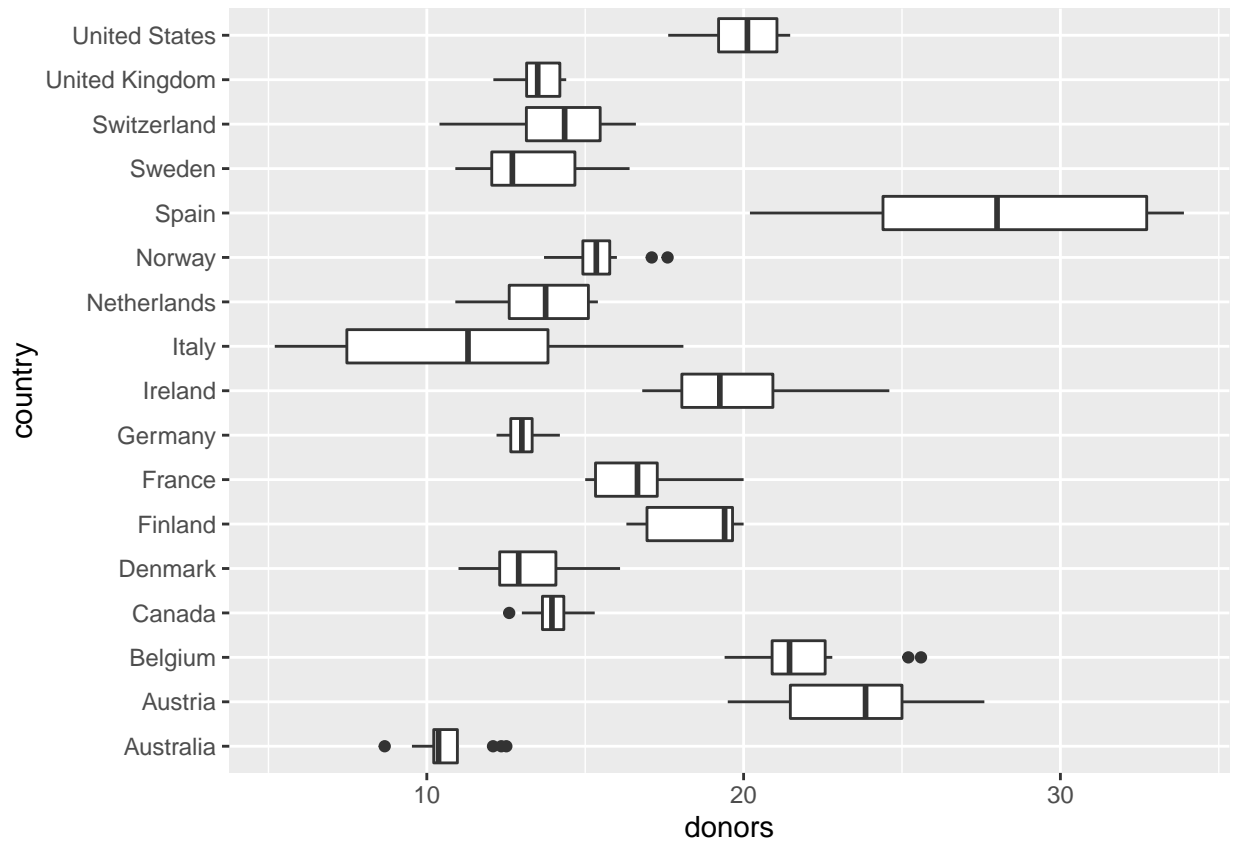
```
p <- ggplot(data = organdata,
            mapping = aes(x = country, y = donors))
p + geom_boxplot()
```

```
## Warning: Removed 34 rows containing non-finite values (stat_boxplot).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = country, y = donors))
p + geom_boxplot() + coord_flip()
```


```
## Warning: Removed 34 rows containing non-finite values (stat_boxplot).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = reorder(country, donors, na.rm=TRUE),
                          y = donors))

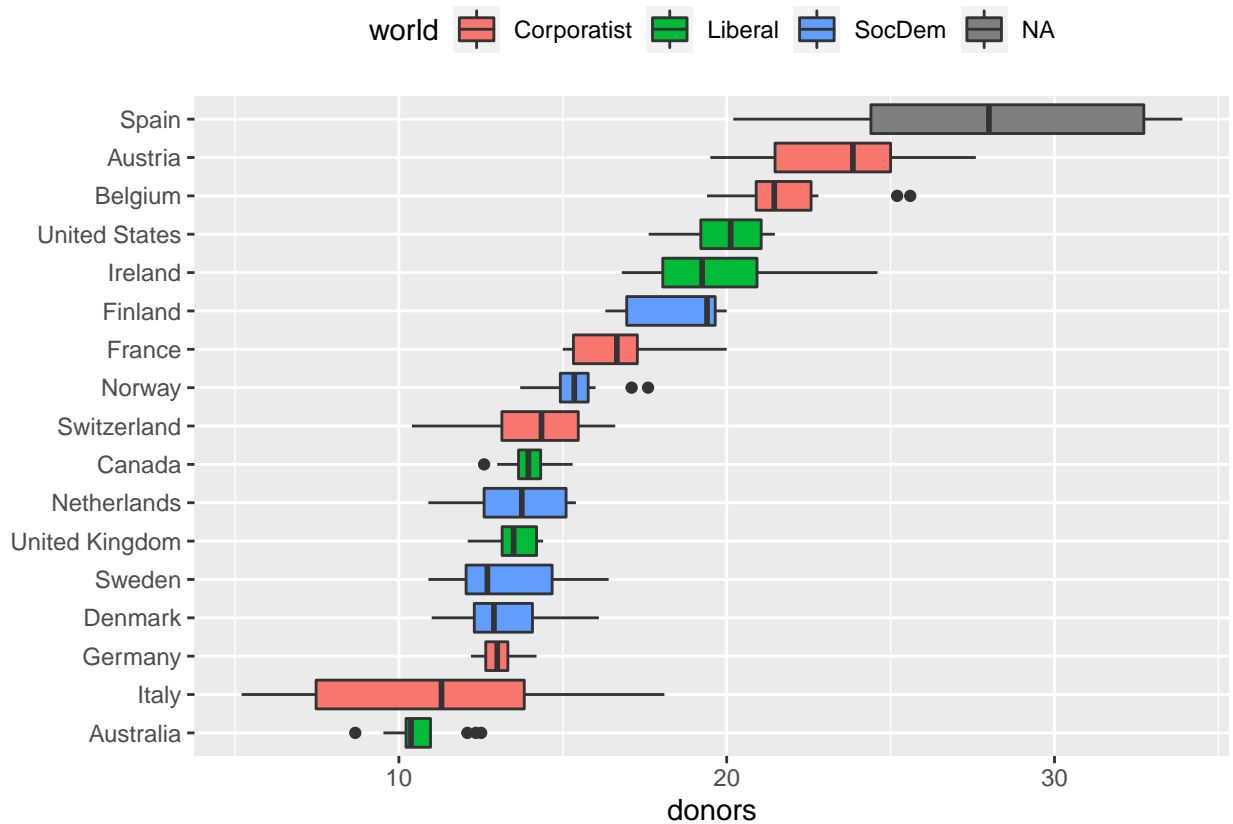
p + geom_boxplot() +
  labs(x=NULL) +
  coord_flip()
```

Warning: Removed 34 rows containing non-finite values (stat_boxplot).




```
p <- ggplot(data = organdata,
            mapping = aes(x = reorder(country, donors, na.rm=TRUE),
                          y = donors, fill = world))
p + geom_boxplot() + labs(x=NULL) +
  coord_flip() + theme(legend.position = "top")
```

```
## Warning: Removed 34 rows containing non-finite values (stat_boxplot).
```




```
p <- ggplot(data = organdata,
  mapping = aes(x = reorder(country, donors, na.rm=TRUE),
    y = donors, color = world))
p + geom_point() + labs(x=NULL) +
  coord_flip() + theme(legend.position = "top")
```

```
## Warning: Removed 34 rows containing missing values (geom_point).
```




```
p <- ggplot(data = organdata,
  mapping = aes(x = reorder(country, donors, na.rm=TRUE),
    y = donors, color = world))
p + geom_jitter() + labs(x=NULL) +
  coord_flip() + theme(legend.position = "top")
```

```
## Warning: Removed 34 rows containing missing values (geom_point).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = reorder(country, donors, na.rm=TRUE),
                          y = donors, color = world))
p + geom_jitter(position = position_jitter(width=0.15)) +
  labs(x=NULL) + coord_flip() + theme(legend.position = "top")
```

```
## Warning: Removed 34 rows containing missing values (geom_point).
```

```
by_country <- organdata %>% group_by(consent_law, country) %>%
  summarize(donors_mean= mean(donors, na.rm = TRUE),
            donors_sd = sd(donors, na.rm = TRUE),
            gdp_mean = mean(gdp, na.rm = TRUE),
            health_mean = mean(health, na.rm = TRUE),
            roads_mean = mean(roads, na.rm = TRUE),
            cerebvas_mean = mean(cerebvas, na.rm = TRUE))
```

`summarise()` has grouped output by 'consent_law'. You can override using the `.groups` argument.

```
by_country
```

```
## # A tibble: 17 x 8
## # Groups:   consent_law [2]
##   consent_law country      donors_mean donors_sd gdp_mean health_mean roads_mean
##   <chr>      <chr>          <dbl>     <dbl>   <dbl>     <dbl>     <dbl>
## 1 Informed   Australia         10.6      1.14   22179.    1958.     105.
## 2 Informed   Canada            14.0      0.751  23711.    2272.     109.
## 3 Informed   Denmark           13.1      1.47   23722.    2054.     102.
## 4 Informed   Germany            13.0      0.611  22163.    2349.     113.
## 5 Informed   Ireland            19.8      2.48   20824.    1480.     118.
## 6 Informed   Netherlands        13.7      1.55   23013.    1993.      76.1
## 7 Informed   United Kin~        13.5      0.775  21359.    1561.      67.9
## 8 Informed   United Sta~        20.0      1.33   29212.    3988.     155.
## 9 Presumed   Austria            23.5      2.42   23876.    1875.     150.
## 10 Presumed  Belgium            21.9      1.94   22500.    1958.     155.
## 11 Presumed  Finland            18.4      1.53   21019.    1615.      93.6
```

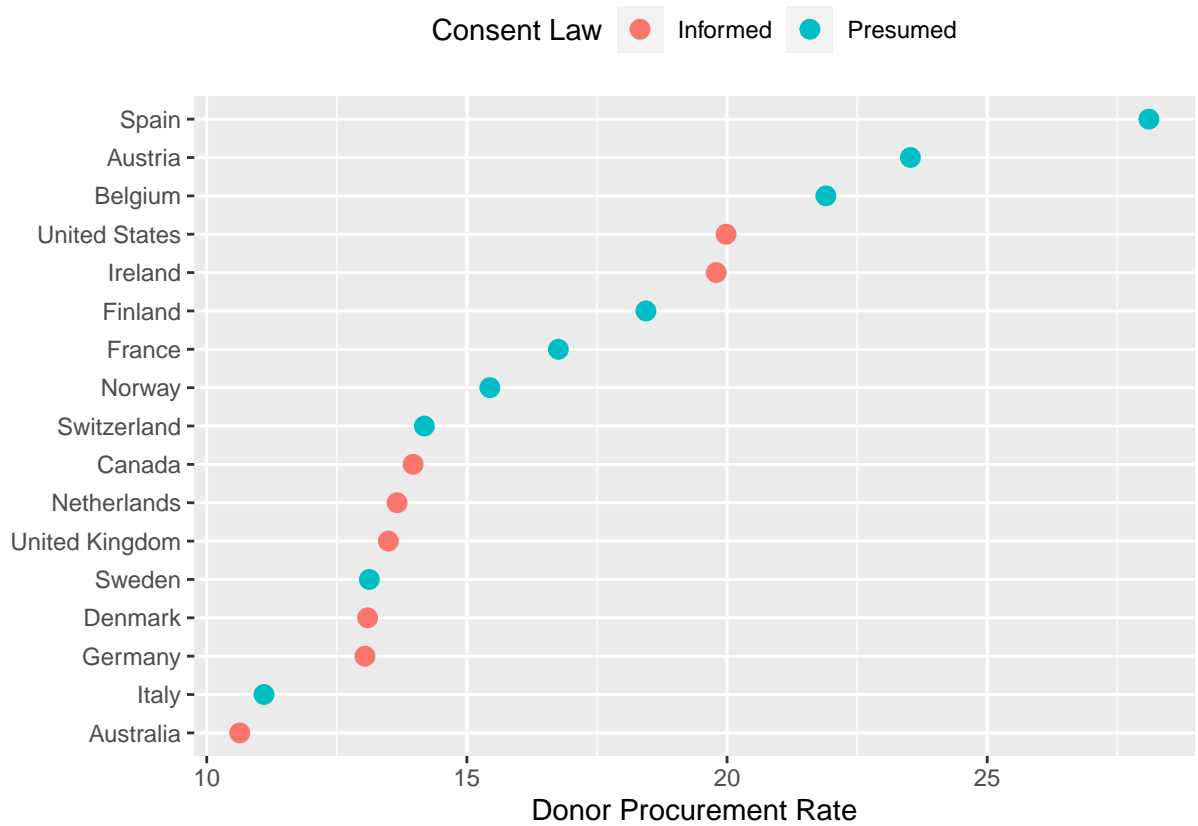
```
## 12 Presumed France 16.8 1.60 22603. 2160. 156.
## 13 Presumed Italy 11.1 4.28 21554. 1757 122.
## 14 Presumed Norway 15.4 1.11 26448. 2217. 70.0
## 15 Presumed Spain 28.1 4.96 16933 1289. 161.
## 16 Presumed Sweden 13.1 1.75 22415. 1951. 72.3
## 17 Presumed Switzerland 14.2 1.71 27233 2776. 96.4
## # ... with 1 more variable: cerebvas_mean <dbl>
```

```
by_country <- organdata %>%
  group_by(consent_law, country) %>%
    summarize_if(is.numeric,
      list(~ mean(., na.rm = TRUE),
           ~ sd(., na.rm = TRUE))) %>%
    ungroup()
```

```
by_country
```

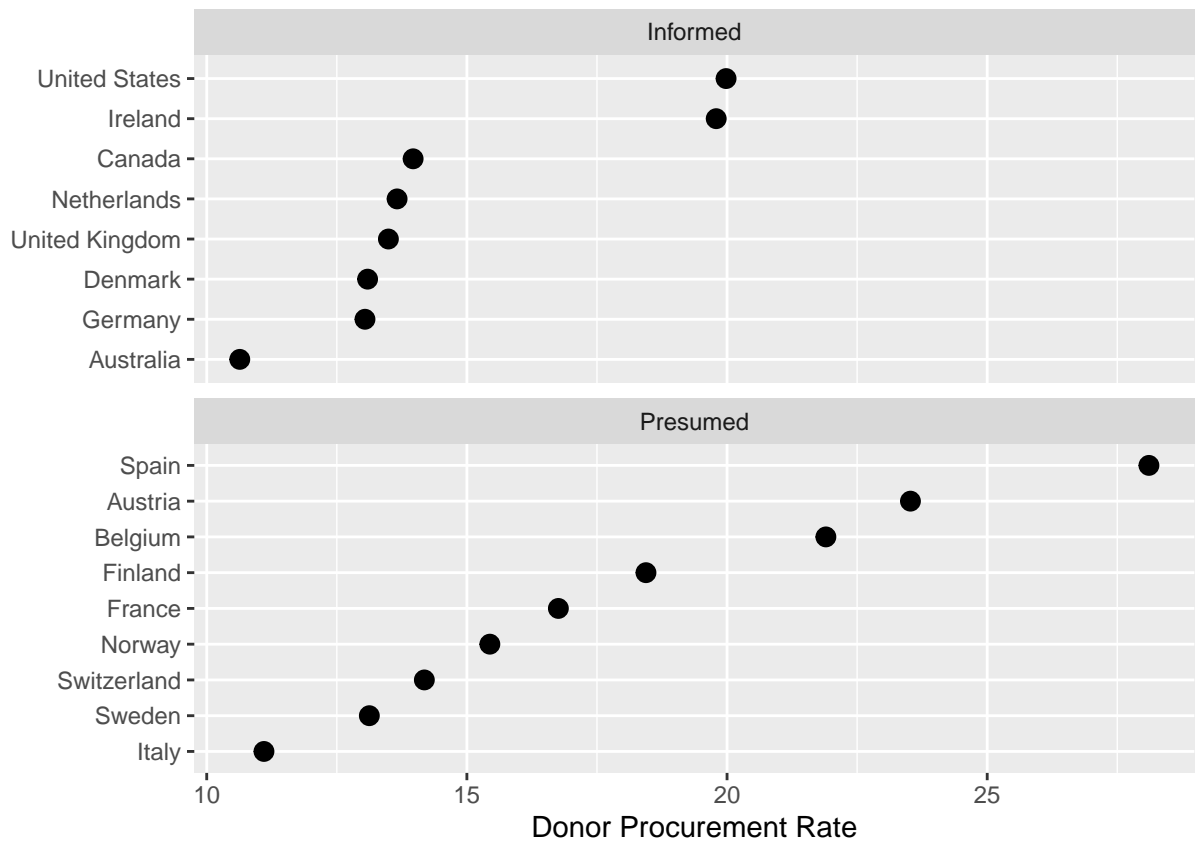
```
## # A tibble: 17 x 28
##   consent_law country donors_mean pop_mean pop_dens_mean gdp_mean gdp_lag_mean
##   <chr>      <chr>      <dbl>    <dbl>        <dbl>    <dbl>    <dbl>
## 1 Informed  Austral~    10.6  18318.        0.237  22179.  21779.
## 2 Informed  Canada     14.0  29608.        0.297  23711.  23353.
## 3 Informed  Denmark    13.1   5257.        12.2   23722.  23275
## 4 Informed  Germany    13.0  80255.        22.5   22163.  21938.
## 5 Informed  Ireland    19.8   3674.         5.23  20824.  20154.
## 6 Informed  Netherl~    13.7  15548.        37.4   23013.  22554.
## 7 Informed  United ~    13.5  58187.        24.0   21359.  20962.
## 8 Informed  United ~    20.0 269330.         2.80  29212.  28699.
## 9 Presumed  Austria    23.5   7927.         9.45  23876.  23415.
## 10 Presumed Belgium    21.9  10153.        30.7   22500.  22096.
## 11 Presumed Finland    18.4   5112.         1.51  21019.  20763
## 12 Presumed France    16.8  58056.        10.5   22603.  22211.
## 13 Presumed Italy     11.1  57360.        19.0   21554.  21195.
## 14 Presumed Norway    15.4   4386.         1.35  26448.  25769.
## 15 Presumed Spain     28.1  39666.         7.84  16933  16584.
## 16 Presumed Sweden    13.1   8789.         1.95  22415.  22094
## 17 Presumed Switzer~   14.2   7037.        17.0   27233  26931.
## # ... with 21 more variables: health_mean <dbl>, health_lag_mean <dbl>,
## #   pubhealth_mean <dbl>, roads_mean <dbl>, cerebvas_mean <dbl>,
## #   assault_mean <dbl>, external_mean <dbl>, txp_pop_mean <dbl>,
## #   donors_sd <dbl>, pop_sd <dbl>, pop_dens_sd <dbl>, gdp_sd <dbl>,
## #   gdp_lag_sd <dbl>, health_sd <dbl>, health_lag_sd <dbl>, pubhealth_sd <dbl>,
## #   roads_sd <dbl>, cerebvas_sd <dbl>, assault_sd <dbl>, external_sd <dbl>,
## #   txp_pop_sd <dbl>
```

```
p <- ggplot(data = by_country,
  mapping = aes(x = donors_mean,
                y = reorder(country, donors_mean),
                color = consent_law))
p + geom_point(size=3) +
  labs(x = "Donor Procurement Rate",
       y = "", color = "Consent Law") +
  theme(legend.position="top")
```



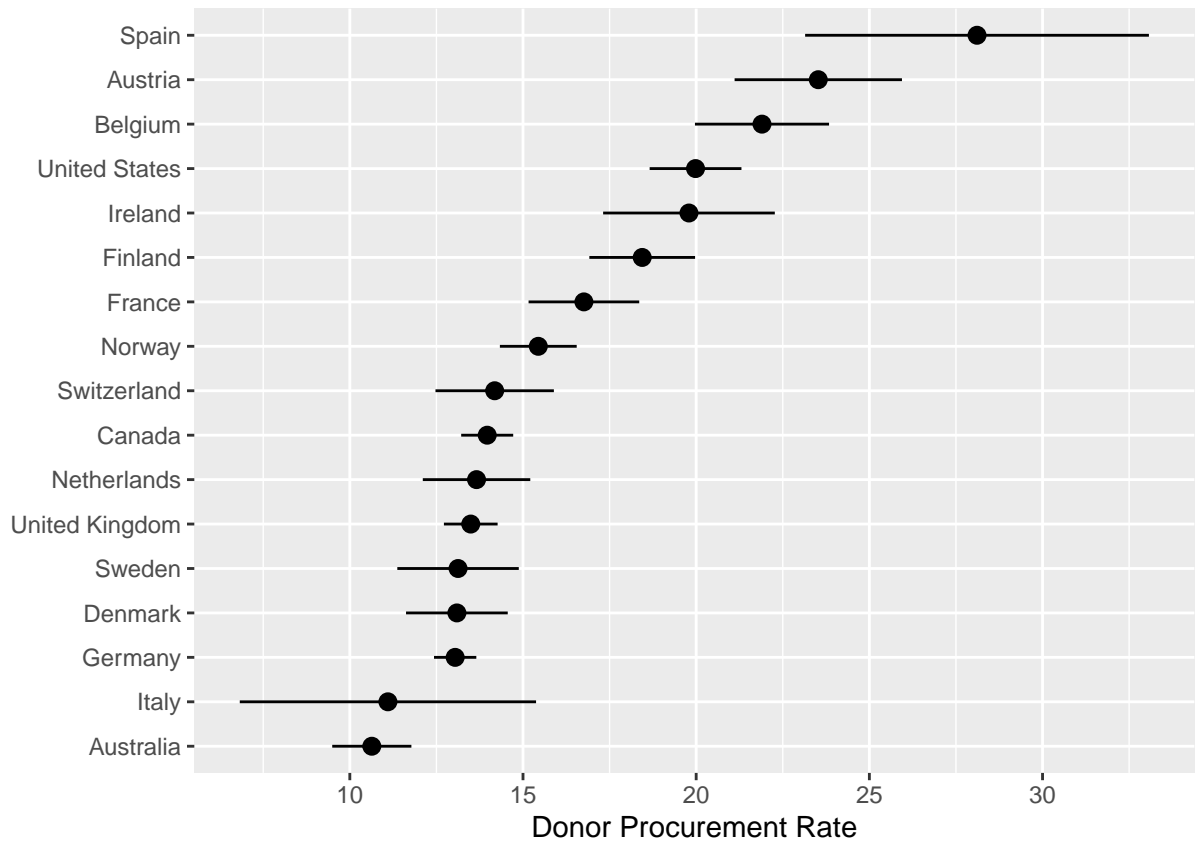
```
p <- ggplot(data = by_country,
            mapping = aes(x = donors_mean,
                          y = reorder(country, donors_mean)))

p + geom_point(size=3) +
  facet_wrap(~ consent_law, scales = "free_y", ncol = 1) +
  labs(x= "Donor Procurement Rate",
       y= "")
```



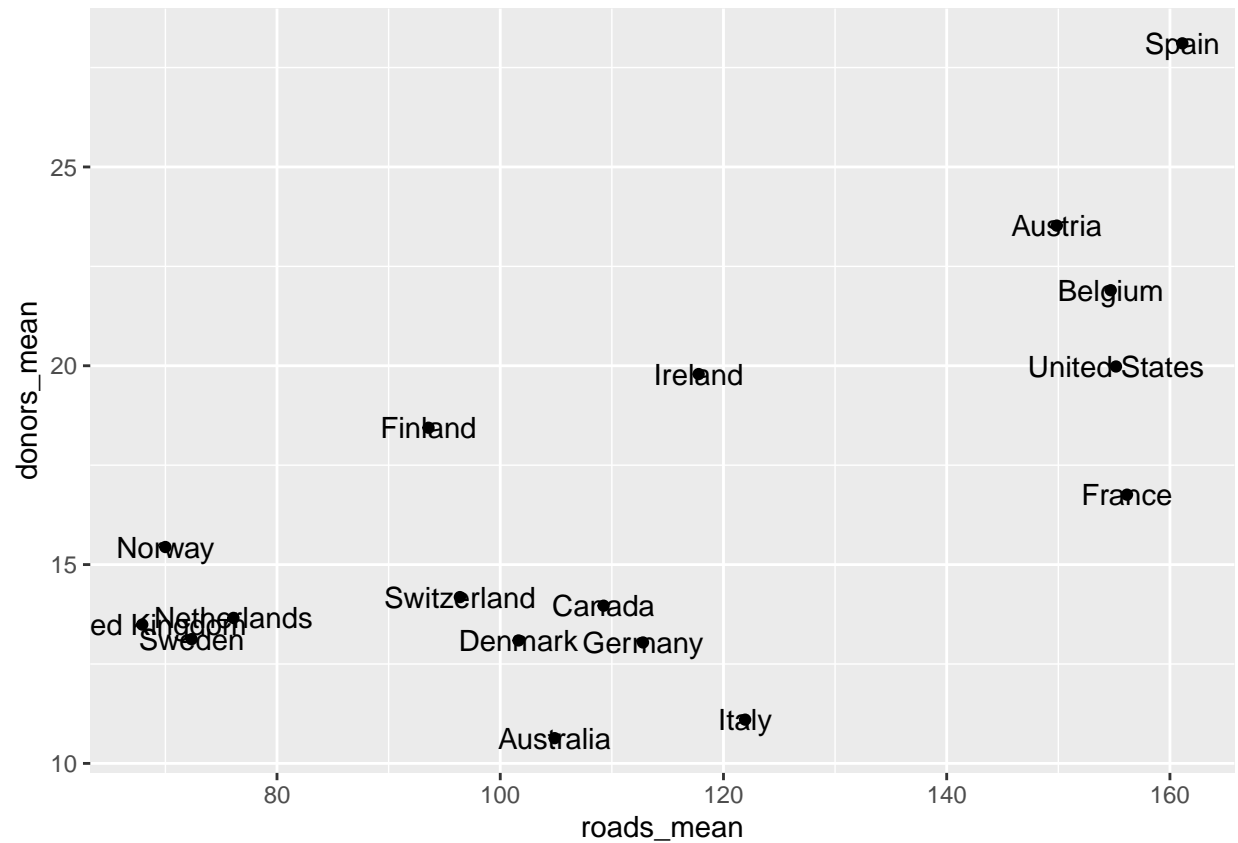
```
p <- ggplot(data = by_country, mapping = aes(x = reorder(country,
  donors_mean), y = donors_mean))

p + geom_pointrange(mapping = aes(ymin = donors_mean - donors_sd,
  ymax = donors_mean + donors_sd)) +
  labs(x= "", y= "Donor Procurement Rate") + coord_flip()
```



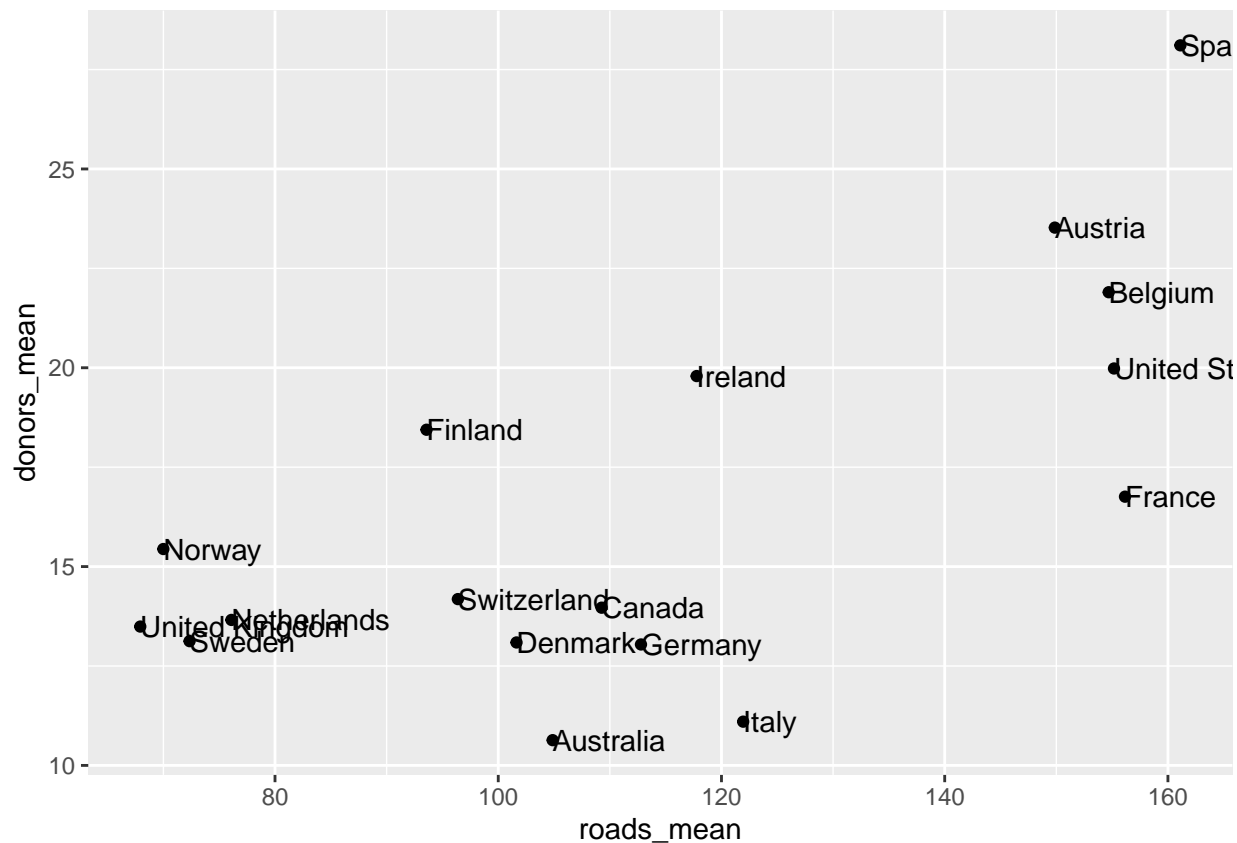
Plot text directly

```
p <- ggplot(data = by_country,
            mapping = aes(x = roads_mean, y = donors_mean))
p + geom_point() + geom_text(mapping = aes(label = country))
```



```
p <- ggplot(data = by_country,
            mapping = aes(x = roads_mean, y = donors_mean))

p + geom_point() + geom_text(mapping = aes(label = country), hjust = 0)
```



```
library(ggrepel)
```

```
## Warning: package 'ggrepel' was built under R version 3.6.3
```

```
elections_historic %>% select(2:7)
```

```
## # A tibble: 49 x 6
```

	year	winner	win_party	ec_pct	popular_pct	popular_margin
	<int>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
## 1	1824	John Quincy Adams	D.-R.	0.322	0.309	-0.104
## 2	1828	Andrew Jackson	Dem.	0.682	0.559	0.122
## 3	1832	Andrew Jackson	Dem.	0.766	0.547	0.178
## 4	1836	Martin Van Buren	Dem.	0.578	0.508	0.142
## 5	1840	William Henry Harrison	Whig	0.796	0.529	0.0605
## 6	1844	James Polk	Dem.	0.618	0.495	0.0145
## 7	1848	Zachary Taylor	Whig	0.562	0.473	0.0479
## 8	1852	Franklin Pierce	Dem.	0.858	0.508	0.0695
## 9	1856	James Buchanan	Dem.	0.588	0.453	0.122
## 10	1860	Abraham Lincoln	Rep.	0.594	0.396	0.101

```
## # ... with 39 more rows
```

```
p_title <- "Presidential Elections: Popular & Electoral College Margins"
```

```
p_subtitle <- "1824-2016"
```

```
p_caption <- "Data for 2016 are provisional."
```

```
x_label <- "Winner's share of Popular Vote"
```

```
y_label <- "Winner's share of Electoral College Votes"
```


```
p <- ggplot(elections_historic, aes(x = popular_pct, y = ec_pct,
```

```

    label = winner_label))

p + geom_hline(yintercept = 0.5, size = 1.4, color = "gray80") +
  geom_vline(xintercept = 0.5, size = 1.4, color = "gray80") +
  geom_point() +
  geom_text_repel() +
  scale_x_continuous(labels = scales::percent) +
  scale_y_continuous(labels = scales::percent) +
  labs(x = x_label, y = y_label, title = p_title, subtitle = p_subtitle,
       caption = p_caption)

```




Selective labels

```

p <- ggplot(data = by_country,
            mapping = aes(x = gdp_mean, y = health_mean))


p + geom_point() +
  geom_text_repel(data = subset(by_country, gdp_mean > 25000),
                  mapping = aes(label = country))

```

```
p <- ggplot(data = by_country,
            mapping = aes(x = gdp_mean, y = health_mean))


p + geom_point() +
  geom_text_repel(data = subset(by_country,
                                gdp_mean > 25000 | health_mean < 1500 |
                                country %in% "Belgium"),
                 mapping = aes(label = country))
```



```
organdata$ind <- organdata$ccode %in% c("Ita", "Spa") &
  organdata$year > 1998
```

```
p <- ggplot(data = organdata,
  mapping = aes(x = roads,
    y = donors, color = ind))
p + geom_point() +
  geom_text_repel(data = subset(organdata, ind),
    mapping = aes(label = ccode)) +
  guides(label = FALSE, color = FALSE)
```

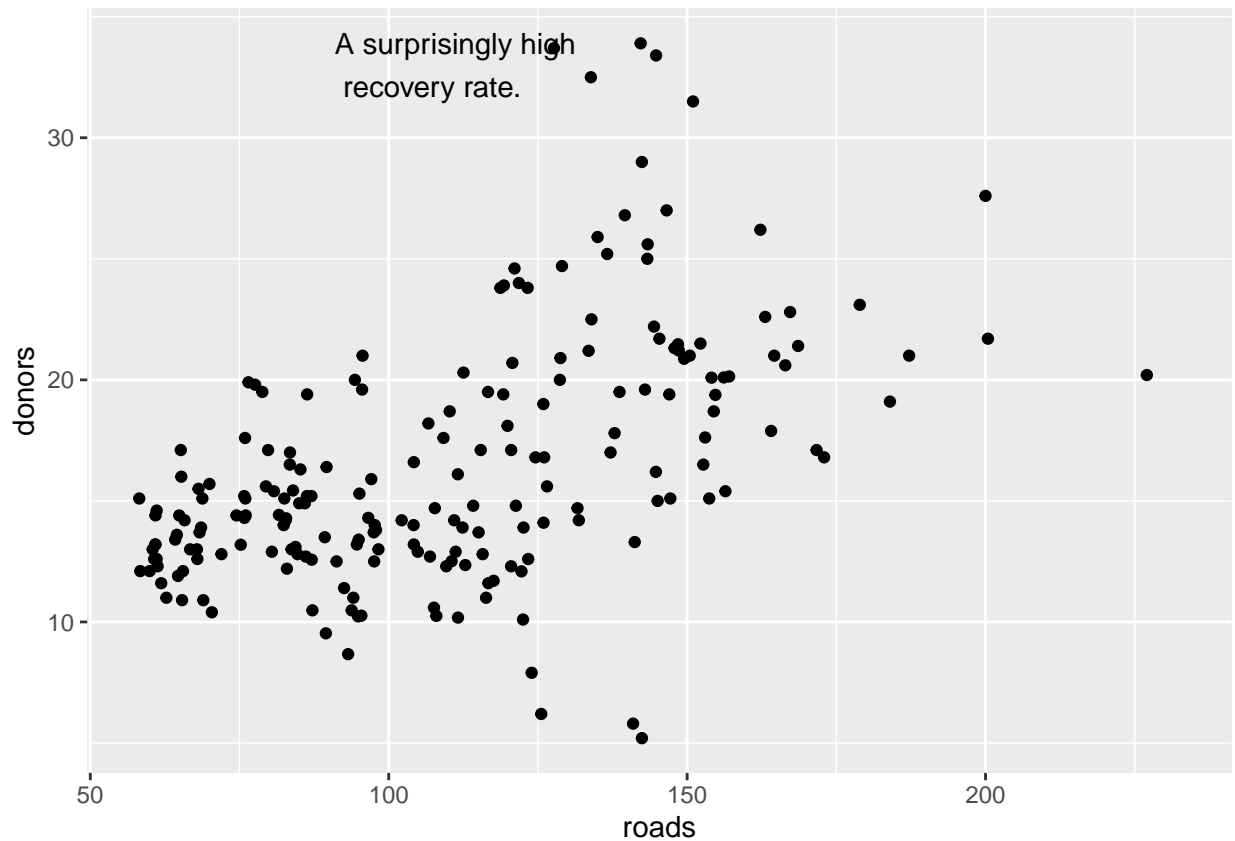
```
## Warning: Removed 34 rows containing missing values (geom_point).
```



Arbitrary annotation


```
p <- ggplot(data = organdata, mapping = aes(x = roads, y = donors))
p + geom_point() + annotate(geom = "text", x = 91, y = 33,
  label = "A surprisingly high \n recovery rate.",
  hjust = 0)
```

```
## Warning: Removed 34 rows containing missing values (geom_point).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = roads, y = donors))
p + geom_point() +
  annotate(geom = "rect", xmin = 125, xmax = 155,
          ymin = 30, ymax = 35, fill = "red", alpha = 0.2) +
  annotate(geom = "text", x = 157, y = 33,
          label = "A surprisingly high \n recovery rate.", hjust = 0)
```

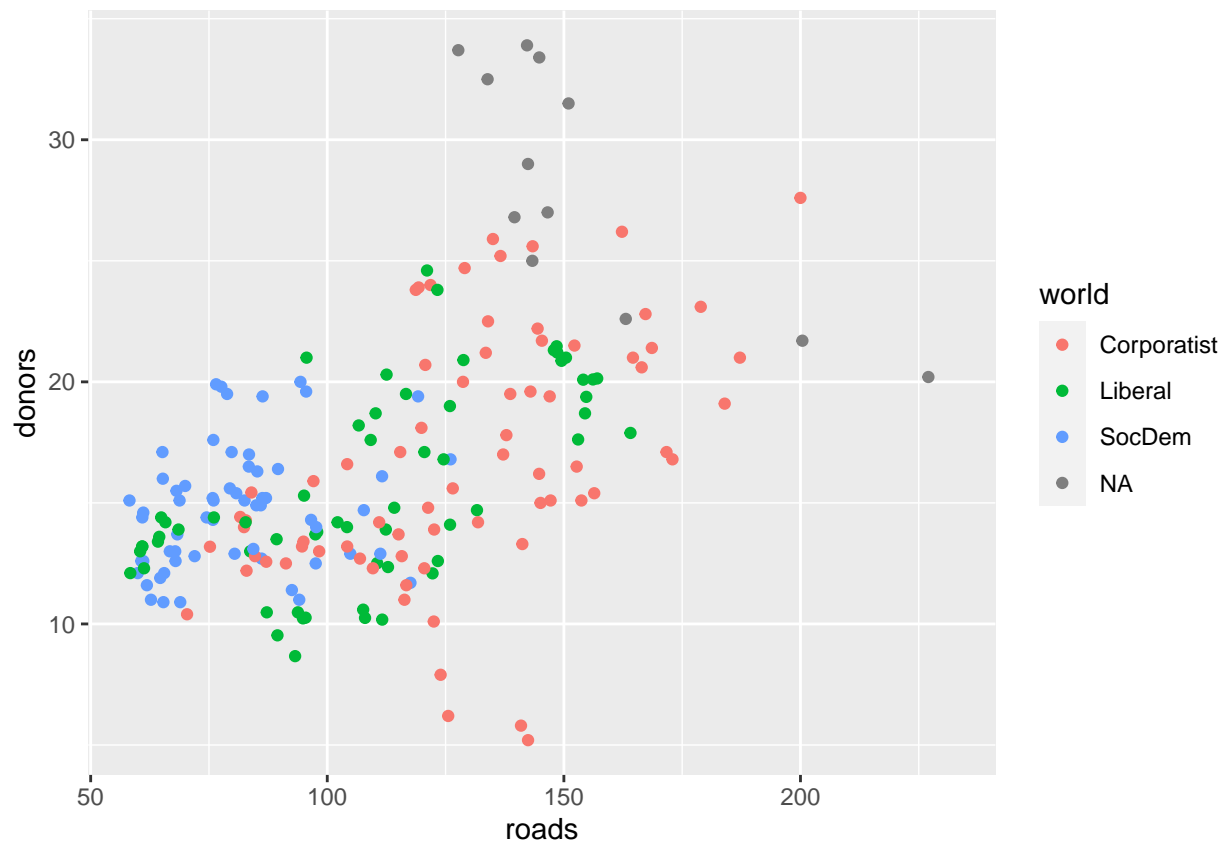
```
## Warning: Removed 34 rows containing missing values (geom_point).
```



Scales and Guides

```
p <- ggplot(data = organdata,  
            mapping = aes(x = roads,  
                          y = donors,  
                          color = world))  
p + geom_point()
```

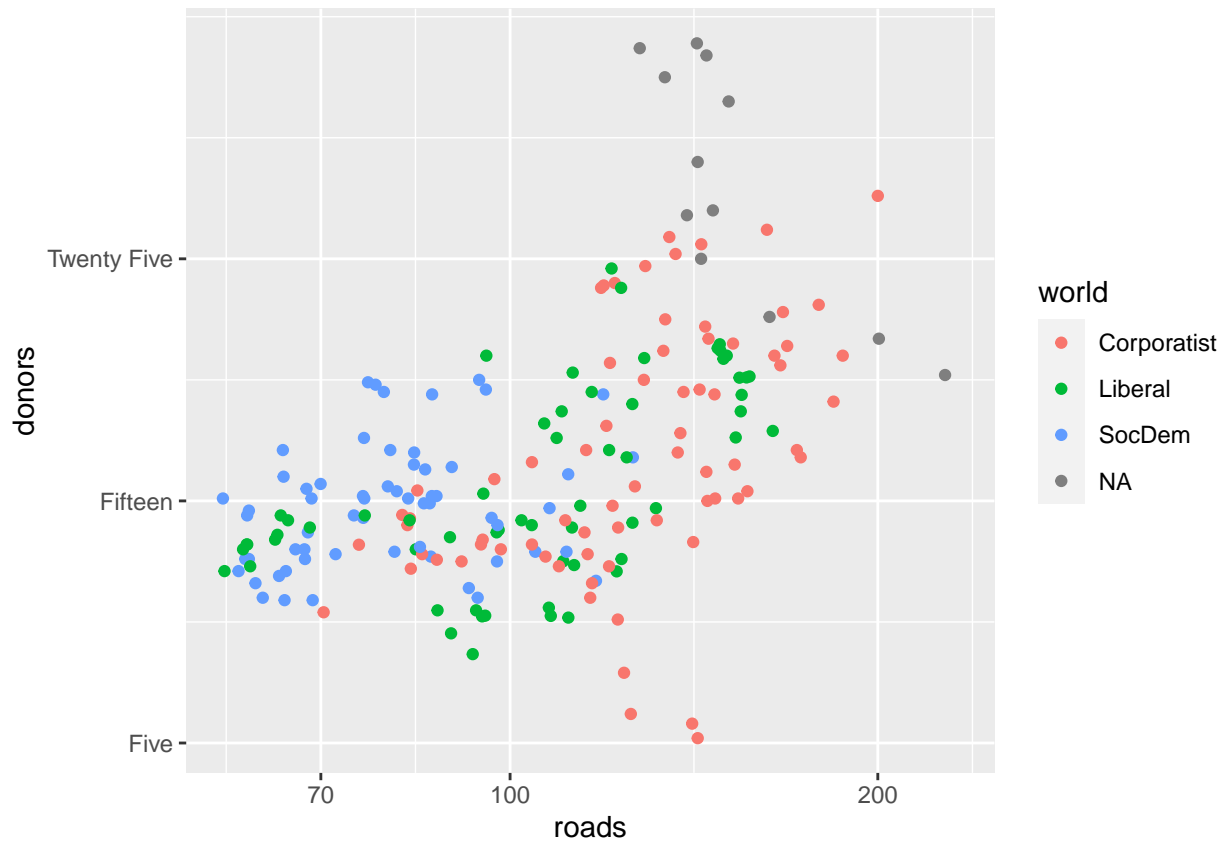
```
## Warning: Removed 34 rows containing missing values (geom_point).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = roads,
                          y = donors,
                          color = world))

p + geom_point() +
  scale_x_log10() +
  scale_y_continuous(breaks = c(5, 15, 25),
                     labels = c("Five", "Fifteen", "Twenty Five"))
```

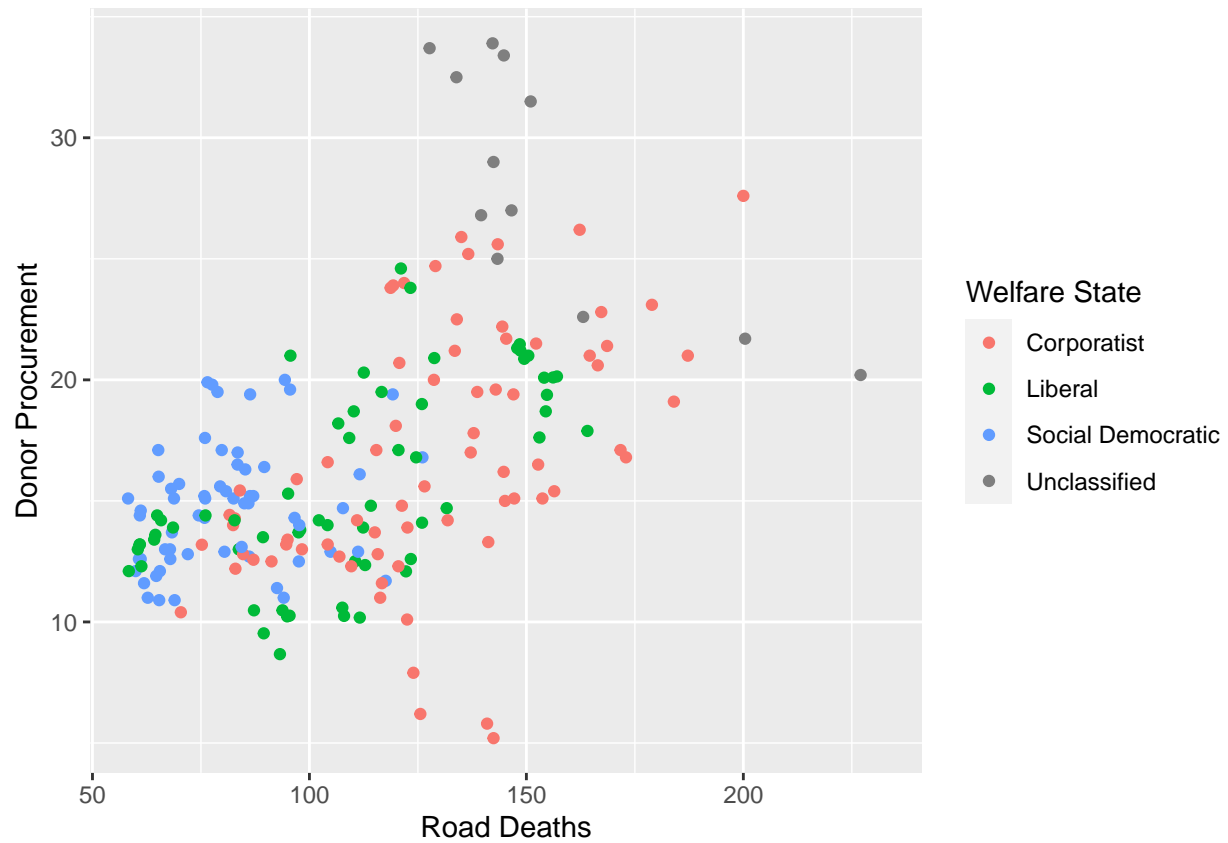
```
## Warning: Removed 34 rows containing missing values (geom_point).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = roads,
                          y = donors,
                          color = world))

p + geom_point() +
  scale_color_discrete(labels =
    c("Corporatist", "Liberal",
      "Social Democratic", "Unclassified")) +
  labs(x = "Road Deaths",
       y = "Donor Procurement",
       color = "Welfare State")
```

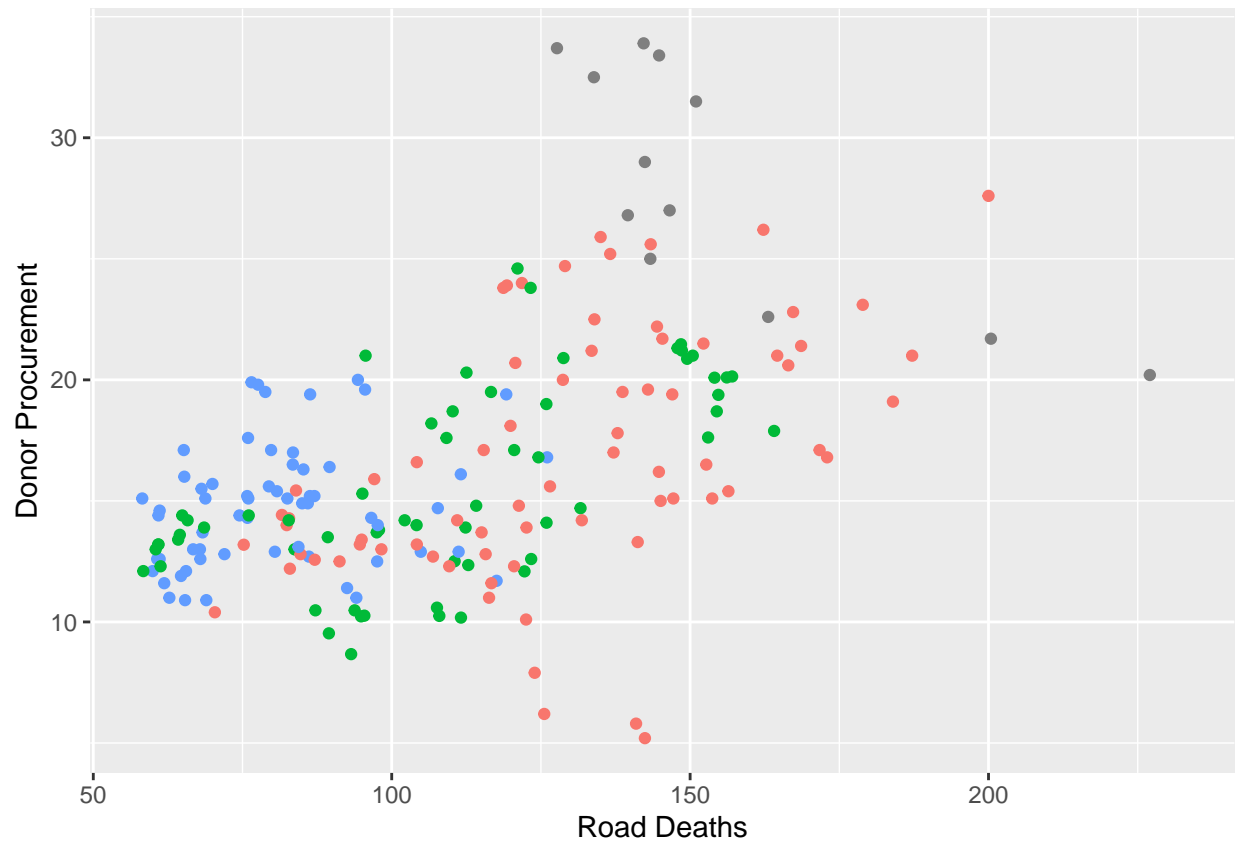
```
## Warning: Removed 34 rows containing missing values (geom_point).
```



```
p <- ggplot(data = organdata,
            mapping = aes(x = roads,
                          y = donors,
                          color = world))

p + geom_point() +
  labs(x = "Road Deaths",
       y = "Donor Procurement") +
  guides(color = FALSE)
```

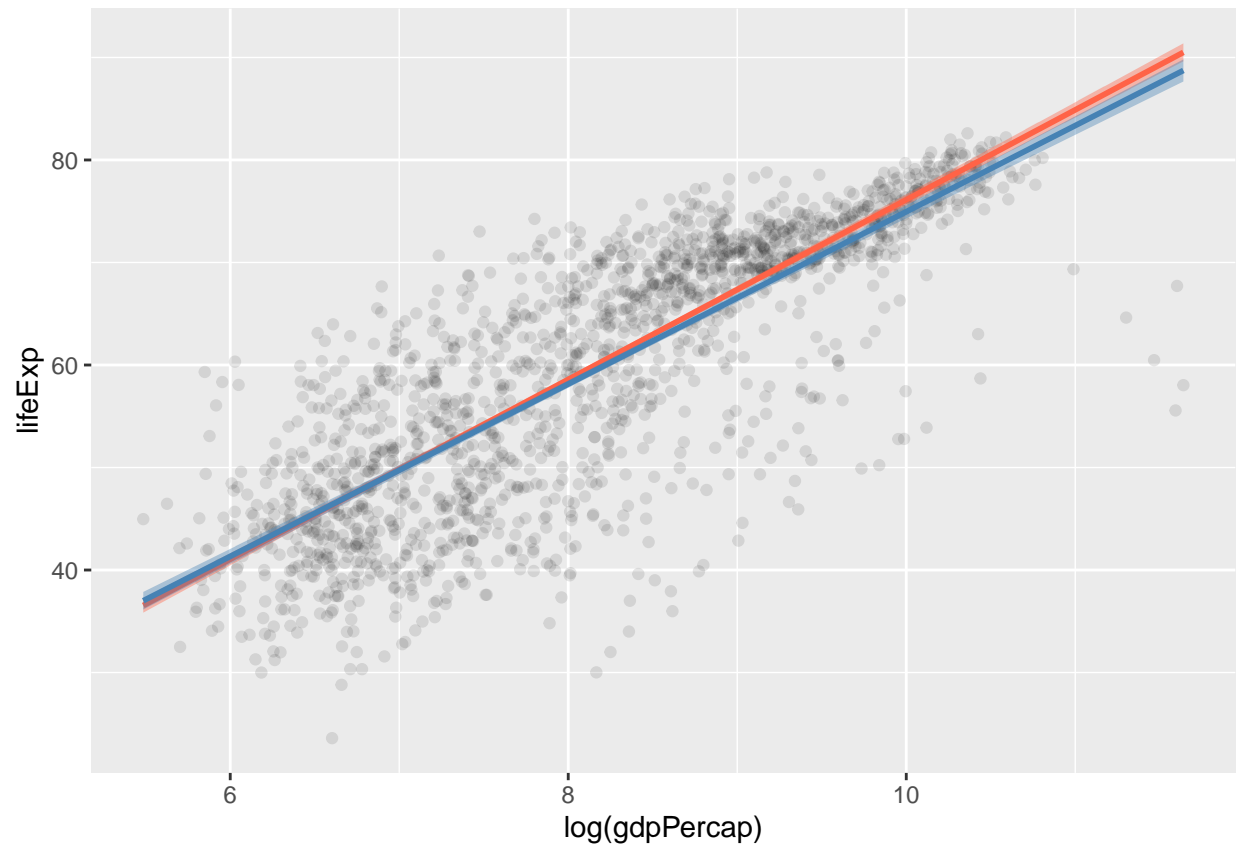
```
## Warning: Removed 34 rows containing missing values (geom_point).
```

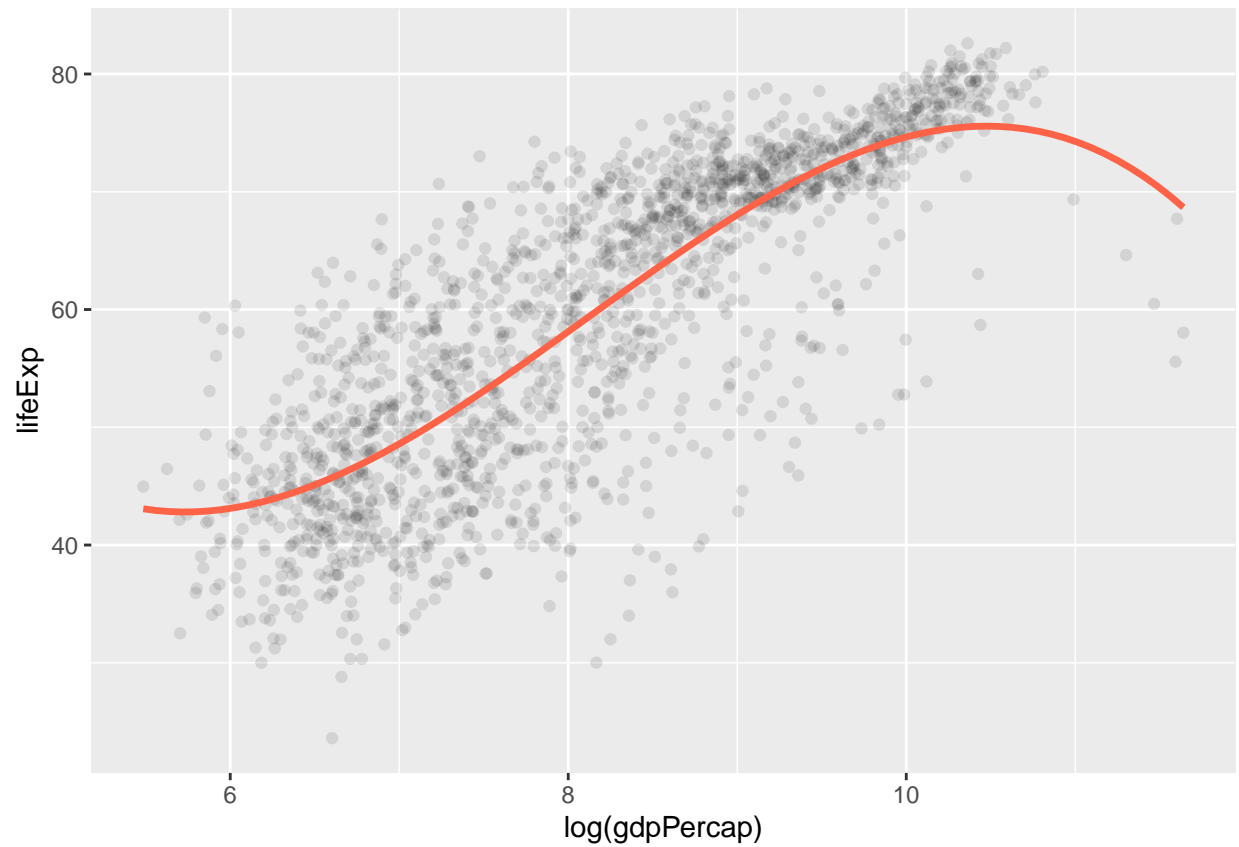
```
p <- ggplot(data = gapminder,
            mapping = aes(x = log(gdpPercap), y = lifeExp))

p + geom_point(alpha=0.1) +
  geom_smooth(color = "tomato", fill="tomato", method = MASS::rlm) +
  geom_smooth(color = "steelblue", fill="steelblue", method = "lm")

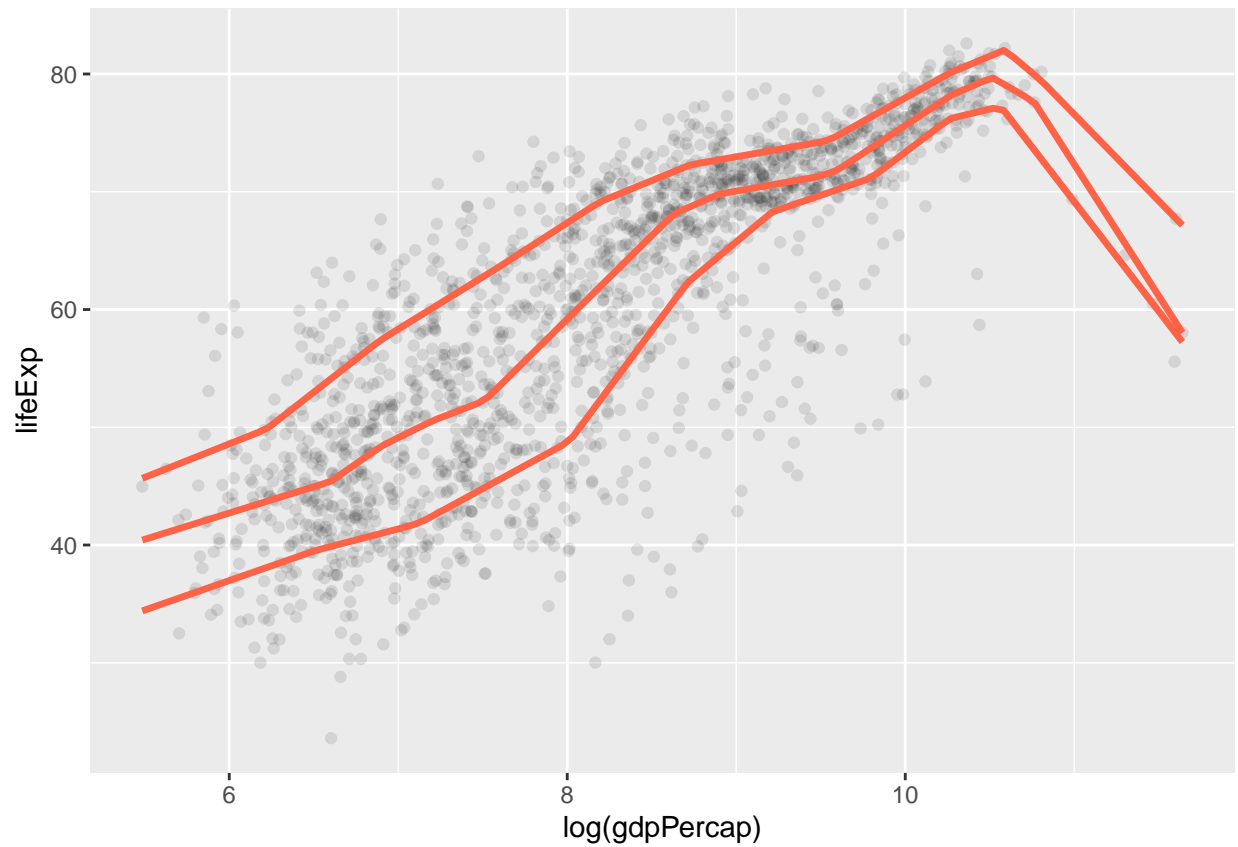
## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```



```
p + geom_point(alpha=0.1) +  
  geom_smooth(color = "tomato", method = "lm", size = 1.2,  
             formula = y ~ splines::bs(x, 3), se = FALSE)
```



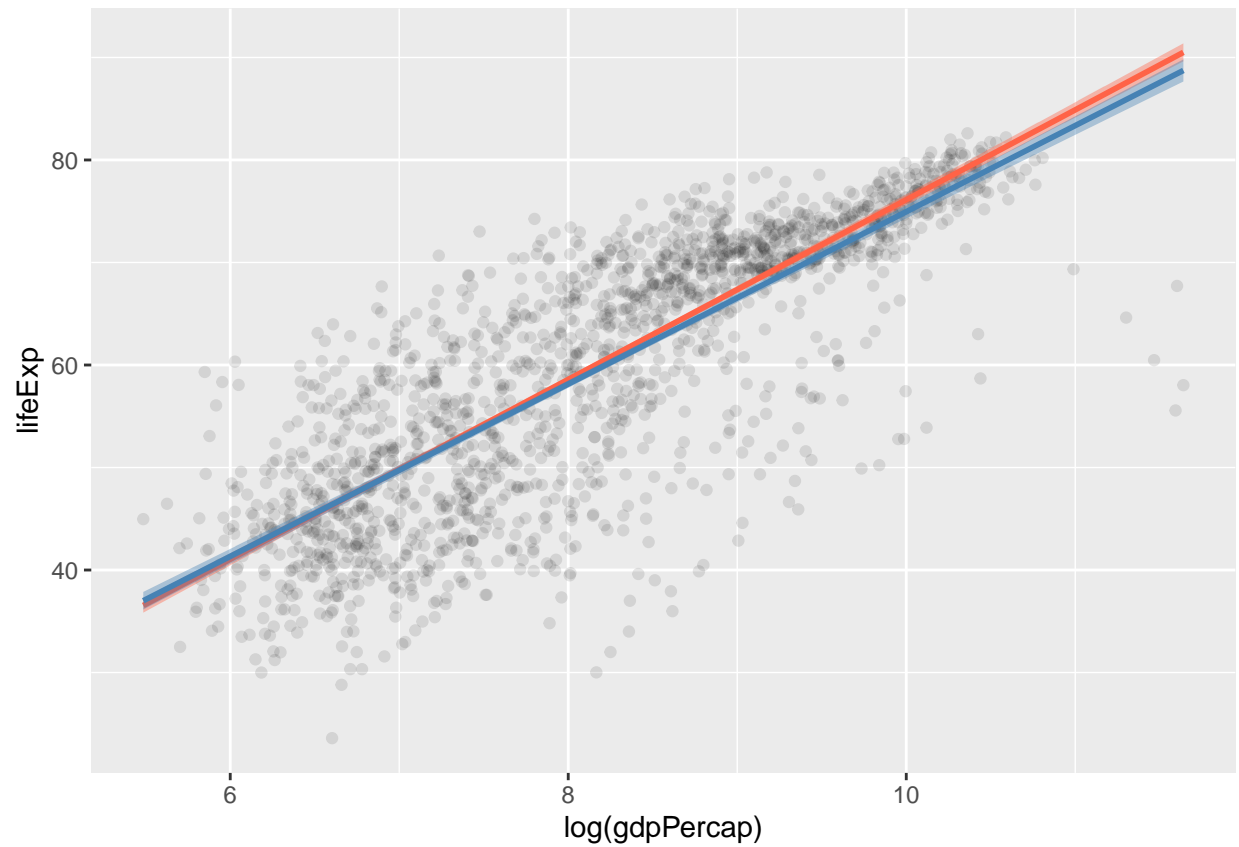
```
p + geom_point(alpha=0.1) +  
  geom_quantile(color = "tomato", size = 1.2, method = "rqss",  
               lambda = 1, quantiles = c(0.20, 0.5, 0.85))  
  
## Smoothing formula not specified. Using: y ~ qss(x, lambda = 1)
```




```
p <- ggplot(data = gapminder,
            mapping = aes(x = log(gdpPercap), y = lifeExp))

p + geom_point(alpha=0.1) +
  geom_smooth(color = "tomato", fill="tomato", method = MASS::rlm) +
  geom_smooth(color = "steelblue", fill="steelblue", method = "lm")


## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```



```
p + geom_point(alpha=0.1) +  
  geom_smooth(color = "tomato", method = "lm", size = 1.2,  
             formula = y ~ splines::bs(x, 3), se = FALSE)
```



```
p + geom_point(alpha=0.1) +  
  geom_quantile(color = "tomato", size = 1.2, method = "rqss",  
               lambda = 1, quantiles = c(0.20, 0.5, 0.85))  
  
## Smoothing formula not specified. Using: y ~ qss(x, lambda = 1)
```



Show several fits at once, with a legend

```
model_colors <- RColorBrewer::brewer.pal(3, "Set1")
model_colors

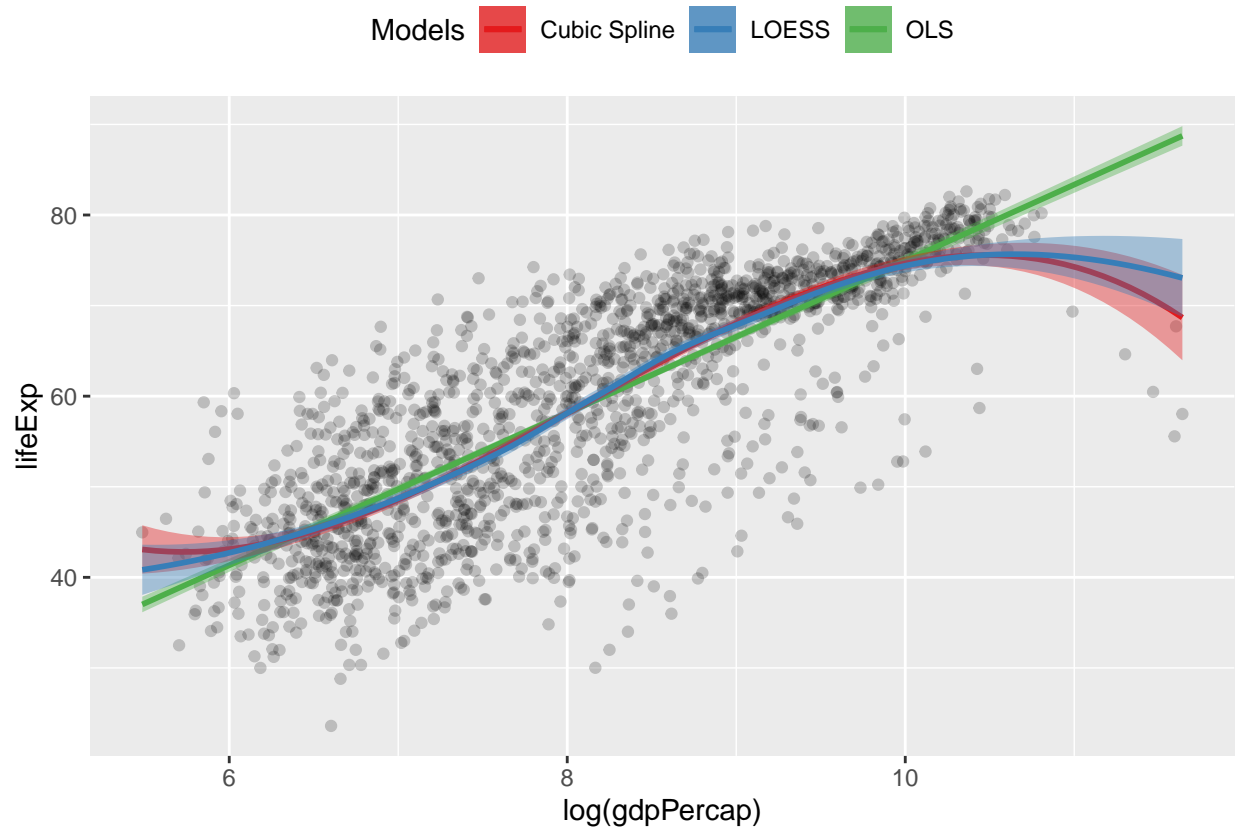
## [1] "#E41A1C" "#377EB8" "#4DAF4A"

p0 <- ggplot(data = gapminder,
             mapping = aes(x = log(gdpPercap), y = lifeExp))

p1 <- p0 + geom_point(alpha = 0.2) +
  geom_smooth(method = "lm", aes(color = "OLS", fill = "OLS")) +
  geom_smooth(method = "lm", formula = y ~ splines::bs(x, df = 3),
             aes(color = "Cubic Spline", fill = "Cubic Spline")) +
  geom_smooth(method = "loess",
             aes(color = "LOESS", fill = "LOESS"))

p1 + scale_color_manual(name = "Models", values = model_colors) +
  scale_fill_manual(name = "Models", values = model_colors) +
  theme(legend.position = "top")

## `geom_smooth()` using formula 'y ~ x'
## `geom_smooth()` using formula 'y ~ x'
```



Look inside model objects

```
## tibble [1,704 x 6] (S3: tbl_df/tbl/data.frame)
## $ country : Factor w/ 142 levels "Afghanistan",...: 1 1 ...
## $ continent: Factor w/ 5 levels "Africa","Americas",...: 3 3 ...
## $ year : int [1:1704] 1952 1957 ...
## $ lifeExp : num [1:1704] 28.8 ...
## $ pop : int [1:1704] 8425333 9240934 ...
## $ gdpPercap: num [1:1704] 779 ...
```

```
out <- lm(formula = lifeExp ~ gdpPercap + pop + continent,
          data = gapminder)
```

```
summary(out)
```

```
##
## Call:
## lm(formula = lifeExp ~ gdpPercap + pop + continent, data = gapminder)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
##	-49.161	-4.486	0.297	5.110	25.175

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	4.781e+01	3.395e-01	140.819	< 2e-16 ***
## gdpPercap	4.495e-04	2.346e-05	19.158	< 2e-16 ***


```
## pop                6.570e-09  1.975e-09   3.326 0.000901 ***
## continentAmericas 1.348e+01  6.000e-01  22.458 < 2e-16 ***
## continentAsia     8.193e+00  5.712e-01  14.342 < 2e-16 ***
## continentEurope   1.747e+01  6.246e-01  27.973 < 2e-16 ***
## continentOceania  1.808e+01  1.782e+00  10.146 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.365 on 1697 degrees of freedom
## Multiple R-squared:  0.5821, Adjusted R-squared:  0.5806
## F-statistic: 393.9 on 6 and 1697 DF,  p-value: < 2.2e-16
```

Generate predictions to graph

```
min_gdp <- min(gapminder$gdpPercap)
max_gdp <- max(gapminder$gdpPercap)
med_pop <- median(gapminder$pop)

pred_df <- expand.grid(gdpPercap = (seq(from = min_gdp,
                                       to = max_gdp,
                                       length.out = 100)),
                     pop = med_pop,
                     continent = c("Africa", "Americas",
                                   "Asia", "Europe", "Oceania"))

dim(pred_df)
```

```
## [1] 500   3
```

```
head(pred_df)
```

```
##   gdpPercap    pop continent
## 1  241.1659 7023596    Africa
## 2 1385.4282 7023596    Africa
## 3 2529.6905 7023596    Africa
## 4 3673.9528 7023596    Africa
## 5 4818.2150 7023596    Africa
## 6 5962.4773 7023596    Africa
```

```
pred_out <- predict(object = out,
                   newdata = pred_df,
                   interval = "predict")
head(pred_out)
```

```
##      fit      lwr      upr
## 1 47.96863 31.54775 64.38951
## 2 48.48298 32.06231 64.90365
## 3 48.99733 32.57670 65.41797
## 4 49.51169 33.09092 65.93245
## 5 50.02604 33.60497 66.44711
## 6 50.54039 34.11885 66.96193
```

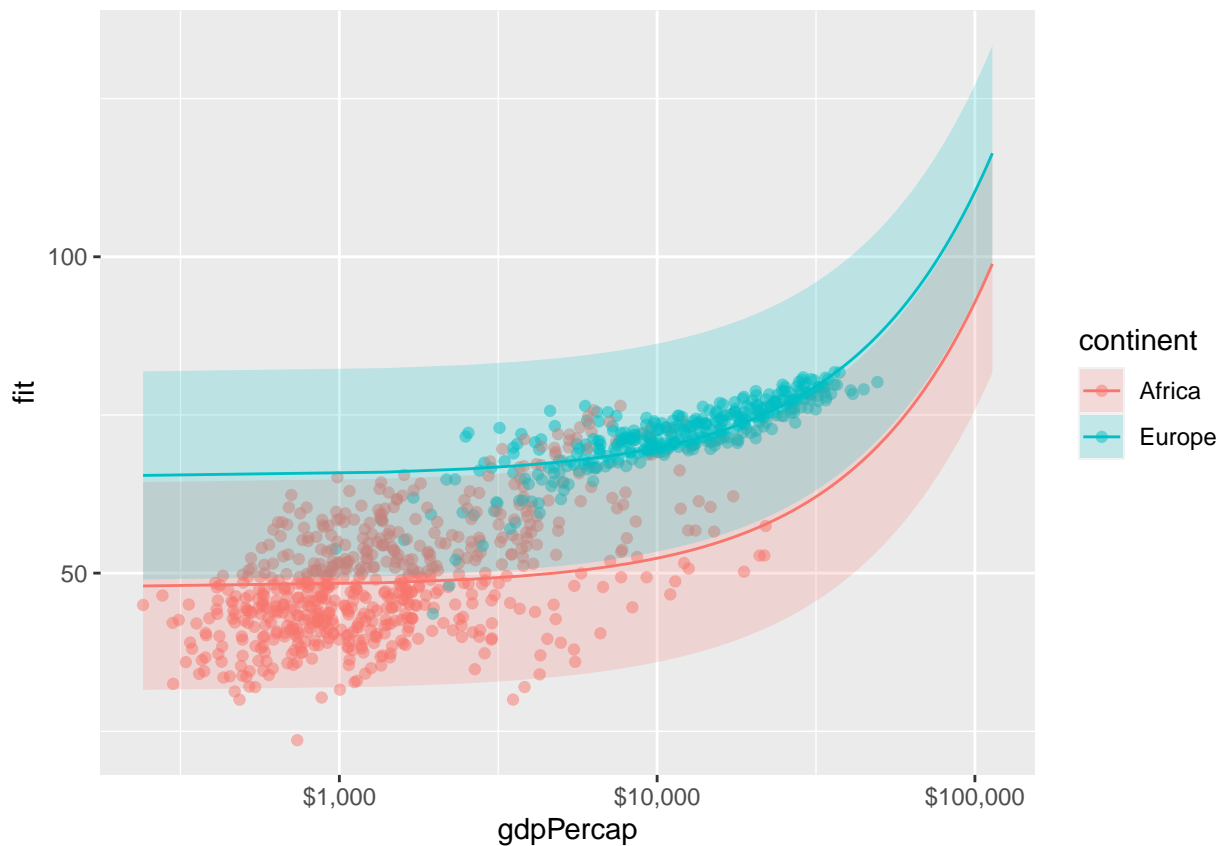
```
pred_df <- cbind(pred_df, pred_out)
head(pred_df)
```

```
##   gdpPercap    pop continent      fit      lwr      upr
```

```
## 1 241.1659 7023596 Africa 47.96863 31.54775 64.38951
## 2 1385.4282 7023596 Africa 48.48298 32.06231 64.90365
## 3 2529.6905 7023596 Africa 48.99733 32.57670 65.41797
## 4 3673.9528 7023596 Africa 49.51169 33.09092 65.93245
## 5 4818.2150 7023596 Africa 50.02604 33.60497 66.44711
## 6 5962.4773 7023596 Africa 50.54039 34.11885 66.96193
```

```
p <- ggplot(data = subset(pred_df, continent %in% c("Europe", "Africa")),
  aes(x = gdpPerCap,
      y = fit, ymin = lwr, ymax = upr,
      color = continent,
      fill = continent,
      group = continent))

p + geom_point(data = subset(gapminder,
  continent %in% c("Europe", "Africa")),
  aes(x = gdpPerCap, y = lifeExp,
      color = continent),
  alpha = 0.5,
  inherit.aes = FALSE) +
  geom_line() +
  geom_ribbon(alpha = 0.2, color = FALSE) +
  scale_x_log10(labels = scales::dollar)
```



Tidy model objects with broom

```
library(broom)
```

```
## Warning: package 'broom' was built under R version 3.6.3
```


```
out_comp <- tidy(out)
out_comp %>% round_df()
```

```
## # A tibble: 7 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	47.8	0.34	141.	0
## 2	gdpPercap	0	0	19.2	0
## 3	pop	0	0	3.33	0
## 4	continentAmericas	13.5	0.6	22.5	0
## 5	continentAsia	8.19	0.57	14.3	0
## 6	continentEurope	17.5	0.62	28.0	0
## 7	continentOceania	18.1	1.78	10.2	0

```
p <- ggplot(out_comp, mapping = aes(x = term,
                                     y = estimate))
```

```
p + geom_point() + coord_flip()
```




```
out_conf <- tidy(out, conf.int = TRUE)
out_conf %>% round_df()
```

```
## # A tibble: 7 x 7
##   term                estimate std.error statistic p.value conf.low conf.high
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)        47.8      0.34     141.      0      47.2     48.5
## 2 gdpPercap           0         0        19.2      0         0         0
## 3 pop                 0         0         3.33      0         0         0
## 4 continentAmericas  13.5      0.6      22.5      0      12.3     14.6
## 5 continentAsia       8.19     0.57     14.3      0       7.07     9.31
## 6 continentEurope    17.5      0.62     28.0      0      16.2     18.7
## 7 continentOceania   18.1      1.78     10.2      0      14.6     21.6
```

```
## out_conf <- subset(out_conf, term %nin% "(Intercept)")
## out_conf$nicelabs <- prefix_strip(out_conf$term, "continent")
```

```
out_conf <- out_conf %>%
  filter(term %nin% "(Intercept)") %>%
  mutate(nicelabs = prefix_strip(term, "continent")) %>%
  select(nicelabs, everything())
```

```
p <- ggplot(out_conf, mapping = aes(x = reorder(nicelabs, estimate),
                                           y = estimate, ymin = conf.low, ymax = conf.high))
p + geom_pointrange() + coord_flip() + labs(x="", y="OLS Estimate")
```



Get observation-level statistics with `augment()`

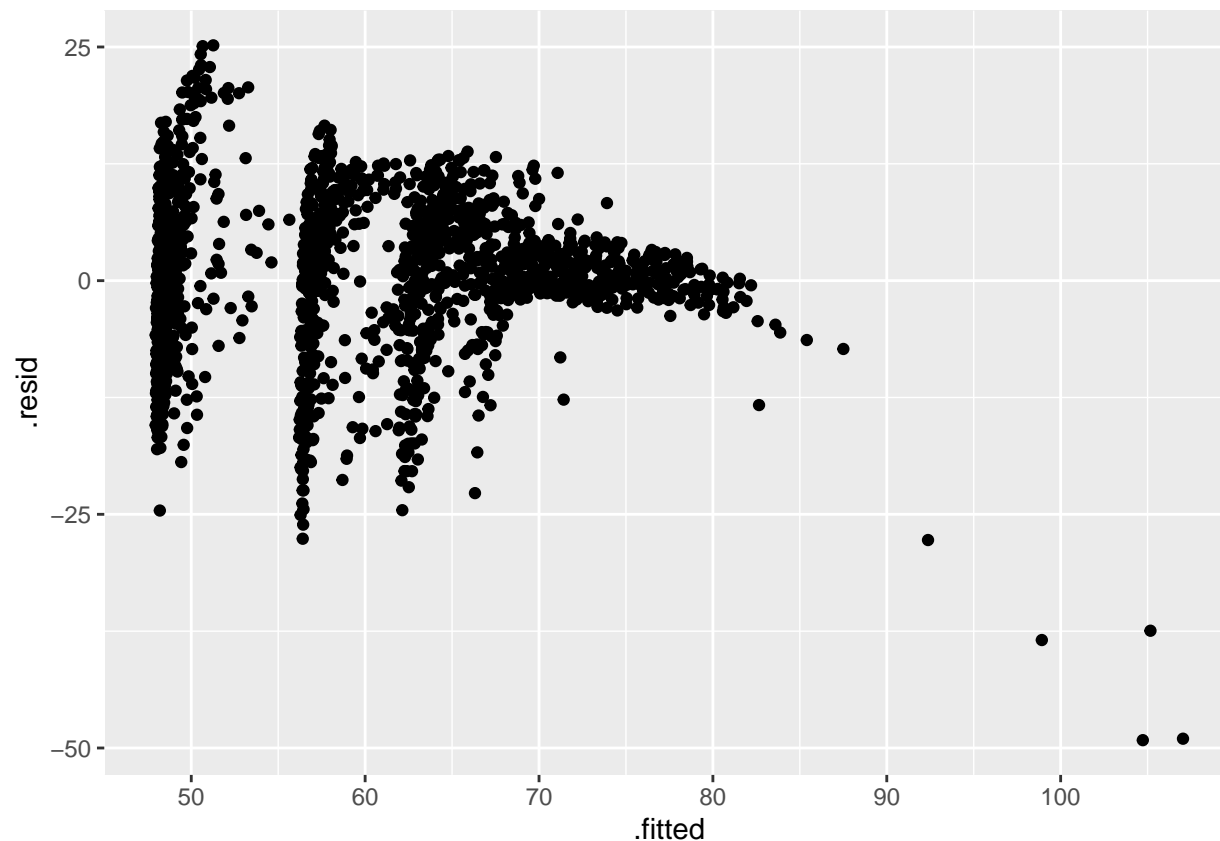
```
out_aug <- augment(out)
head(out_aug) %>% round_df()
```

```
## # A tibble: 6 x 10
##   lifeExp gdpPercap      pop continent .fitted .resid .hat .sigma .cooks
##   <dbl>    <dbl>    <dbl> <fct>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  28.8      779.  8425333 Asia     56.4 -27.6    0  8.34  0.01
## 2  30.3      821.  9240934 Asia     56.4 -26.1    0  8.34   0
## 3  32        853. 10267083 Asia     56.5 -24.5    0  8.35   0
## 4  34.0      836. 11537966 Asia     56.5 -22.4    0  8.35   0
## 5  36.1      740. 13079460 Asia     56.4 -20.3    0  8.35   0
## 6  38.4      786. 14880372 Asia     56.5 -18.0    0  8.36   0
## # ... with 1 more variable: .std.resid <dbl>
```

```
out_aug <- augment(out, data = gapminder)
head(out_aug) %>% round_df()
```

```
## # A tibble: 6 x 12
##   country continent year lifeExp      pop gdpPercap .fitted .resid .hat .sigma
##   <fct>    <fct>    <dbl>  <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Afghanis~ Asia     1952   28.8 8.43e6    779.    56.4 -27.6    0  8.34
## 2 Afghanis~ Asia     1957   30.3 9.24e6    821.    56.4 -26.1    0  8.34
## 3 Afghanis~ Asia     1962   32   1.03e7    853.    56.5 -24.5    0  8.35
## 4 Afghanis~ Asia     1967   34.0 1.15e7    836.    56.5 -22.4    0  8.35
## 5 Afghanis~ Asia     1972   36.1 1.31e7    740.    56.4 -20.3    0  8.35
## 6 Afghanis~ Asia     1977   38.4 1.49e7    786.    56.5 -18.0    0  8.36
## # ... with 2 more variables: .cooks <dbl>, .std.resid <dbl>
```

```
p <- ggplot(data = out_aug,
            mapping = aes(x = .fitted, y = .resid))
p + geom_point()
```



Get model-level statistics with `glance()`

```
glance(out) %>% round_df()
```

```
## # A tibble: 1 x 12
##   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
##   <dbl>      <dbl> <dbl>    <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     0.58        0.58  8.37     394.     0     6 -6034. 12084. 12127.
## # ... with 3 more variables: deviance <dbl>, df.residual <dbl>, nobs <dbl>
```


```
library(survival)
```

```
## Warning: package 'survival' was built under R version 3.6.3
```

```
out_cph <- coxph(Surv(time, status) ~ age + sex, data = lung)
out_surv <- survfit(out_cph)
```

```
out_tidy <- tidy(out_surv)
```

```
p <- ggplot(data = out_tidy, mapping = aes(time, estimate))
p + geom_line() +
  geom_ribbon(mapping = aes(ymin = conf.low, ymax = conf.high), alpha = .2)
```



Grouped analysis and list columns

```
eu77 <- gapminder %>% filter(continent == "Europe", year == 1977)
```

```
fit <- lm(lifeExp ~ log(gdpPercap), data = eu77)
summary(fit)
```

```
##
## Call:
## lm(formula = lifeExp ~ log(gdpPercap), data = eu77)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.4956 -1.0306  0.0935  1.1755  3.7125
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    29.489     7.161   4.118 0.000306 ***
## log(gdpPercap)  4.488     0.756   5.936 2.17e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.114 on 28 degrees of freedom
## Multiple R-squared:  0.5572, Adjusted R-squared:  0.5414
## F-statistic: 35.24 on 1 and 28 DF,  p-value: 2.173e-06
```

```
out_le <- gapminder %>%
  group_by(continent, year) %>%
  nest()
```

```
out_le
```

```
## # A tibble: 60 x 3
## # Groups:   continent, year [60]
##   continent year data
##   <fct>      <int> <list>
## 1 Asia      1952 <tibble [33 x 4]>
## 2 Asia      1957 <tibble [33 x 4]>
## 3 Asia      1962 <tibble [33 x 4]>
## 4 Asia      1967 <tibble [33 x 4]>
## 5 Asia      1972 <tibble [33 x 4]>
## 6 Asia      1977 <tibble [33 x 4]>
## 7 Asia      1982 <tibble [33 x 4]>
## 8 Asia      1987 <tibble [33 x 4]>
## 9 Asia      1992 <tibble [33 x 4]>
## 10 Asia     1997 <tibble [33 x 4]>
## # ... with 50 more rows
```

```
out_le %>% filter(continent == "Europe" & year == 1977) %>%
  unnest(cols = c(data))
```

```
## # A tibble: 30 x 6
## # Groups:   continent, year [1]
##   continent year country      lifeExp      pop gdpPercap
##   <fct>      <int> <fct>      <dbl>      <int>      <dbl>
## 1 Europe    1977 Albania      68.9    2509048      3533.
## 2 Europe    1977 Austria      72.2    7568430     19749.
## 3 Europe    1977 Belgium      72.8    9821800     19118.
## 4 Europe    1977 Bosnia and Herzegovina 69.9    4086000      3528.
## 5 Europe    1977 Bulgaria      70.8    8797022      7612.
## 6 Europe    1977 Croatia      70.6    4318673     11305.
## 7 Europe    1977 Czech Republic 70.7   10161915     14800.
## 8 Europe    1977 Denmark      74.7    5088419     20423.
## 9 Europe    1977 Finland      72.5    4738902     15605.
## 10 Europe   1977 France      73.8   53165019     18293.
## # ... with 20 more rows
```

```
fit_ols <- function(df) {
  lm(lifeExp ~ log(gdpPercap), data = df)
}
```

```
out_le <- gapminder %>%
  group_by(continent, year) %>%
  nest() %>%
  mutate(model = map(data, fit_ols))
```

```
out_le
```

```
## # A tibble: 60 x 4
## # Groups:   continent, year [60]
##   continent year data      model
```



```
##      <fct>      <int> <list>          <list>
## 1 Asia        1952 <tibble [33 x 4]> <lm>
## 2 Asia        1957 <tibble [33 x 4]> <lm>
## 3 Asia        1962 <tibble [33 x 4]> <lm>
## 4 Asia        1967 <tibble [33 x 4]> <lm>
## 5 Asia        1972 <tibble [33 x 4]> <lm>
## 6 Asia        1977 <tibble [33 x 4]> <lm>
## 7 Asia        1982 <tibble [33 x 4]> <lm>
## 8 Asia        1987 <tibble [33 x 4]> <lm>
## 9 Asia        1992 <tibble [33 x 4]> <lm>
## 10 Asia       1997 <tibble [33 x 4]> <lm>
## # ... with 50 more rows

fit_ols <- function(df) {
  lm(lifeExp ~ log(gdpPercap), data = df)
}

out_tidy <- gapminder %>%
  group_by(continent, year) %>%
  nest() %>%
  mutate(model = map(data, fit_ols),
         tidied = map(model, tidy)) %>%
  unnest(cols = c(tidied)) %>%
  filter(term %nin% "(Intercept)" &
         continent %nin% "Oceania")

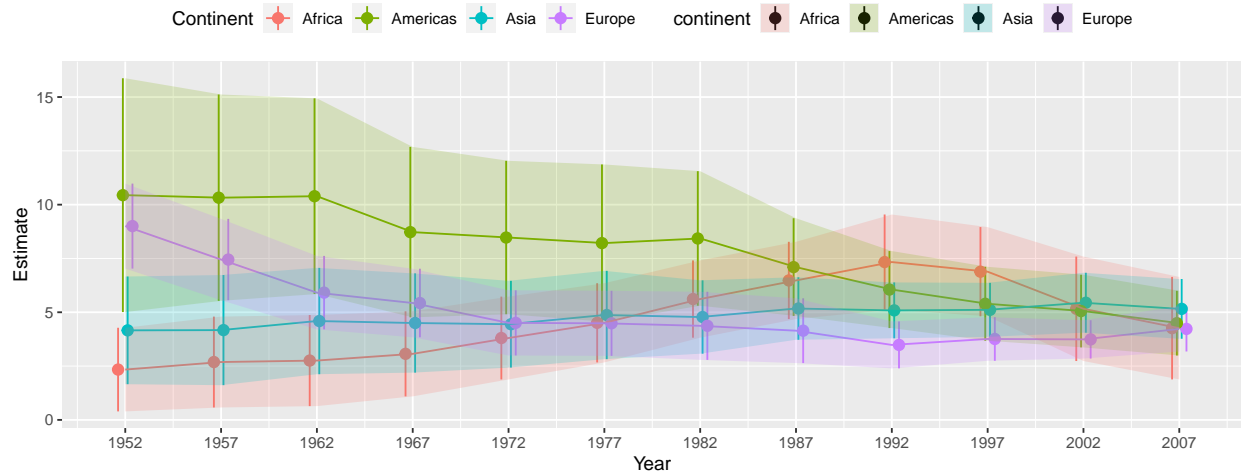
out_tidy %>%
  ungroup() %>%
  sample_n(5)

## # A tibble: 5 x 9
##   continent year data      model term      estimate std.error statistic p.value
##   <fct>      <int> <list>    <list> <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Africa    2002 <tibble ~ <lm> log(gd~    5.16      1.21       4.25 9.20e-5
## 2 Africa    1987 <tibble ~ <lm> log(gd~    6.48      0.898     7.22 2.75e-9
## 3 Americas  1997 <tibble ~ <lm> log(gd~    5.39      0.859     6.28 2.08e-6
## 4 Asia      1952 <tibble ~ <lm> log(gd~    4.16      1.25      3.33 2.28e-3
## 5 Americas  1967 <tibble ~ <lm> log(gd~    8.72      1.98      4.41 2.05e-4

p <- ggplot(data = out_tidy,
            mapping = aes(x = year, y = estimate,
                          ymin = estimate - 2*std.error,
                          ymax = estimate + 2*std.error,
                          color = continent, group = continent,
                          fill = continent))

p + geom_pointrange(position = position_dodge(width = 1)) +
  geom_line() +
  geom_ribbon(mapping = aes(x = year,
                          ymin = estimate - 2*std.error,
                          ymax = estimate + 2*std.error,
                          group = continent,
                          fill = continent),
            alpha = 0.2,
```

```
inherit.aes = FALSE) +
scale_x_continuous(breaks = unique(gapminder$year)) +
theme(legend.position = "top") +
labs(x = "Year", y = "Estimate", color = "Continent")
```



Grouped Analysis: PCA Example

On the full dataset ...

```
mw_pca <- midwest %>%
  group_by(state) %>%
  select_if(is.numeric) %>%
  select(-PID)
```

```
mw_pca
```

```
## # A tibble: 437 x 25
## # Groups:   state [5]
##   state area poptotal popdensity popwhite popblack popamerindian popasian
##   <chr> <dbl> <int> <dbl> <int> <int> <int> <int>
## 1 IL 0.052 66090 1271. 63917 1702 98 249
## 2 IL 0.014 10626 759 7054 3496 19 48
## 3 IL 0.022 14991 681. 14477 429 35 16
## 4 IL 0.017 30806 1812. 29344 127 46 150
## 5 IL 0.018 5836 324. 5264 547 14 5
## 6 IL 0.05 35688 714. 35157 50 65 195
## 7 IL 0.017 5322 313. 5298 1 8 15
## 8 IL 0.027 16805 622. 16519 111 30 61
## 9 IL 0.024 13437 560. 13384 16 8 23
## 10 IL 0.058 173025 2983. 146506 16559 331 8033
## # ... with 427 more rows, and 17 more variables: popother <int>,
## # percwhite <dbl>, percblack <dbl>, percamerindian <dbl>, percasian <dbl>,
## # percother <dbl>, popadults <int>, perchs <dbl>, percollege <dbl>,
## # percprof <dbl>, poppovertyknown <int>, percpovertyknown <dbl>,
## # percbelowpoverty <dbl>, percchildbelowpovert <dbl>, percadultpoverty <dbl>,
## # percelderlypoverty <dbl>, inmetro <int>
```

```
do_pca <- function(df){
  prcomp(df,
    center = TRUE, scale = TRUE)
}

out_pca <- mw_pca %>%
  ungroup() %>%
  select(-state) %>%
  do_pca()

summary(out_pca)

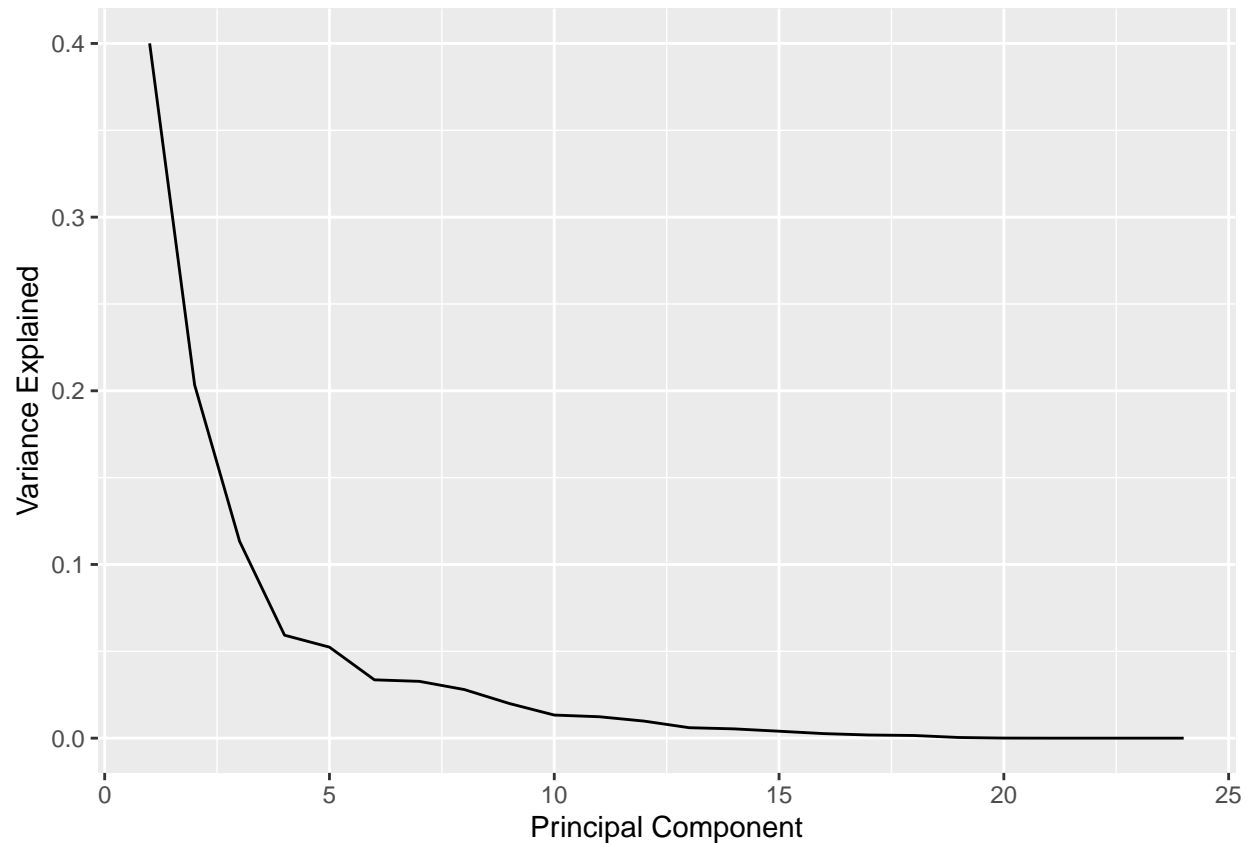
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.0986 2.2096 1.6495 1.19289 1.12159 0.89776 0.8859
## Proportion of Variance 0.4001 0.2034 0.1134 0.05929 0.05241 0.03358 0.0327
## Cumulative Proportion 0.4001 0.6035 0.7168 0.77614 0.82856 0.86214 0.8948
##              PC8      PC9      PC10      PC11      PC12      PC13      PC14
## Standard deviation  0.81948 0.69212 0.5650 0.54394 0.48541 0.38000 0.35833
## Proportion of Variance 0.02798 0.01996 0.0133 0.01233 0.00982 0.00602 0.00535
## Cumulative Proportion 0.92283 0.94278 0.9561 0.96841 0.97823 0.98425 0.98960
##              PC15      PC16      PC17      PC18      PC19      PC20      PC21
## Standard deviation  0.30948 0.25009 0.20879 0.19244 0.09654 0.03473 0.01328
## Proportion of Variance 0.00399 0.00261 0.00182 0.00154 0.00039 0.00005 0.00001
## Cumulative Proportion 0.99359 0.99619 0.99801 0.99955 0.99994 0.99999 1.00000
##              PC22      PC23      PC24
## Standard deviation  0.003862 2.886e-09 5.193e-16
## Proportion of Variance 0.000000 0.000e+00 0.000e+00
## Cumulative Proportion 1.000000 1.000e+00 1.000e+00

tidy_pca <- tidy(out_pca, matrix = "pcs")

tidy_pca

## # A tibble: 24 x 4
##       PC std.dev percent cumulative
##   <dbl>   <dbl>   <dbl>   <dbl>
## 1     1     3.10    0.400     0.400
## 2     2     2.21    0.203     0.603
## 3     3     1.65    0.113     0.717
## 4     4     1.19    0.0593    0.776
## 5     5     1.12    0.0524    0.829
## 6     6     0.898    0.0336    0.862
## 7     7     0.886    0.0327    0.895
## 8     8     0.819    0.0280    0.923
## 9     9     0.692    0.0200    0.943
## 10    10     0.565    0.0133    0.956
## # ... with 14 more rows

tidy_pca %>%
  ggplot(aes(x = PC, y = percent)) +
  geom_line() +
  labs(x = "Principal Component", y = "Variance Explained")
```



... or nested by state

```
mw_pca <- mw_pca %>%
  group_by(state) %>%
  nest()
```

```
mw_pca
```

```
## # A tibble: 5 x 2
## # Groups:   state [5]
##   state data
##   <chr> <list>
## 1 IL    <tibble [102 x 24]>
## 2 IN    <tibble [92 x 24]>
## 3 MI    <tibble [83 x 24]>
## 4 OH    <tibble [88 x 24]>
## 5 WI    <tibble [72 x 24]>
```

```
state_pca <- mw_pca %>%
  mutate(pca = map(data, do_pca))
```

```
state_pca
```

```
## # A tibble: 5 x 3
## # Groups:   state [5]
##   state data      pca
```

```

##   <chr> <list>           <list>
## 1 IL   <tibble [102 x 24]> <prcomp>
## 2 IN   <tibble [92 x 24]>  <prcomp>
## 3 MI   <tibble [83 x 24]>  <prcomp>
## 4 OH   <tibble [88 x 24]>  <prcomp>
## 5 WI   <tibble [72 x 24]>  <prcomp>

do_tidy <- function(pr){
  broom::tidy(pr, matrix = "pcs")
}

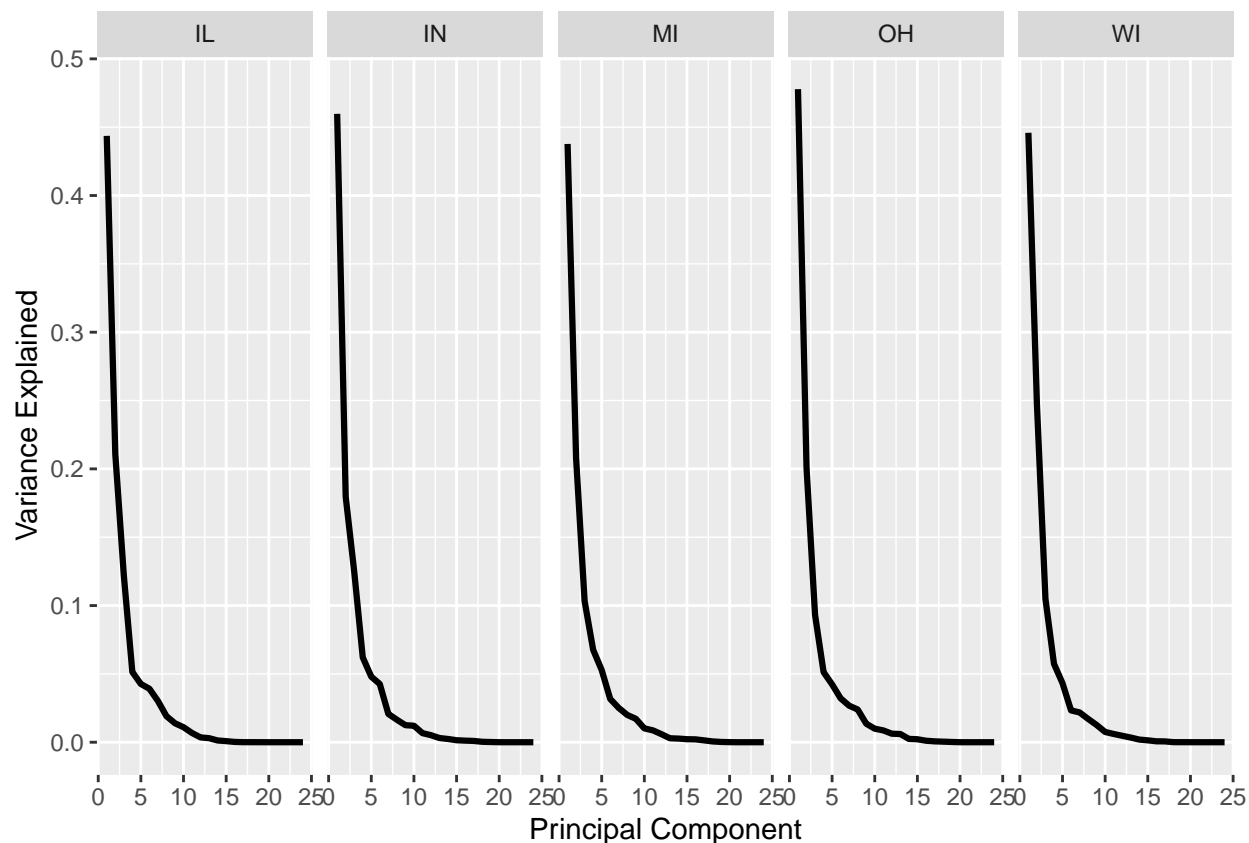
state_pca <- mw_pca %>%
  mutate(pca = map(data, do_pca),
         pcs = map(pca, do_tidy))

state_pca

## # A tibble: 5 x 4
## # Groups:   state [5]
##   state data          pca      pcs
##   <chr> <list>         <list> <list>
## 1 IL   <tibble [102 x 24]> <prcomp> <tibble [24 x 4]>
## 2 IN   <tibble [92 x 24]> <prcomp> <tibble [24 x 4]>
## 3 MI   <tibble [83 x 24]> <prcomp> <tibble [24 x 4]>
## 4 OH   <tibble [88 x 24]> <prcomp> <tibble [24 x 4]>
## 5 WI   <tibble [72 x 24]> <prcomp> <tibble [24 x 4]>

state_pca %>%
  unnest(cols = c(pcs)) %>%
  ggplot(aes(x = PC, y = percent)) +
  geom_line(size = 1.1) +
  facet_wrap(~ state, nrow = 1) +
  labs(x = "Principal Component",
       y = "Variance Explained")

```



```
do_aug <- function(pr){
  broom::augment(pr)
}


state_pca <- mw_pca %>%
  mutate(pca = map(data, do_pca),
         pcs = map(pca, do_tidy),
         fitted = map(pca, do_aug))

state_pca

## # A tibble: 5 x 5
## # Groups:   state [5]
##   state data          pca      pcs      fitted
##   <chr> <list>         <list> <list>    <list>
## 1 IL   <tibble [102 x 24]> <prcomp> <tibble [24 x 4]> <tibble [102 x 25]>
## 2 IN   <tibble [92 x 24]> <prcomp> <tibble [24 x 4]> <tibble [92 x 25]>
## 3 MI   <tibble [83 x 24]> <prcomp> <tibble [24 x 4]> <tibble [83 x 25]>
## 4 OH   <tibble [88 x 24]> <prcomp> <tibble [24 x 4]> <tibble [88 x 25]>
## 5 WI   <tibble [72 x 24]> <prcomp> <tibble [24 x 4]> <tibble [72 x 25]>


state_pca %>%
  unnest(cols = c(fitted)) %>%
  ggplot(aes(x = .fittedPC1,
             y = .fittedPC2)) +
  geom_point() +
```

```
facet_wrap(~ state) +
labs(x = "First Principal Component",
     y = "Second Principal Component")
```



Grouped PCA in a single sequence

```
midwest %>%
  group_by(state) %>%
  select_if(is.numeric) %>%
  select(-PID) %>%
  nest() %>%
  mutate(pca = map(data, do_pca),
         pcs = map(pca, do_tidy),
         fitted = map(pca, do_aug)) %>%
  unnest(cols = c(fitted)) %>%
  add_column(county = midwest$county) %>%
  ggplot(mapping = aes(x = .fittedPC2,
                      y = .fittedPC3,
                      label = county)) +
  geom_text(size = 1.1) +
  labs(x = "Second Principal Component",
       y = "Third Principal Component") +
  theme_minimal() + facet_wrap(~ state, ncol = 2)
```



Plot marginal effects

Note that calculating marginal effects can take some time!

```
library(margins)
```

```
## Warning: package 'margins' was built under R version 3.6.3
```

```
gss_sm$polviews_m <- releval(gss_sm$polviews, ref = "Moderate")
```

```
out_bo <- glm(obama ~ polviews_m + sex*race,
              family = "binomial", data = gss_sm)
summary(out_bo)
```

```
##
```

```
## Call:
```

```
## glm(formula = obama ~ polviews_m + sex * race, family = "binomial",
##      data = gss_sm)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2.9045  -0.5541   0.1772   0.5418   2.2437
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.296493   0.134091   2.211  0.02703 *
## polviews_mExtremely Liberal  2.372950   0.525045   4.520 6.20e-06 ***
```



```
## polviews_mLiberal          2.600031    0.356666    7.290 3.10e-13 ***
## polviews_mSlightly Liberal    1.293172    0.248435    5.205 1.94e-07 ***
## polviews_mSlightly Conservative -1.355277    0.181291   -7.476 7.68e-14 ***
## polviews_mConservative      -2.347463    0.200384  -11.715 < 2e-16 ***
## polviews_mExtremely Conservative -2.727384    0.387210   -7.044 1.87e-12 ***
## sexFemale                   0.254866    0.145370    1.753 0.07956 .
## raceBlack                   3.849526    0.501319    7.679 1.61e-14 ***
## raceOther                   -0.002143    0.435763   -0.005 0.99608
## sexFemale:raceBlack         -0.197506    0.660066   -0.299 0.76477
## sexFemale:raceOther         1.574829    0.587657    2.680 0.00737 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2247.9 on 1697 degrees of freedom
## Residual deviance: 1345.9 on 1686 degrees of freedom
## (1169 observations deleted due to missingness)
## AIC: 1369.9
##
## Number of Fisher Scoring iterations: 6
```

```
bo_m <- margins(out_bo)
summary(bo_m)
```

```
##           factor      AME      SE      z      p      lower
## polviews_mConservative -0.4119 0.0283 -14.5394 0.0000 -0.4674
## polviews_mExtremely Conservative -0.4538 0.0420 -10.7971 0.0000 -0.5361
## polviews_mExtremely Liberal 0.2681 0.0295 9.0996 0.0000 0.2103
## polviews_mLiberal 0.2768 0.0229 12.0736 0.0000 0.2319
## polviews_mSlightly Conservative -0.2658 0.0330 -8.0596 0.0000 -0.3304
## polviews_mSlightly Liberal 0.1933 0.0303 6.3896 0.0000 0.1340
## raceBlack 0.4032 0.0173 23.3568 0.0000 0.3694
## raceOther 0.1247 0.0386 3.2297 0.0012 0.0490
## sexFemale 0.0443 0.0177 2.5073 0.0122 0.0097
## upper
## -0.3564
## -0.3714
## 0.3258
## 0.3218
## -0.2011
## 0.2526
## 0.4371
## 0.2005
## 0.0789
```

```
bo_gg <- as_tibble(summary(bo_m))
prefixes <- c("polviews_m", "sex")
bo_gg$factor <- prefix_strip(bo_gg$factor, prefixes)
bo_gg$factor <- prefix_replace(bo_gg$factor, "race", "Race: ")


bo_gg %>% select(factor, AME, lower, upper)
```

```
## # A tibble: 9 x 4
##   factor      AME    lower    upper
```

```
##   <chr>                <dbl>   <dbl>   <dbl>
## 1 Conservative        -0.412  -0.467  -0.356
## 2 Extremely Conservative -0.454  -0.536  -0.371
## 3 Extremely Liberal     0.268   0.210   0.326
## 4 Liberal              0.277   0.232   0.322
## 5 Slightly Conservative -0.266  -0.330  -0.201
## 6 Slightly Liberal      0.193   0.134   0.253
## 7 Race: Black           0.403   0.369   0.437
## 8 Race: Other           0.125   0.0490  0.200
## 9 Female                0.0443  0.00967 0.0789
```

```
p <- ggplot(data = bo_gg, aes(x = reorder(factor, AME),
                             y = AME, ymin = lower, ymax = upper))

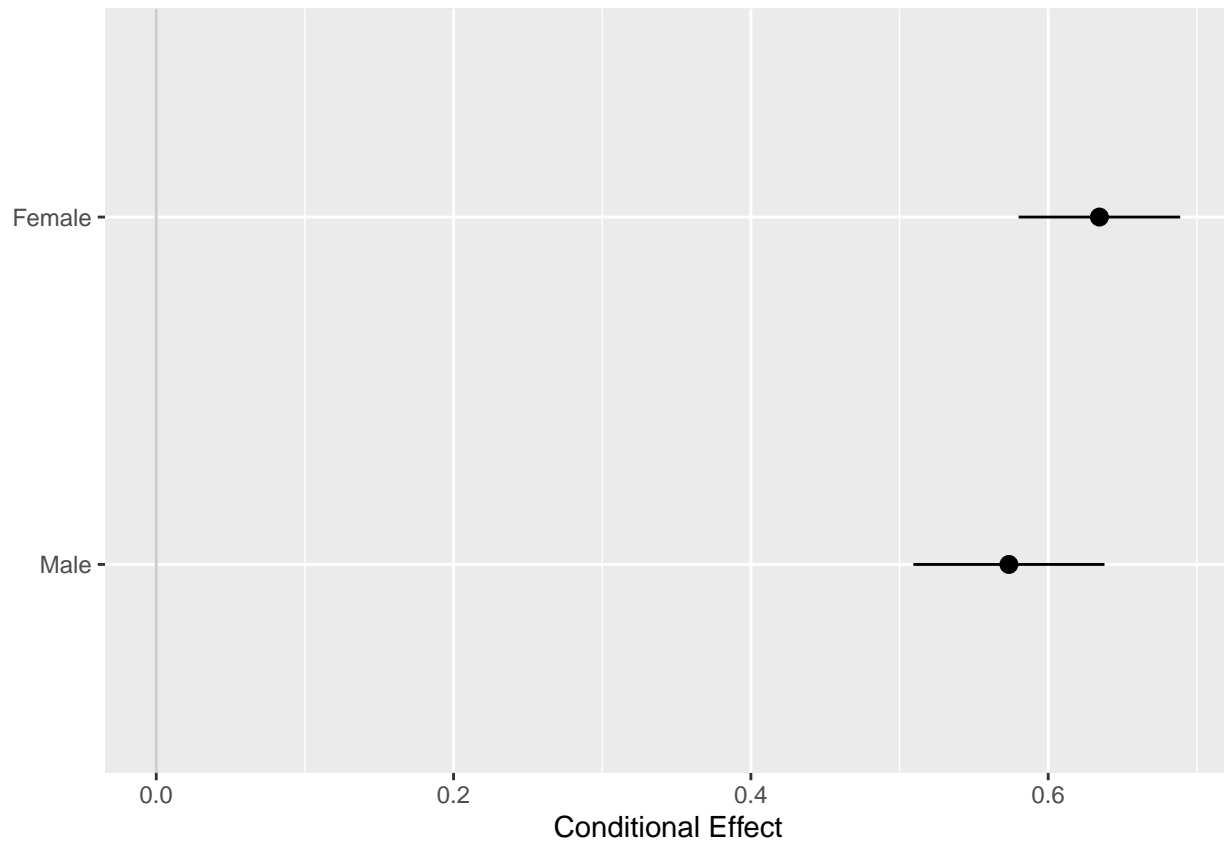
p + geom_hline(yintercept = 0, color = "gray80") +
  geom_pointrange() + coord_flip() +
  labs(x = NULL, y = "Average Marginal Effect")
```



```
pv_cp <- cplot(out_bo, x = "sex", draw = FALSE)

p <- ggplot(data = pv_cp, aes(x = reorder(xvals, yvals),
                             y = yvals, ymin = lower, ymax = upper))

p + geom_hline(yintercept = 0, color = "gray80") +
  geom_pointrange() + coord_flip() +
  labs(x = NULL, y = "Conditional Effect")
```



Plots from complex surveys

```
library(survey)
```

```
## Warning: package 'survey' was built under R version 3.6.3
```

```
## Loading required package: grid
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##     expand, pack, unpack
```

```
##
```

```
## Attaching package: 'survey'
```

```
## The following object is masked from 'package:graphics':
```

```
##
```

```
##     dotchart
```

```
library(srvyr)
```

```
## Warning: package 'srvyr' was built under R version 3.6.3
```

```
##
```

```
## Attaching package: 'srvyr'
```

```
## The following object is masked from 'package:stats':
##
## filter
```

```
options(survey.lonely.psu = "adjust")
options(na.action="na.pass")
```

```
gss_wt <- subset(gss_lon, year > 1974) %>%
  mutate(stratvar = interaction(year, vstrat)) %>%
  as_survey_design(ids = vpsu,
                  strata = stratvar,
                  weights = wtssall,
                  nest = TRUE)
```

```
out_grp <- gss_wt %>%
  filter(year %in% seq(1976, 2016, by = 4)) %>%
  group_by(year, race, degree) %>%
  summarize(prop = survey_mean(na.rm = TRUE))
```

```
out_grp
```

```
## # A tibble: 162 x 5
## # Groups:   year, race [30]
##   year race degree      prop prop_se
##   <dbl> <fct> <fct>      <dbl> <dbl>
## 1 1976 White Lt High School 0.327 0.0160
## 2 1976 White High School 0.517 0.0161
## 3 1976 White Junior College 0.0128 0.00298
## 4 1976 White Bachelor 0.101 0.00955
## 5 1976 White Graduate 0.0392 0.00642
## 6 1976 White <NA> 0.00285 0.00151
## 7 1976 Black Lt High School 0.558 0.0603
## 8 1976 Black High School 0.335 0.0476
## 9 1976 Black Junior College 0.0423 0.0192
## 10 1976 Black Bachelor 0.0577 0.0238
## # ... with 152 more rows
```

```
out_mrg <- gss_wt %>%
  filter(year %in% seq(1976, 2016, by = 4)) %>%
  mutate(racedeg = interaction(race, degree)) %>%
  group_by(year, racedeg) %>%
  summarize(prop = survey_mean(na.rm = TRUE))
```

```
out_mrg
```

```
## # A tibble: 155 x 4
## # Groups:   year [10]
##   year racedeg      prop prop_se
##   <dbl> <fct>      <dbl> <dbl>
## 1 1976 White.Lt High School 0.297 0.0146
## 2 1976 Black.Lt High School 0.0470 0.00837
## 3 1976 Other.Lt High School 0.00194 0.00138
## 4 1976 White.High School 0.469 0.0159
## 5 1976 Black.High School 0.0282 0.00593
## 6 1976 Other.High School 0.00324 0.00166
## 7 1976 White.Junior College 0.0117 0.00268
```

```
## 8 1976 Black.Junior College 0.00356 0.00162
## 9 1976 White.Bachelor      0.0916 0.00883
## 10 1976 Black.Bachelor     0.00486 0.00213
## # ... with 145 more rows

out_mrg <- gss_wt %>%
  filter(year %in% seq(1976, 2016, by = 4)) %>%
  mutate(racedeg = interaction(race, degree)) %>%
  group_by(year, racedeg) %>%
  summarize(prop = survey_mean(na.rm = TRUE)) %>%
  separate(racedeg, sep = "\\.", into = c("race", "degree"))
```

```
out_mrg
```

```
## # A tibble: 155 x 5
## # Groups:   year [10]
##   year race degree      prop prop_se
##   <dbl> <chr> <chr>    <dbl>  <dbl>
## 1 1976 White Lt High School 0.297 0.0146
## 2 1976 Black Lt High School 0.0470 0.00837
## 3 1976 Other Lt High School 0.00194 0.00138
## 4 1976 White High School 0.469 0.0159
## 5 1976 Black High School 0.0282 0.00593
## 6 1976 Other High School 0.00324 0.00166
## 7 1976 White Junior College 0.0117 0.00268
## 8 1976 Black Junior College 0.00356 0.00162
## 9 1976 White Bachelor 0.0916 0.00883
## 10 1976 Black Bachelor 0.00486 0.00213
## # ... with 145 more rows
```

```
p <- ggplot(data = subset(out_grp, race %nin% "Other"),
  mapping = aes(x = degree, y = prop,
    ymin = prop - 2*prop_se,
    ymax = prop + 2*prop_se,
    fill = race,
    color = race,
    group = race))
```

```
dodge <- position_dodge(width=0.9)
```

```
p + geom_col(position = dodge, alpha = 0.2) +
  geom_errorbar(position = dodge, width = 0.2) +
  scale_x_discrete(labels = scales::wrap_format(10)) +
  scale_y_continuous(labels = scales::percent) +
  scale_color_brewer(type = "qual", palette = "Dark2") +
  scale_fill_brewer(type = "qual", palette = "Dark2") +
  labs(title = "Educational Attainment by Race",
    subtitle = "GSS 1976-2016",
    fill = "Race",
    color = "Race",
    x = NULL, y = "Percent") +
  facet_wrap(~ year, ncol = 2) +
  theme(legend.position = "top")
```

Educational Attainment by Race GSS 1976–2016

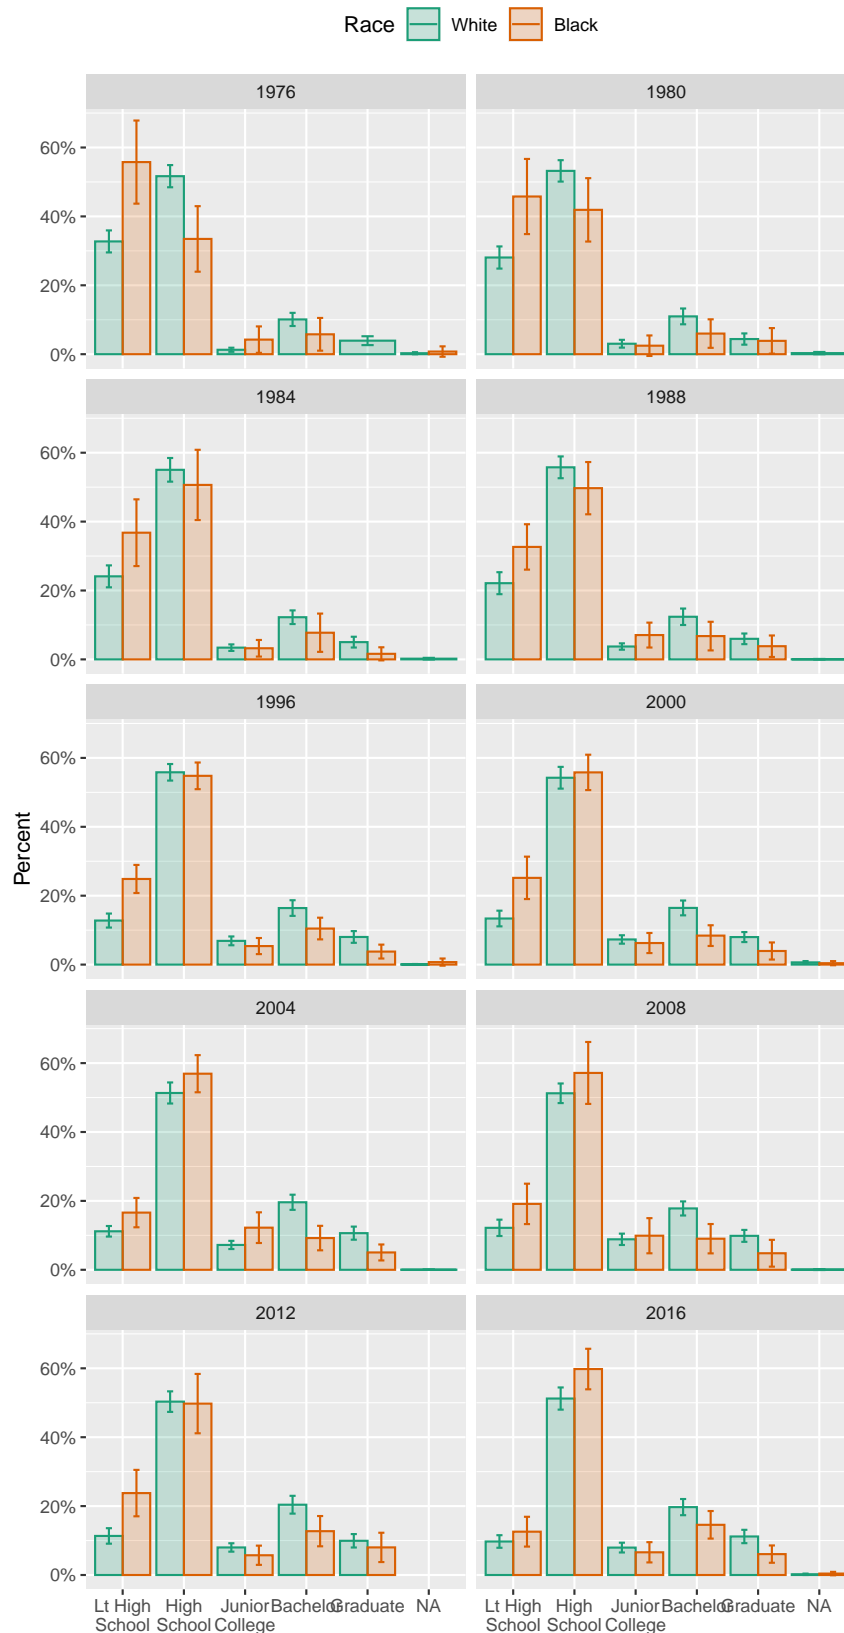


Figure 16: Weighted estimates of educational attainment for Whites and Blacks, GSS selected years 1976–2016. Faceting barplots is often a bad idea, and the more facets there are the worse an idea it is. With a small-multiple plot the viewer wants to compare across panels (in this case, over time), but this is difficult to do when the data inside the panels are categorical comparisons shown as bars (in this case, education level by group).

```

p <- ggplot(data = subset(out_grp, race %nin% "Other"),
  mapping = aes(x = year, y = prop, ymin = prop - 2*prop_se,
    ymax = prop + 2*prop_se, fill = race, color = race,
    group = race))

p + geom_ribbon(alpha = 0.3, aes(color = NULL)) +
  geom_line() +
  facet_wrap(~ degree, ncol = 1) +
  scale_y_continuous(labels = scales::percent) +
  scale_color_brewer(type = "qual", palette = "Dark2") +
  scale_fill_brewer(type = "qual", palette = "Dark2") +
  labs(title = "Educational Attainment\nby Race",
    subtitle = "GSS 1976-2016", fill = "Race",
    color = "Race", x = NULL, y = "Percent") +
  theme(legend.position = "top")

```

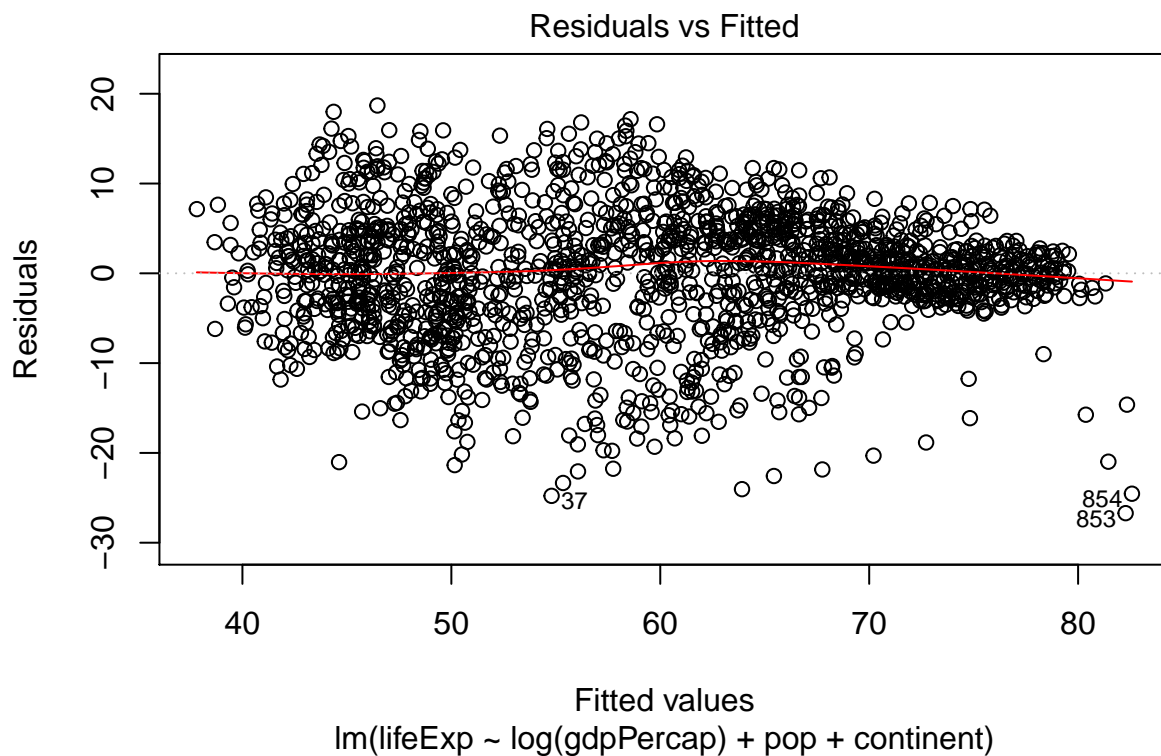
Default plots for models

```

out <- lm(formula = lifeExp ~ log(gdpPercap) + pop + continent, data = gapminder)

plot(out, which = c(1,2), ask=FALSE)

```



Educational Attainment by Race GSS 1976–2016

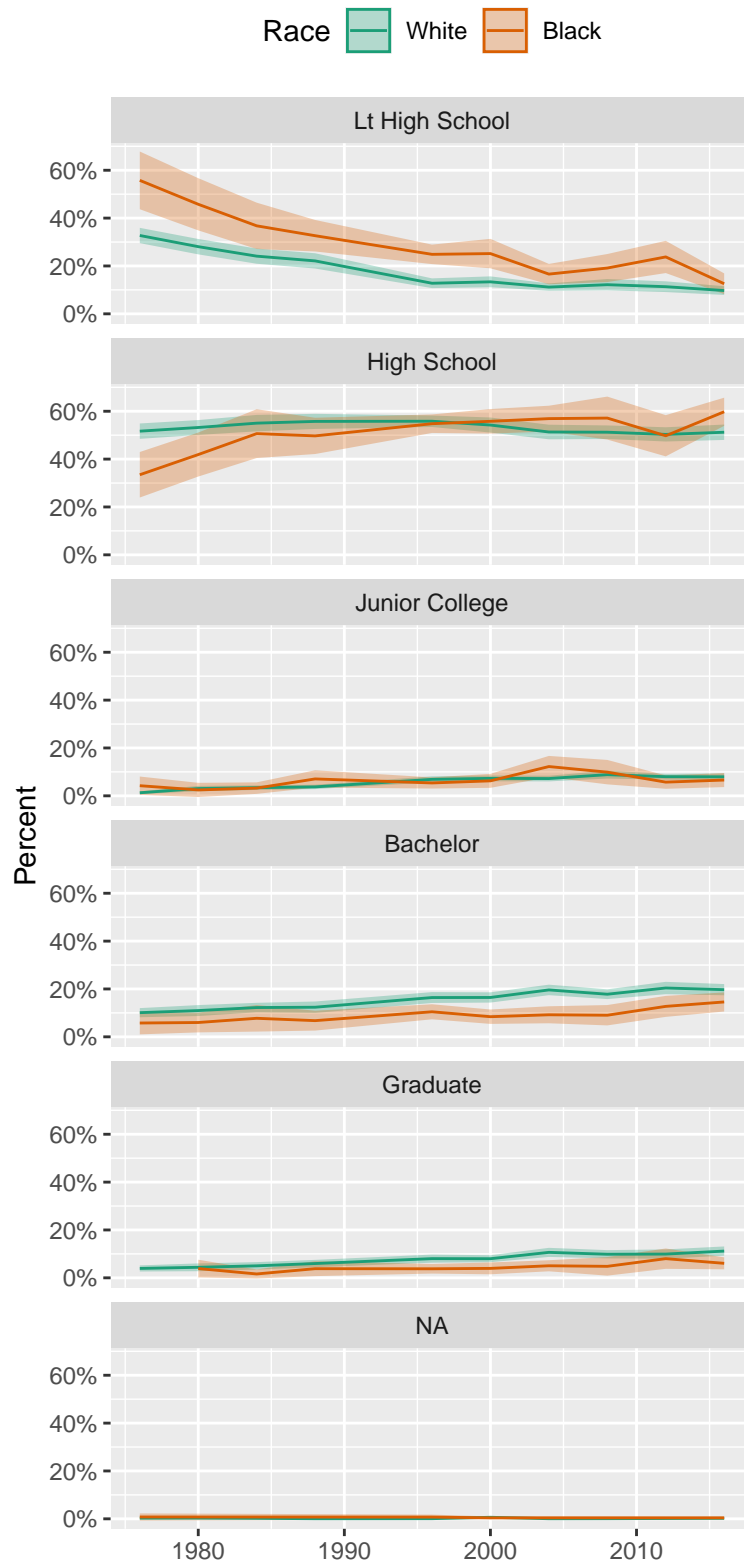


Figure 17: Faceting by education instead.

